

Chapter 1.

UNIX System Overview



朱金辉

csjhzhu@scut.edu.cn

华南理工大学软件学院



1. Introduction

- Tour of UNIX from a programmer's perspective
- Brief descriptions and examples of terms and concepts
- Overview of services provided by UNIX

2. Unix Architecture




Figure 1.1 Architecture of the UNIX operating system



3. Logging In

- Login name
- Password
- /etc/passwd
 - name,
 - encrypted password or “x” (password is in /etc/shadow),
 - numeric user ID,
 - numeric group ID,
 - real name,
 - home directory,
 - shell program



Logging In

- **Shells**: a command line interpreter that reads user input and executes commands

Name	Path	FreeBSD 8.0	Linux 3.2.0	Mac OS X 10.6.8	Solaris 10
Bourne shell	/bin/sh	•	•	copy of bash	•
Bourne-again shell	/bin/bash	optional	•	•	•
C shell	/bin/csh	link to tcsh	optional	link to tcsh	•
Korn shell	/bin/ksh	optional	optional	•	•
TENEX C shell	/bin/tcsh	•	optional	•	•

Figure 1.2 Common shells used on UNIX systems



4. Files and Directories

- **Filesystem**: hierarchical arrangement of directories and files
- Root directory: /
- **File attributes**: type, size, owner, permissions, last modification time, ...
- `stat()`, `fstat()`: return file attribute struct



Files and Directories

- Filename
 - Chars not allowed: (/) and (NULL)
- Two filenames automatically created whenever a new dir is created:
 - . (dot) current directory
 - .. (dot-dot) parent directory
- What is .. in root directory (/)?



Files and Directories

- Pathname
 - A sequence of zero or more filenames, separated by slashes (/), and optionally starting with a slash
- Absolute pathname
- Relative pathname



Program 1.3:

(bare bones implementation of ls command)

```
#include      "apue.h"
#include      <dirent.h>

int main(int argc, char *argv[]){
    DIR      *dp;
    struct dirent *dirp;

    if (argc != 2)
        err_quit("usage: ls directory_name");

    if ( (dp = opendir(argv[1])) == NULL)
        err_sys("can't open %s", argv[1]);
    while ( (dirp = readdir(dp)) != NULL)
        printf("%s\n", dirp->d_name);
    closedir(dp);
    exit(0);
}
```



Program 1.3

- Edit and save in myls.c

cc myls.c (output: a.out)

./a.out /dev (output:)

\$./a.out /etc/ssl/private

can't open /etc/ssl/private: Permission denied

a.out /dev/tty

can't open /dev/tty: Not a directory




5. Input and Output

- File Descriptors
 - small nonnegative integers that kernel uses to identify files being accessed by a process
- Standard Input
- Standard Output
- Standard Error



Input and Output

- `ls`
 - `stdin, stdout, stderr`: ☾ terminal
- `ls > myfile.abc`
 - Stdout: `myfile.abc`
- How to redirect **stderr** to a file?
- How to redirect **stdin** from a file?
- Unbuffered I/O
 - `open()`, `read()`, `write()`, `lseek()`, `close()`



Program 1.4: stdin stdout

```
#include      "apue.h"
#define       BUFSIZE      4096
int main(void) {
    int       n;
    char      buf[BUFSIZE];
    while ( (n = read(STDIN_FILENO, buf, BUFSIZE)) > 0)
        if (write(STDOUT_FILENO, buf, n) != n)
            err_sys("write error");
    if (n < 0)
        err_sys("read error");
    exit(0);
}
```



Standard I/O

- A buffered interface
- No need to worry about BUFSIZE
- Deal with “lines of input”
 - fgets() reads an entire line
 - read() reads a specified # of bytes
- printf() (#include <stdio.h>)



Program 1.5:

stdin stdout using standard I/O

```
#include "apue.h"
int main(void) {
    int          c;
    while ( (c = getc(stdin)) != EOF)
        if (putc(c, stdout) == EOF)
            err_sys("output error");
    if (ferror(stdin))
        err_sys("input error");
    exit(0);
}
```



6. Programs and Processes

- Program: an **executable file** in disk
- Process: an **executing instance** of a program
- Process also called “**task**” by some OS
- Unique nonnegative integer identifier for each process (**pid**)



Program 1.6: process ID

```
#include "apue.h"

int main(void) {
    printf("hello world from process
    ID %d\n", getpid());
    exit(0);
}
```



Process Control

- Three functions
 - `fork()`
 - `exec()`: 6 variants
 - `waitpid()`

Program 1.7: exec stdin cmds

```
#include "apue.h"
#include <sys/wait.h>
int
main(void)
{
    char buf[MAXLINE];      /* from apue.h */
    pid_t pid;
    int      status;

    printf("%% ");          /* print prompt (printf requires %% to print %) */
    while (fgets(buf, MAXLINE, stdin) != NULL) {
        if (buf[strlen(buf) - 1] == '\n')
            buf[strlen(buf) - 1] = 0; /* replace newline with null */

        if ((pid = fork()) < 0) {
            err_sys("fork error");
        } else if (pid == 0) { /* child */
            execlp(buf, buf, (char *)0);
            err_ret("couldn't execute: %s", buf);
            exit(127);
        }

        /* parent */
        if ((pid = waitpid(pid, &status, 0)) < 0)
            err_sys("waitpid error");
        printf("%% ");
    }
    exit(0);
}
```



Threads

- All threads within a process share the same address space, file descriptors, stacks, and process related attributes.
- Each thread executes on its own stack.
- Threads are identified by IDs.



7. Error Handling

- Negative return value when error occurs
- `#include <errno.h>`
- `errno` variable
 - never cleared if error does not occur
 - never set to 0 by any function



Error Handling (contd)

- 2 functions for printing error messages:

```
#include <string.h>  
char *strerror(int errnum);
```

```
#include <stdio.h>  
void perror(const char *msg);
```

- strerror() returns a string
- perror() outputs "msg: <error_msg>"




Program 1.8: use of error func

```
#include      <errno.h>
#include      "apue.h"
```

```
int main(int argc, char *argv[]) {
    fprintf(stderr, "EACCES: %s\n", strerror(EACCES));

    errno = ENOENT;
    perror(argv[0]);

    exit(0);
}
```



Program 1.8: results

\$ a.out

EACCES: Permission denied

a.out: No such file or directory

- (argv[0] passed as arg to perror())



8. User Identification

- User ID: numeric identifier of a user
- Group ID: numeric identifier of a group

```
#include "apue.h"
```

```
int main(void) {  
    printf("uid = %d, gid = %d\n",  
           getuid(), getgid());  
    exit(0);  
}
```



9. Signals

- A technique to **notify** a process that some condition has occurred
- E.g.: divide by zero ☾ SIGFPE
- Process **response** to a signal
 - Ignore the signal, OR
 - Let the default action occur, OR
 - Provide a function to handle the signal.



Signals Example: shell2.c

```
#include "apue.h"
```

```
#include <sys/wait.h>
```

```
static void          sig_int(int);          /* our signal-catching function */
```

```
int main(void) {
```

```
    char  buf[MAXLINE];
```

```
    pid_t pid;
```

```
    int    status;
```

```
    if (signal(SIGINT, sig_int) == SIG_ERR)
```

```
        err_sys("signal error");
```

```
    printf("%0% "); /* print prompt (printf requires %% to print %) */
```

```
    while (fgets(buf, MAXLINE, stdin) != NULL) {
```

```
        buf[strlen(buf) - 1] = 0; /* replace newline with null */
```

```
        if ( (pid = fork()) < 0)
```

```
            err_sys("fork error");
```



Signals (contd)

```
        else if (pid == 0) {                /* child */
            execlp(buf, buf, (char *) 0);
            err_ret("couldn't execute: %s", buf);
            exit(127);
        }
        /* parent */
        if ( (pid = waitpid(pid, &status, 0)) < 0)
            err_sys("waitpid error");
        printf("%0%0 ");
    }
    exit(0);
}
void sig_int(int signo) {
    printf("interrupt\n%0%0 ");
}
```



10. Time Values

- Two different time values
- **Calendar time:** #seconds since the Epoch, which is 00:00:00 Jan 1, 1970, Coordinated Universal Time (UTC).
- **Process time:** measures CPU resources used by a process, in clock ticks, which is 50, 60, or 100 ticks per second.



Time Values (contd)

- **Execution time** of a process has 3 values:
- **clock time**: total amount of time from process start to finish
- **user CPU time**: CPU time due to user instructions in a process
- **system CPU time**: CPU time due to kernel activities on behalf of the process



Time Values (contd)

- To measure process execution time, use the “time” command as follows:

```
time ls > /dev/null
```

```
real      0m19.81s
```

```
user      0m0.43s
```

```
sys       0m4.53s
```



11. System Calls & Library Functions

- **System Calls:**
Entry points into an OS kernel
- Cannot be changed by user
- A function of the same name in the standard C library
- User just calls those C functions whenever system calls are needed

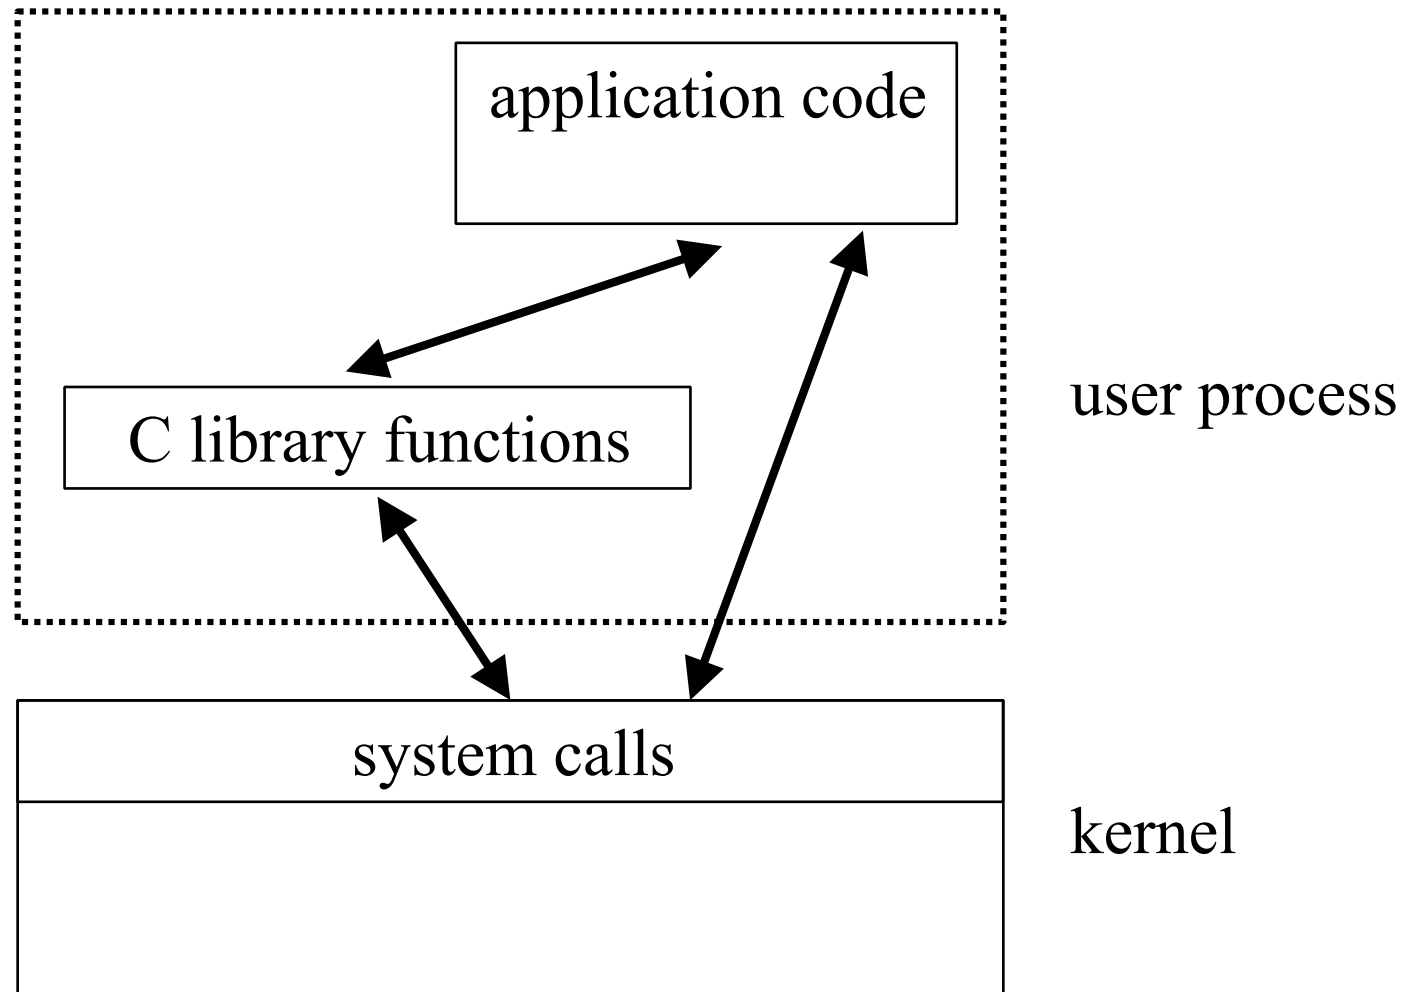


System Calls & Library Functions

- **Library Functions:** not entry points into kernel, just functions, but they may invoke one or more system calls
 - E.g.: `printf()` invokes `write()` system call
 - E.g.: `strcpy()`, `atoi()`: do not invoke any system call
- Implementor view: fundamental diff
- Programmer view: no critical difference

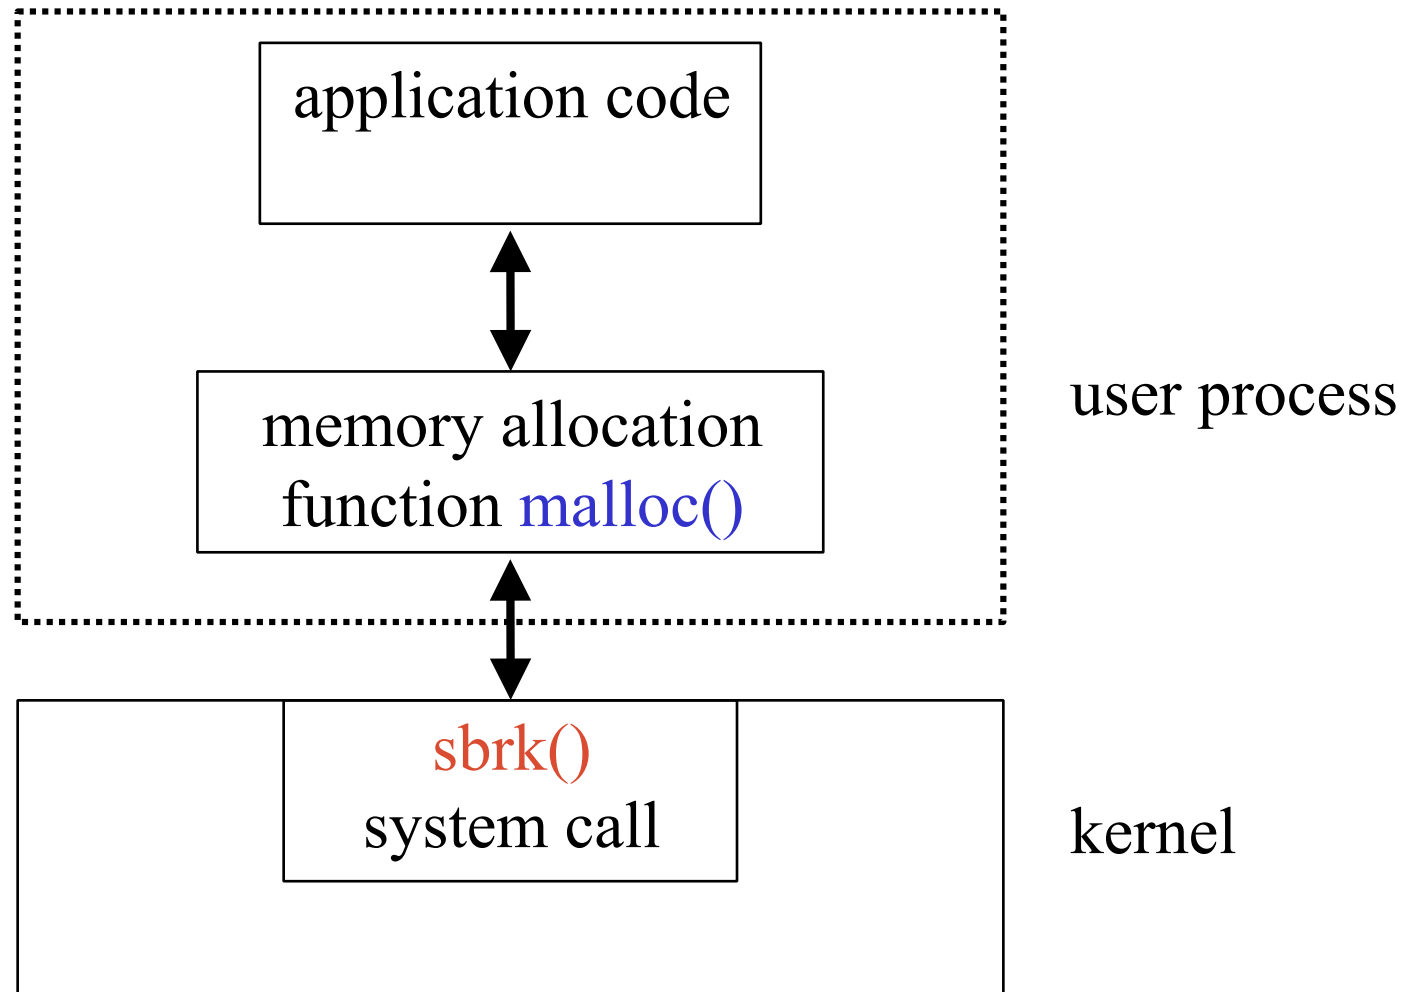


System Calls & Library Functions



System Calls & Library Functions

Example



Example

- execve
- glibc

