

Chapter 5.

Standard I/O Library



朱金辉

华南理工大学软件学院

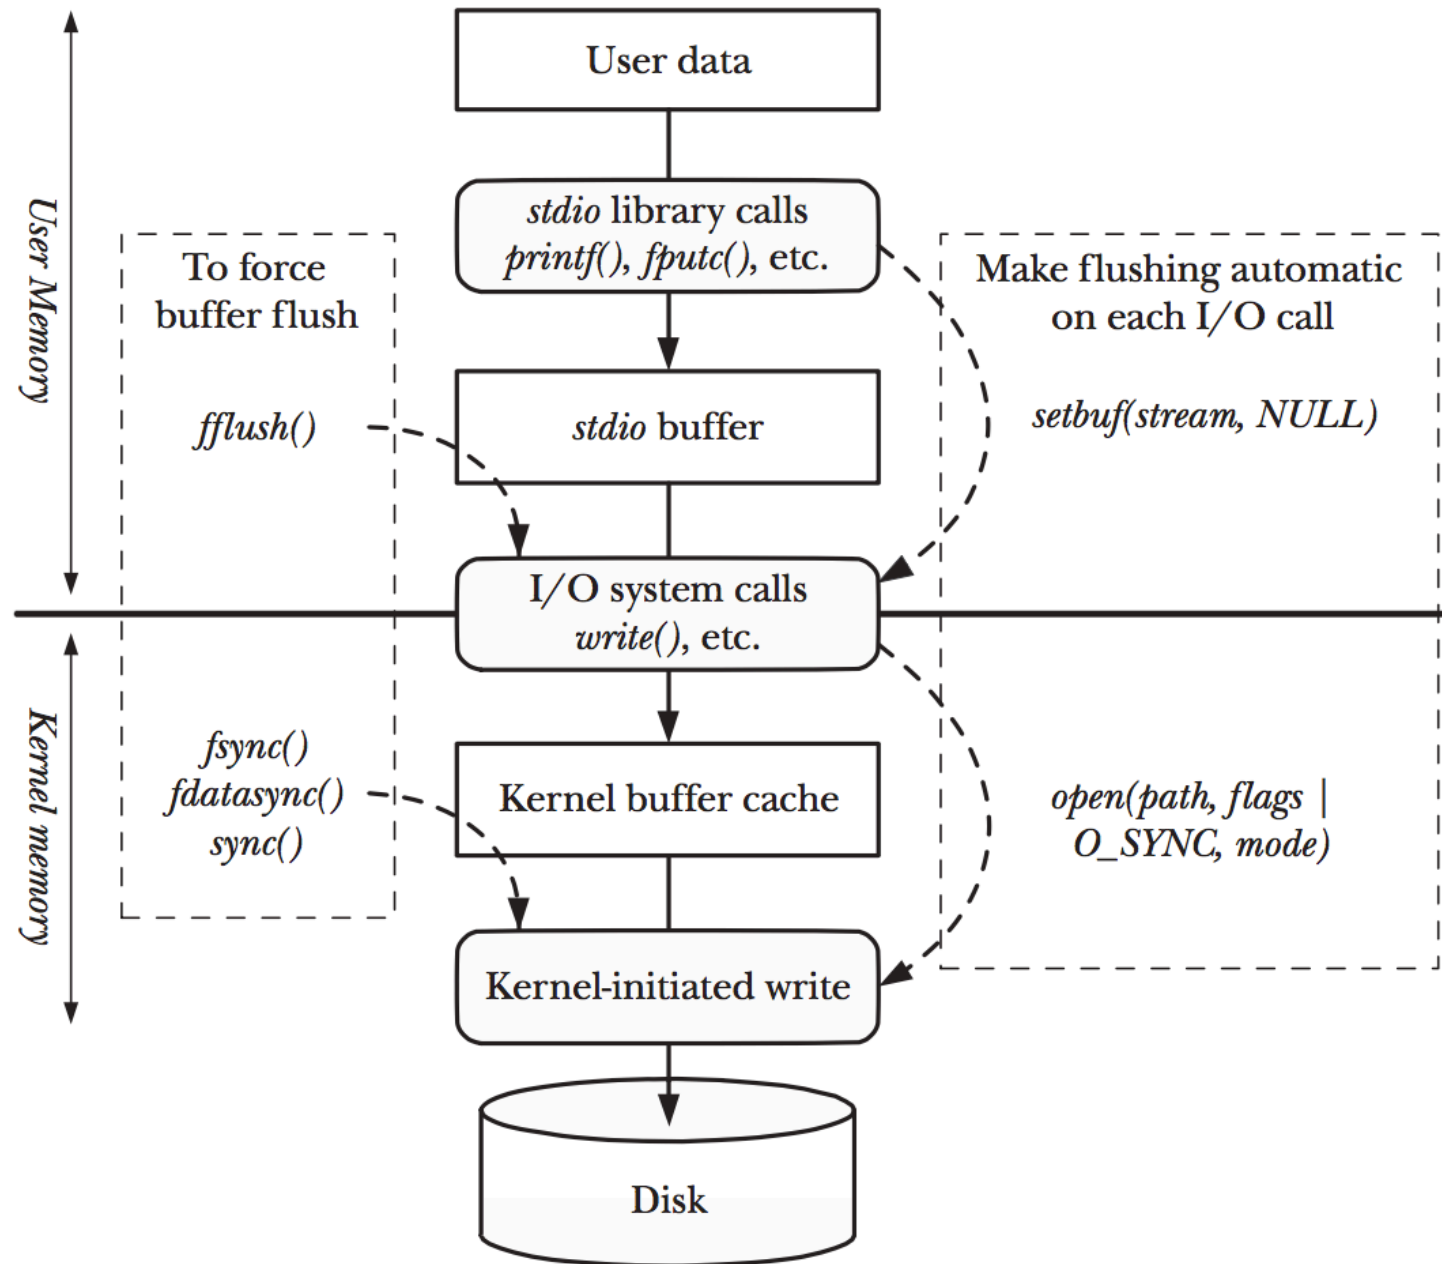


1. Introduction

Standard I/O Library

- Written by Dennis Ritchie in 1975
- Implemented on many OS
- ANSI C standard
- Buffer allocation
- Perform I/O in optimal-sized chunks
- Streams

Summary of I/O buffering





2. Streams

- A stream is associated with a file when we open or create a file using std I/O lib
- `fopen()` returns a **FILE object**
 - file descriptor
 - pointer to stream buffer
 - buffer size
 - #chars in buffer
 - error flag
 - ...



Streams

- Applications need not examine FILE obj
- Pass FILE pointer as **argument** to standard I/O functions
- FILE * is called **file pointer**



3. Standard Input, Output, Error

- 3 streams are automatically created for a process
 - stdin
 - stdout
 - stderr
 - Defined in `<stdio.h>`



4. Buffering

3 types of buffering

- Fully Buffered
- Line Buffered
- Unbuffered



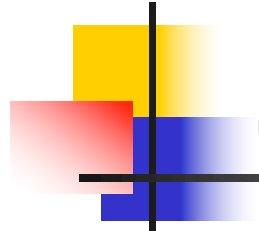
1) Fully Buffered I/O

- Actual I/O takes place when std I/O buffer is FILLED!
- Disk files are fully buffered
- A buffer can be flushed (write out contents of buffer to disk, ..., etc.)
 - `fflush()`: write out buffer contents (std I/O)
 - `tcflush()`: discard buffer data (terminal driver)



2) Line Buffered I/O

- Actual I/O is performed only when a **newline** character is encountered
- Line buffering is used on a stream associated with a **terminal**
- Fixed buffer size → I/O takes place before newline when buffer is full



3) Unbuffered I/O

- Characters unbuffered
- stderr is unbuffered
- hence, error messages are always output very quickly!



Buffering

- ANSI C requirements on buffering:
 - stdin and stdout are fully buffered (if and only if do not refer to an interactive device)
 - stderr is never fully buffered
- SVR4 and 4.3+BSD buffering:
 - stderr is always unbuffered
 - All other streams are line buffered for terminal device; otherwise fully buffered



Buffering

- Type of buffering can be changed:

```
#include <stdio.h>
```

```
void setbuf(    FILE *fp, char *buf    );
```

```
int setvbuf(    FILE *fp, char *buf,  
               int mode, size_t size );
```

- Returns 0 if OK, nonzero on error



setbuf(): turns buffering on or off

- `char *buf:`
 - a buffer of length `BUFSIZ` (`<stdio.h>`)
(Terminal device: line buffered
 otherwise: fully buffered)
 - `NULL` (buffering turned off)



setvbuf(): set buffering type

- `int mode`
 - `_IOFBF`: fully buffered
 - `_IOLBF`: line buffered
 - `_IONBF`: unbuffered
- `char *buf`:
 - (`NULL` → automatic buffer allocation
 - Files: buffer size = `st_blksize` in `stat`
 - Pipes: buffer size = `BUFSIZ`)
- `size_t size`: buffer size



setbuf() v/s setvbuf() (Fig. 5.1)

Function	<i>mode</i>	<i>buf</i>	Buffer and length	Type of buffering
setbuf		non-null	user <i>buf</i> of length <code>BUFSIZ</code>	fully buffered or line buffered
		NULL	(no buffer)	unbuffered
setvbuf	_IOFBF	non-null	user <i>buf</i> of length <i>size</i>	fully buffered
		NULL	system buffer of appropriate length	
	_IOLBF	non-null	user <i>buf</i> of length <i>size</i>	line buffered
		NULL	system buffer of appropriate length	
	_IONBF	(ignored)	(no buffer)	unbuffered

Figure 5.1 Summary of the setbuf and setvbuf functions



Stream Flushing

- A stream can be flushed:

```
#include <stdio.h>
```

```
int fflush(FILE *fp);
```

- Returns: 0 if OK, EOF on error
- Unwritten data passed to kernel
- `fp = NULL` → ALL output streams flushed!



5. Opening a Stream

```
#include <stdio.h>
```

```
FILE *fopen(    const char *pathname,  
               const char *type );
```

```
FILE *freopen(const char *pathname,  
              const char *type, FILE *fp);
```

```
FILE *fdopen(  int fildes,  
             const char *type );
```

- Return: **file pointer** if OK, NULL on error



fopen, freopen, fdopen


- **fopen** opens a specified file
- **freopen** opens a specified file on a specified stream (to open a specified file as stdin, stdout, or stderr)
- **fdopen** associates a stream with a file descriptor (for pipes, network comm channels, which cannot be opened by fopen)



type argument (Fig. 5.2)

<i>type</i>	Description	open(2) Flags
r or rb w or wb a or ab	open for reading truncate to 0 length or create for writing append; open for writing at end of file, or create for writing	O_RDONLY O_WRONLY O_CREAT O_TRUNC O_WRONLY O_CREAT O_APPEND
r+ or r+b or rb+ w+ or w+b or wb+	open for reading and writing truncate to 0 length or create for reading and writing	O_RDWR O_RDWR O_CREAT O_TRUNC
a+ or a+b or ab+	open or create for reading and writing at end of file	O_RDWR O_CREAT O_APPEND


Figure 5.2 The type argument for opening a standard I/O stream



6 ways to open stream (Fig. 5.3)

Restriction	r	w	a	r+	w+	a+
file must already exist	•			•		
previous contents of file discarded		•			•	
stream can be read	•			•	•	•
stream can be written		•	•	•	•	•
stream can be written only at end			•			•

Figure 5.3 Six ways to open a standard I/O stream



To close an open stream

```
#include <stdio.h>
```

```
int fclose(FILE *fp);
```


- Returns: **0** if OK, **EOF** on error
- Before closing:
 - Output buffer is flushed
 - Input buffer is discarded



6. Reading / Writing a Stream

3 types of **Unformatted** I/O:

1. **Character-at-a-time** I/O:
getc(), fgetc(), getchar()
2. **Line-at-a-time** I/O:
fgets(), fputs()
3. **Direct** I/O: fread(), fwrite()
(read/write some number of objects)



Char-at-a-time: Input Functions

```
#include <stdio.h>
```

```
int getc(FILE *fp);
```

```
int fgetc(FILE *fp);
```

```
int getchar(void); (== getc(stdin))
```

- Return: next char if OK, EOF on EOF or error

Usually, getc are implemented as **macros**, but in the GNU C library implementation the macro simply expands to a function call.



File Error or EOF?

- `#include <stdio.h>`
`int ferror(FILE *fp);`
`int feof(FILE *fp);`
- Return: nonzero if true, 0 otherwise
- `void clearerr(FILE *fp);`
- Clears 2 flags in FILE object:
 - an error flag
 - an EOF flag




ungetc()

```
#include <stdio.h>
```

```
int ungetc(int c, FILE *fp);
```

- Returns: c if OK, EOF on error
- Chars are pushed back into stream
- Need not push back same char as read
- Cannot push EOF



Char-at-a-time: Output Functions

```
#include <stdio.h>
```

```
int putc(int c, FILE *fp);
```

```
int fputc(int c, FILE *fp);
```

```
int putchar(int c); (== putc(c, stdout))
```

- Return: c if OK, EOF on error



7. Line-at-a-Time: Input Functions

```
#include <stdio.h>
```

```
char *fgets(char *buf, int n, FILE *fp);
```


deprecated

(stores: **line** + **\n** + **NULL**)

```
char *gets(char *buf);
```

(stores **line** only, without **\n** and **NULL**)

- Return: buf if OK, NULL on EOF or error




Line-at-a-time: Output Functions

```
#include <stdio.h>
```

```
int fputs(const char *str, FILE *fp);
```

```
int puts(const char *str);
```



- Return: nonnegative value if OK,
EOF on error



Line-at-a-time I/O

Suggestion:

- Use:
 - `fgets()` and `fputs()`
- Remember:
 - we always have to deal with the **newline character** at the end of each line!



8. Standard I/O Efficiency

- How does the three types of unformatted I/O compare in efficiency?
- How does unformatted I/O compare with direct read() and write()?
- Are macros more efficient than functions? Why?



stdin → stdout (getc, putc) Program 5.4

```
#include    "apue.h"  
int main(void) {  
    int      c;  
    while ( (c = getc(stdin)) != EOF)  
        if (putc(c, stdout) == EOF)  
            err_sys("output error");  
    if (ferror(stdin))  
        err_sys("input error");  
    exit(0);  
}
```



stdin → stdout (fgets(), fputs)

Program 5.5

```
#include    "apue.h"
int main(void) {
    char    buf[MAXLINE];
    while (fgets(buf, MAXLINE, stdin) != NULL)
        if (fputs(buf, stdout) == EOF)
            err_sys("output error");
    if (ferror(stdin))
        err_sys("input error");
    exit(0);
}
```



Timing Results using std I/O

Function	User CPU (seconds)	System CPU (seconds)	Clock time (seconds)	Bytes of program text
best time from Figure 3.6	0.05	0.29	3.18	
fgets, fputs	2.27	0.30	3.49	143
getc, putc	8.45	0.29	10.33	114
fgetc, fputc	8.16	0.40	10.18	114
single byte time from Figure 3.6	134.61	249.94	394.95	

Figure 5.6 Timing results using standard I/O routines

```
→ fileio git:(master) x time ./test_fileio 1 <~/test.data >/dev/null
./test_fileio 1 < ~/test.data > /dev/null 18.84s user 139.50s system 98% cpu 2:40.01 total
→ fileio git:(master) x time ./test_fileio 1024 <~/test.data >/dev/null
./test_fileio 1024 < ~/test.data > /dev/null 0.02s user 0.21s system 84% cpu 0.275 total
→ fileio git:(master) x time ./test_fileio 4096 <~/test.data >/dev/null
./test_fileio 4096 < ~/test.data > /dev/null 0.00s user 0.11s system 70% cpu 0.158 total
→ fileio git:(master) x time ./test_fileio 524288 <~/test.data >/dev/null
./test_fileio 524288 < ~/test.data > /dev/null 0.00s user 0.08s system 60% cpu 0.139 total

→ stdio git:(master) x time ./getcputc < ~/test.data > /dev/null
./getcputc < ~/test.data > /dev/null 6.79s user 0.14s system 93% cpu 7.378 total
→ stdio git:(master) x time ./fgetsfputs < ~/test.data > /dev/null
./fgetsfputs < ~/test.data > /dev/null 0.09s user 0.11s system 75% cpu 0.260 total
```



9. Binary I/O

```
#include <stdio.h>
```

```
size_t fread( void *ptr, size_t size,  
              size_t nobj, FILE *fp);
```

```
size_t fwrite(const void *ptr, size_t size,  
              size_t nobj, FILE *fp);
```

- Return: #objects read or written



Binary I/O

- Read or write an **entire structure!**
- Cannot use `fgets()` or `fputs()`
- There may be NULLs or newlines in a structure
- Hence, **`fread()`** and **`fwrite()`** are required!



Binary I/O: Read/Write Binary Array

- Write elements 2 through 5 of a floating point array:

```
float data[10]
```

```
if (fwrite(&data[2], sizeof(float), 4, fp) != 4)
```

```
    err_sys( "fwrite error" );
```



Binary I/O: Read/Write a Structure

```
struct {  
    short count;  
    long total;  
    char name[NAMESIZE];  
} item;  
  
if (fwrite(&item, sizeof(item), 1, fp) != 1)  
    err_sys( "fwrite error" );
```

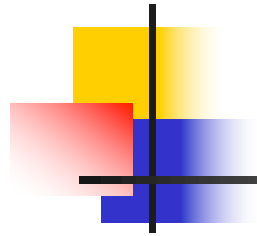


Binary Data Exchange

May be **incompatible** among different machine architectures because:

- **Offset** of a member in a structure can differ between compilers and systems
- **Binary formats** of int and float may differ between different machines

So use a higher level protocol instead!



10. Positioning a Stream

Three ways:

1. `ftell` and `fseek`
2. `ftello` and `fseeko`
3. `fgetpos` and `fsetpos`
 - more portable (ANSI)
 - new data type: `fpos_t`



ftell() and fseek()

```
#include <stdio.h>
```

```
long ftell (FILE *fp);
```

- Returns: curr filepos indicator if OK,
-1L on error

```
int fseek (FILE *fp, long offset,  
int whence);
```

- Returns: 0 if OK, nonzero on error

```
void rewind (FILE *fp);
```




fgetpos() and fsetpos()

```
#include <stdio.h>
```

```
int fgetpos(FILE *fp, fpos_t *pos);
```

```
int fsetpos(FILE *fp, const fpos_t *pos);
```

Return: 0 if OK, nonzero on error



11. Formatted Output

```
#include <stdio.h>
```

```
int printf(const char *format, ...);
```

```
int fprintf(FILE *fp, const char *format, ...);
```

- Return: #char output if OK, <0 on error

```
int sprintf(char *buf, const char *format, ...);
```

- Returns: #char stored in array (without NULL)
- NULL appended to buf



Formatted Output (variants)

```
#include <stdarg.h>
```

```
#include <stdio.h>
```

```
int vprintf(const char *format, va_list arg);
```

```
int fprintf(FILE *fp, const char *format,  
            va_list arg);
```

- Return: #char output if OK, <0 on error

```
int vsprintf(char *buf, const char *format,  
            va_list arg);
```

- Return: #char stored in array



Formatted Input


```
#include <stdio.h>
```

```
int scanf(const char *format, ...);
```

```
int fscanf(FILE *fp, const char *format, ...);
```

```
int sscanf(const char *buf,  
           const char *format, ...);
```

Return: #items assigned, EOF on error or on
EOF



12. Implementation Details

```
#include <stdio.h>
```

```
int fileno(FILE *fp);
```

- Returns: file descriptor from stream
- Needed by `dup()` and `fcntl()` functions



Print stdio buffering (Program 5.11)

```
#include "apue.h"
void      pr_stdio(const char *, FILE *);
int
main(void)
{
    FILE *fp;

    fputs("enter any character\n", stdout);
    if (getchar() == EOF)
        err_sys("getchar error");
    fputs("one line to standard error\n", stderr);

    pr_stdio("stdin", stdin);
    pr_stdio("stdout", stdout);
    pr_stdio("stderr", stderr);
```

```

if ( (fp = fopen("/etc/motd", "r")) == NULL)
    err_sys("fopen error");
if (getc(fp) == EOF)
    err_sys("getc error");
pr_stdio("/etc/motd", fp);
exit(0);
}

```

void

```

pr_stdio(const char *name, FILE *fp)
{
    printf("stream = %s, ", name);
        /* following is nonportable */
    if (fp->_flag & _IONBF)        printf("unbuffered");
    else if (fp->_flag & _IOLBF)    printf("line buffered");
    else /* if neither of above */ printf("fully buffered");
    printf(", buffer size = %d\n", fp->_bufsiz);
}

```



Program 5.11: Output (1)

\$ a.out

enter any character

one line to standard error

stream = stdin, line buffered, buffer size = 128

stream = stdout, line buffered, buffer size = 128

stream = stderr, unbuffered, buffer size = 8

stream = /etc/motd, fully buffered, buffer size=8192



Program 5.11: Output (2)

```
$ a.out < /etc/termcap > std.out 2> std.err
```

```
$ cat std.err
```

one line to standard error

```
$ cat std.out
```

enter any character

stream=stdin, fully buffered, buffer size=8192

stream=stdout, fully buffered, buffer size=8192

stream=stderr, unbuffered, buffer size=8

stream=/etc/motd, fully buffered, buffer size=8192



13. Temporary Files

- NULL → store in static area
- ptr = user-given buffer

```
#include <stdio.h>
```

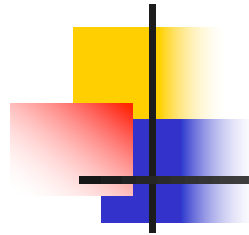
```
char *tmpnam(char *ptr);
```

- Returns: pointer to unique pathname

```
FILE *tmpfile(void);
```

- Returns: file pointer if OK, NULL on error

wb+
(create for
read/write)



tmpfile()

UNIX technique for tmpfile():

- Call **tmpnam()** to create a unique pathname
- **Create** the file
- **Unlink** it immediately

(unlinking is not deleting immediately)



Program 5.12: tmpnam, tmpfile

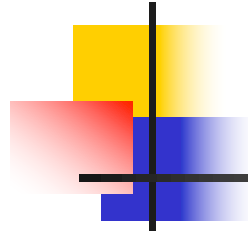
```
#include "apue.h"

int main(void) {
    char name[L_tmpnam], line[MAXLINE];
    FILE *fp;

    printf("%s\n", tmpnam(NULL));    /* first temp name */

    tmpnam(name);                    /* second temp name */
    printf("%s\n", name);

    if ( (fp = tmpfile()) == NULL)   /* create temp file */
        err_sys("tmpfile error");
    fputs("one line of output\n", fp); /* write to temp file */
    rewind(fp);                       /* then read it back */
    if (fgets(line, sizeof(line), fp) == NULL)
        err_sys("fgets error");
    fputs(line, stdout);              /* print the line we wrote */
    exit(0);
}
```



Program 5.12: Output

\$ a.out

/usr/tmp/fileC1Icwc

/usr/tmp/filemSkHSe

one line of output