

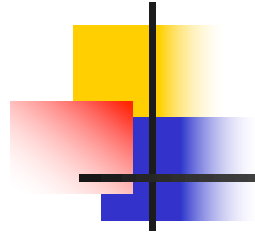
# Chapter 7.

# Unix Process Environment

---

朱金辉

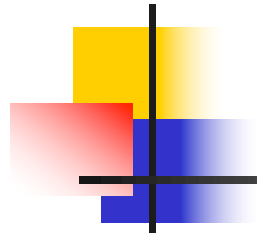
华南理工大学软件学院



# 1. Introduction

---

- How is main() called?
- How are arguments passed?
- Memory layout?
- Memory allocation?
- Environment variables
- Process termination



## 2. main Function

---

- `int main(int argc, char *argv[]);`
- `arc = #arguments`
- `argv[] = arguments`
- Kernel executes a special **START-UP routine** before `main()`
- Start-up routine sets things up before `main()` is called: stack, heap, etc.



## 3. Process Termination

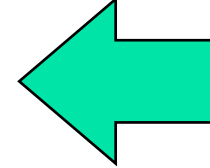
---

- Normal termination:
  - return from main()
  - calling exit()
  - calling \_exit()
  - Return of the last thread from its start routine (Section 11.5)
  - Calling pthread\_exit (Section 11.5) from the last thread
- Abnormal termination
  - calling abort() (Section 11.7)
  - terminated by a signal (Section 10.2)
  - Response of the last thread to a cancellation request (Sections 11.5 and 12.7)



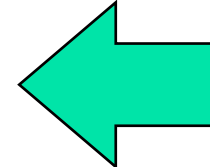
# exit() and \_exit()

- #include <stdlib.h>
- void exit(int *status*);



fclose()  
all open  
streams

- #include <unistd.h>
- void \_exit(int *status*);



return  
to kernel  
immediately

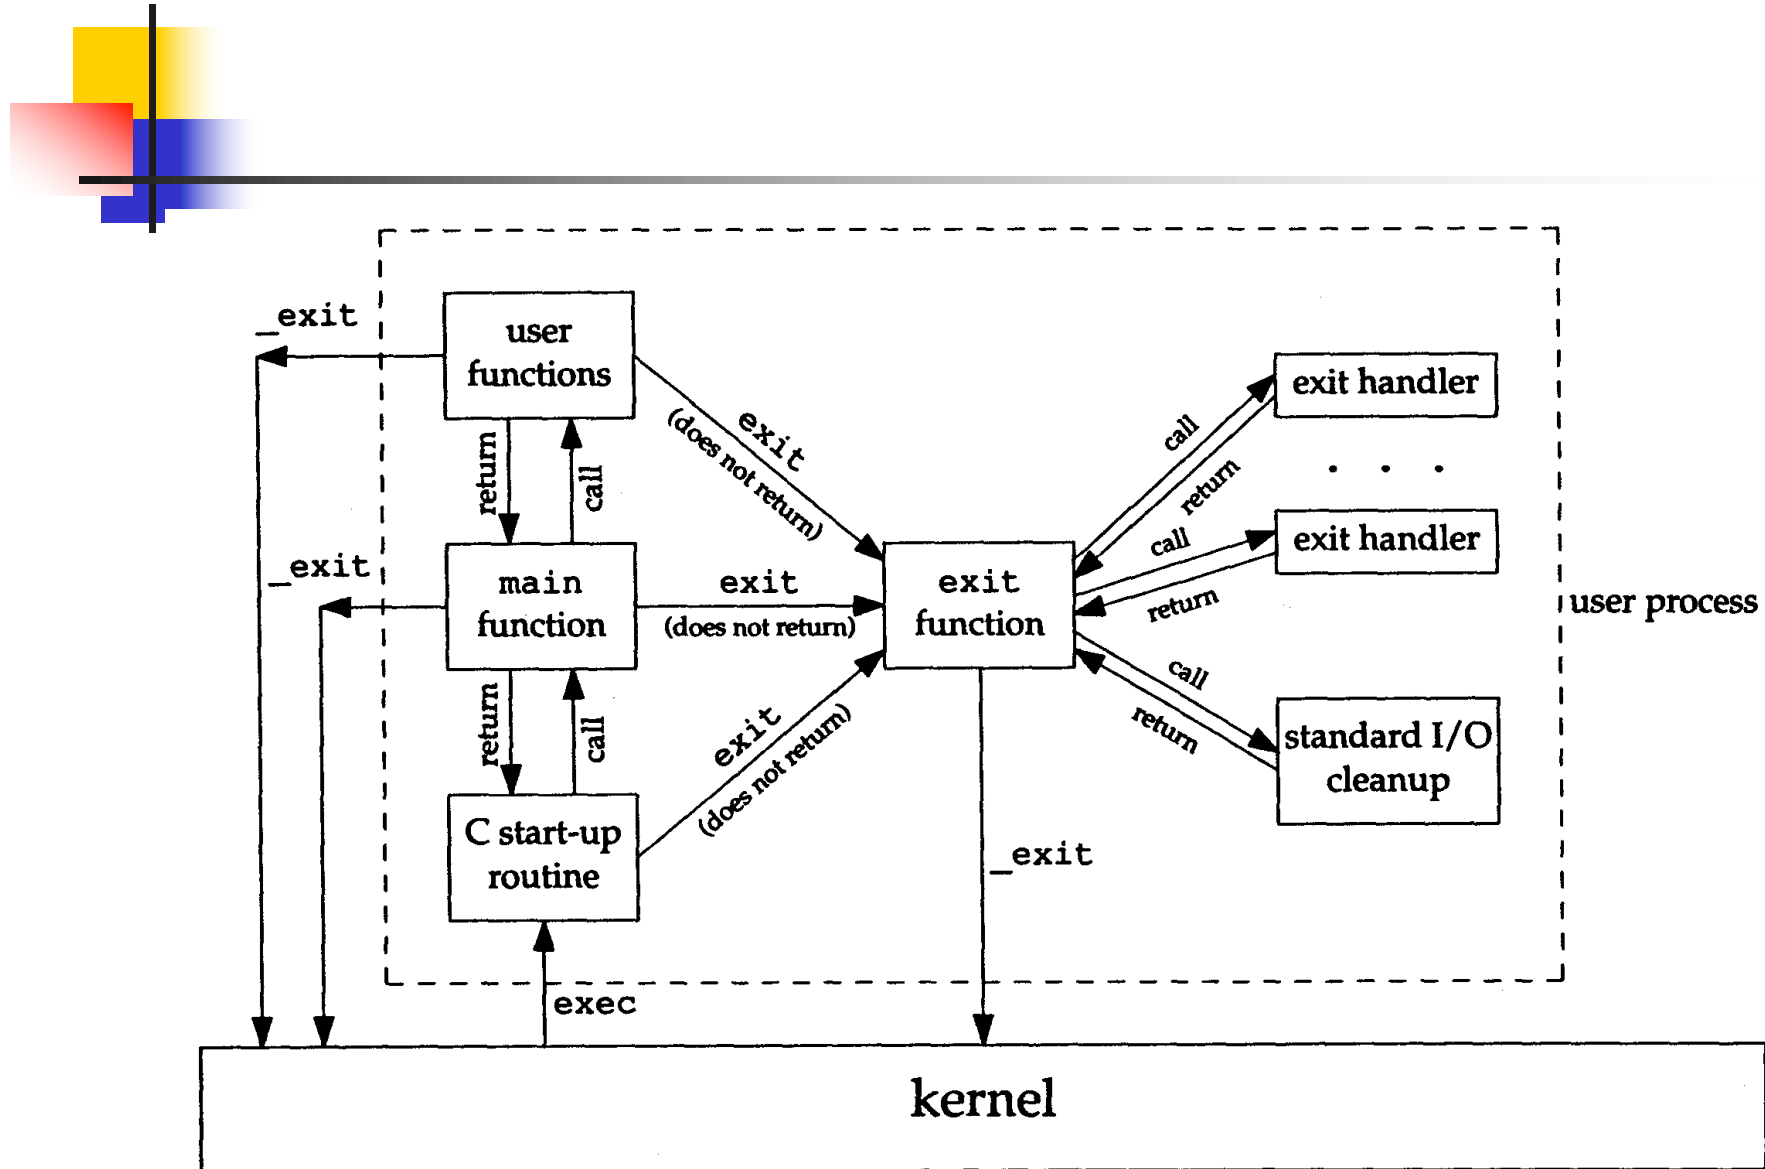


## atexit(): Exit Handler

---

- `#include <stdlib.h>`
- `int atexit(void (*func) (void));`
- Returns: 0 if OK, nonzero on error
- *func* is an exit handler
- `exit()` calls these exit handler functions in the **reverse order** of registration
- `#times called = #times registered`

# Program Start & Termination





## Program 7.3: Exit Handlers

---

```
#include          "apue.h"
static void      my_exit1(void), my_exit2(void);
int main(void) {
    if (atexit(my_exit2) != 0)
        err_sys("can't register my_exit2");
    if (atexit(my_exit1) != 0)
        err_sys("can't register my_exit1");
    if (atexit(my_exit1) != 0)
        err_sys("can't register my_exit1");
    printf("main is done\n");
    return(0);
}
static void my_exit1(void) {
    printf("first exit handler\n");
}
static void my_exit2(void) {
    printf("second exit handler\n");
}
```






## Program 7.3: results

---

- **\$ a.out**
- main is done
- first exit handler
- first exit handler
- second exit handler



## 4. Command-Line Arguments

---

- `exec()` can pass command-line arguments to a new program
- Part of normal operation of Unix shells
- `argv[argc]` is NULL (ANSI, POSIX.1)




## Program 7.4: echo()

---

```
#include    "apue.h"

int
main(int argc, char *argv[])
{
    int      i;

    for (i = 0; i < argc; i++)
        /* echo all command-line args */
        printf("argv[%d]: %s\n", i, argv[i]);
    exit(0);
}
```



## Program 7.4: results

---

- `$ ./echoarg arg1 TEST foo`
- `argv[0]: ./echoarg`
- `argv[1]: arg1`
- `argv[2]: TEST`
- `argv[3]: foo`



# getopt

---

- `#include <unistd.h>`
- `int getopt(int argc, char * const argv[],  
              const char *options);`
- `extern int optind, opterr, optopt;`
- `extern char *optarg;`
- Returns: the next option character, or  
          -1 when all options have been processed
- Example:
  - `command [-i] [-u username] [-z] filename`
  - pass "iu:z" as the *options*

```

1  #include <unistd.h>
2  #include <stdlib.h>
3  #include <stdio.h>
4
5  int main(int argc, char *argv[])
6  {
7      int flags, opt;
8      int nsecs, tfnd;
9
10     nsecs = 0;
11     tfnd = 0;
12     flags = 0;
13     while ((opt = getopt(argc, argv, "nt:")) != -1)
14     {
15         switch (opt)
16         {
17             case 'n':
18                 flags = 1;
19                 break;
20             case 't':
21                 nsecs = atoi(optarg);
22                 tfnd = 1;
23                 break;
24             default: /* '?' */
25                 fprintf(stderr, "Usage: %s [-t nsecs] [-n] name\n",
26                     argv[0]);
27                 exit(EXIT_FAILURE);
28         }
29     }
30     printf("name argument = %s\n", argv[optind]);
31     /* Other code omitted */
32     exit(EXIT_SUCCESS);
33 }

```

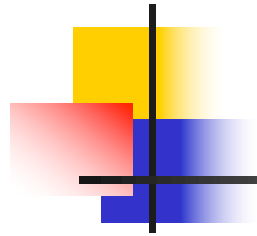
- **Optarg**: If an option takes an argument, getopt sets optarg to point to the option's argument string when an option is processed.
- **Optind**: The index in the argv array of the next string to be processed. It starts at 1 and is incremented for each argument processed by getopt.



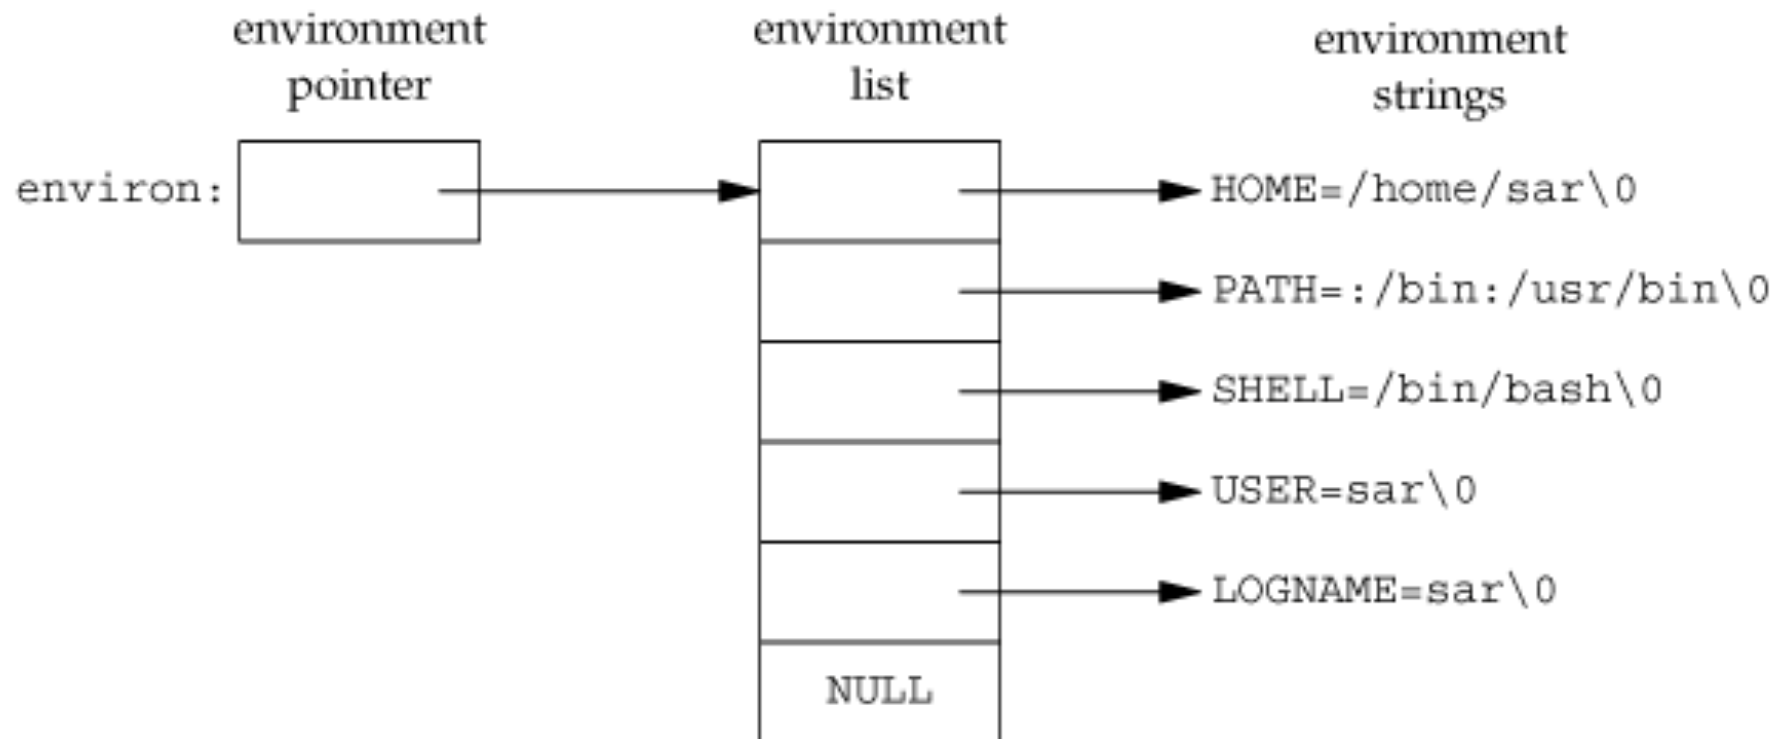
## 5. Environment List

---

- An array of character pointers (null-terminated C strings)
- Array address is in global variable `environ`
- `extern char **environ;`
- `getenv()`: get an environment string
- `putenv()`: set an environment string



# Environment List (Fig. 7.5)








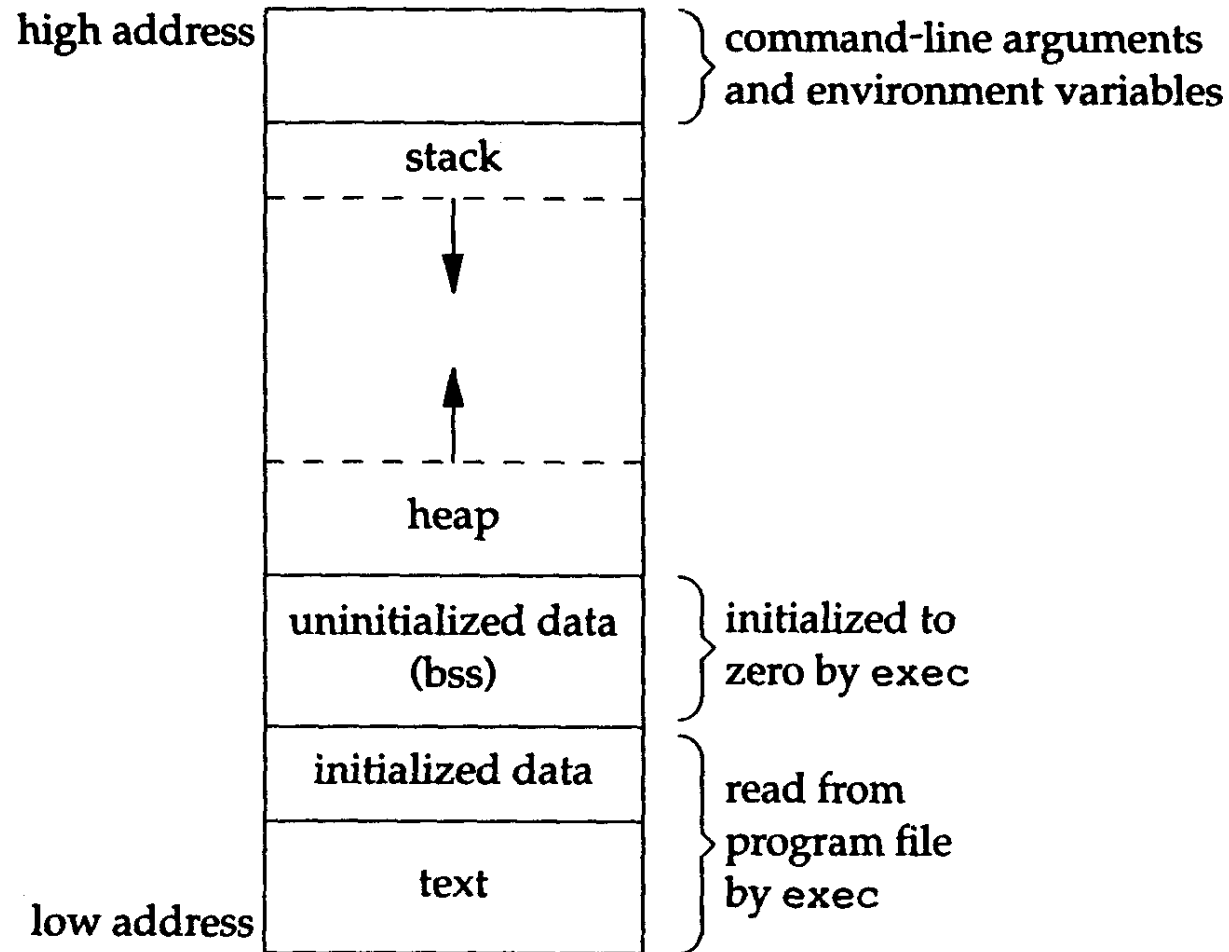
## 6. Memory Layout of a C Program


---

- **Text segment:** Machine instructions  
(read-only, sharable)
- **Initialized data segment:**  
e.g. `int maxcount = 99;` (initialized!)
- **Uninitialized data segment:**  
(bss: block started by symbol)  
e.g. `long sum[1000];`
- **Stack:** automatic variables, function calling information, context-switch information,  
(recursive functions)
- **Heap:** dynamic memory allocation



# Memory Layout (Fig. 7.6)





# size

---

```
$ size /bin/cc /bin/sh
```

text	data	bss	dec	hex	filename
81920	16384	664	98968	18298	/bin/cc
90112	16384	0	106496	1a000	/bin/sh



## 7. Shared Libraries

---

- Common library routines removed from executable files
- Single copy of common library routines in memory is maintained
- No need to re-link edit every program if a library is updated or changed
- Size is smaller, some run-time overhead



# Shared Libraries

Without Shared  
Libraries

```
$ ls -l a.out
```

```
-rwxrwxr-x 1 stevens 104859 Aug 2 14:25 a.out
```

```
$ size a.out
```

text	data	bss	dec	hex
49152	49152	0	98304	18000

```
$ ls -l a.out
```

```
-rwxrwxr-x 1 stevens 24576 Aug 2 14:26 a.out
```

```
$ size a.out
```

text	data	bss	dec	hex
8192	8192	0	16384	4000

With  
Shared  
Libraries



## 8. Memory Allocation

---

- malloc():
  - allocates specified **#bytes**,
  - initial value of memory is indeterminate
- calloc():
  - allocates specified **#objects** of specified size,
  - initialized to all **0** bits
- realloc():
  - changes size of previously allocated memory,
  - initial value of new area is indeterminate



# Memory Allocation

---

```
#include <stdlib.h>
```

```
void *malloc(size_t size);
```

```
void *calloc(size_t nobj, size_t size);
```

```
void *realloc(void *ptr, size_t newsize);
```

Return: nonnull pointer if OK,  
NULL on error

```
void free(void *ptr);
```

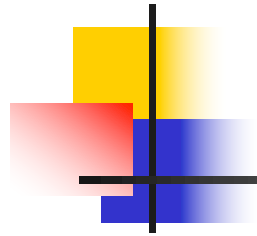


# alloca()

---

- Allocates memory from **stack**, instead of heap
- **Advantage**: No need to free space, automatically freed after function returns
- **Disadvantage**: Some systems do not support `alloca()`





## 9. Environment Variables

---

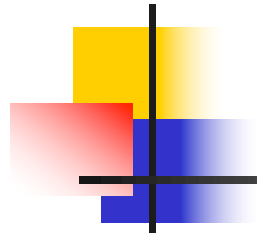
```
#include <stdlib.h>
```

```
char *getenv(const char *name);
```

- Returns: pointer to value associated with name, NULL if not found
- Some environment variables are set automatically by shell upon login
- E.g.: HOME, USER, etc.

# Environment Variables (Fig. 7.7)

Variable	POSIX.1	FreeBSD 5.2.1	Linux 2.4.22	Mac OS X 10.3	Solaris 9	Description
COLUMNS	•	•	•	•	•	terminal width
DATETIME	XSI		•		•	getdate(3) template file pathname
HOME	•	•	•	•	•	home directory
LANG	•	•	•	•	•	name of locale
LC_ALL	•	•	•	•	•	name of locale
LC_COLLATE	•	•	•	•	•	name of locale for collation
LC_CTYPE	•	•	•	•	•	name of locale for character classification
LC_MESSAGES	•	•	•	•	•	name of locale for messages
LC_MONETARY	•	•	•	•	•	name of locale for monetary editing
LC_NUMERIC	•	•	•	•	•	name of locale for numeric editing
LC_TIME	•	•	•	•	•	name of locale for date/time formatting
LINES	•	•	•	•	•	terminal height
LOGNAME	•	•	•	•	•	login name
MSGVERB	XSI	•			•	fmtmsg(3) message components to process
NLSPATH	XSI	•	•	•	•	sequence of templates for message catalogs
PATH	•	•	•	•	•	list of path prefixes to search for executable file
PWD	•	•	•	•	•	absolute pathname of current working directory
SHELL	•	•	•	•	•	name of user's preferred shell
TERM	•	•	•	•	•	terminal type
TMPDIR	•	•	•	•	•	pathname of directory for creating temporary files
TZ	•	•	•	•	•	time zone information



## Setting an environment variable

---

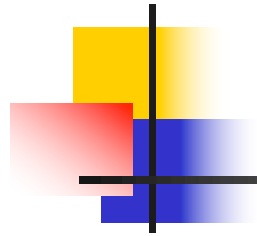
```
#include <stdlib.h>
```

```
int putenv(const char *str);
```

```
int setenv(const char *name, const char  
           *value, int rewrite);
```

- Return: 0 if OK, nonzero on error

```
void unsetenv(const char *name);
```



## 10. setjmp(), longjmp() Functions

---

- In C, we cannot `goto` a label in another function
- `setjmp()` and `longjmp()` must be used
- See Program 7.9 (a skeleton) for command processing
  - read commands,
  - determine commands
  - call functions to process each command




# cmd1.c

---


```
#define TOK_ADD 5
void do_line(char *);
void cmd_add (void);
int get_token(void);

int main(void) {
    char line[MAXLINE];
    while (fgets(line,MAXLINE,stdin )!=NULL)
        do_line(line);
    exit(0);
}

char *tok_ptr;
```



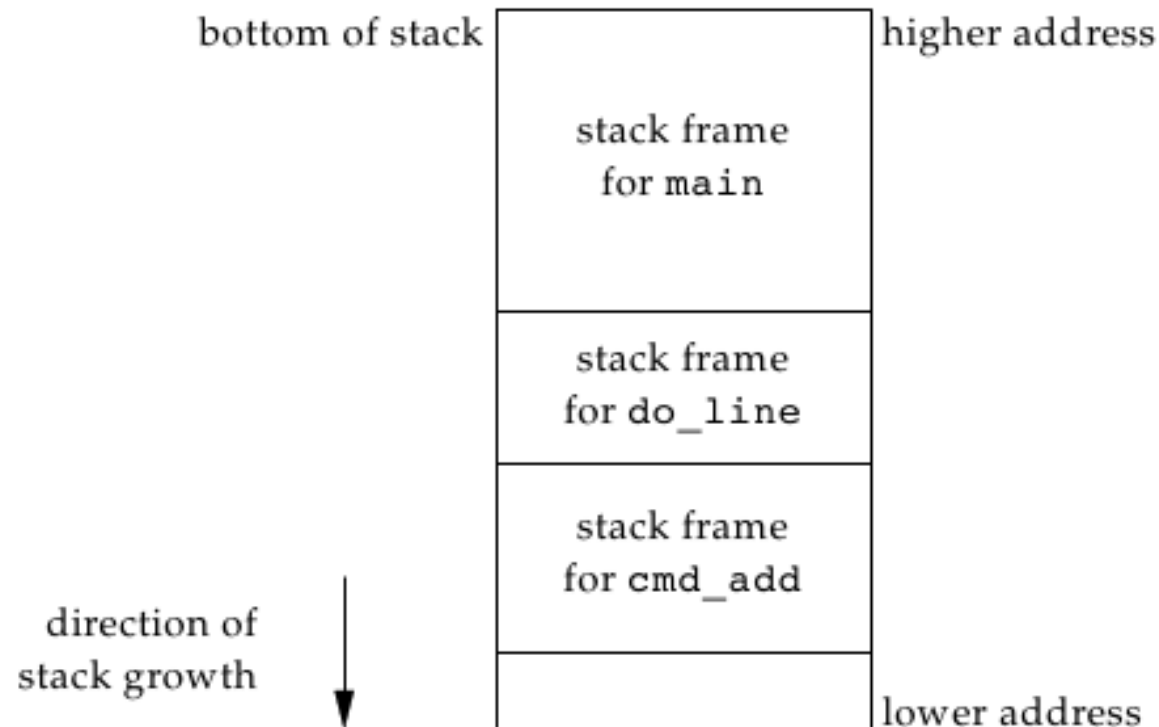
```
do_line(char *ptr) {  
    int cmd;  
    tok_ptr=ptr;  
    while ((cmd=get_token())>0) {  
        switch (cmd) {  
        case TOK_ADD:  
            cmd_add();  
            break;  
        }  
    }  
}
```



```
void cmd_add(void) {  
    int token;  
    token=get_token();  
    /* rest code */  
}  
  
int get_token(void) {  
    /*fetch next token */  
}
```



# After cmd\_add(): stack frame







# setjmp() and longjmp()

---

- Often we are deeply nested,
- An error occurs,
- We want to print an error, ignore rest of input, and return to main()
- Large # of levels → handle return at each level for each error
- Direct nonlocal goto: setjmp, longjmp



# setjmp() and longjmp()

---


- `#include <setjmp.h>`
- `int setjmp(jmp_buf env);`
- Returns: 0 if called directly, nonzero if returning from a call to longjmp
- `void longjmp(jmp_buf env, int val);`



## Program 7.11

---

- `setjmp(jmpbuffer)` stores current state of main at the start of program exec
- `longjmp(jmpbuffer, 1)` unwounds the stacks of `do_line()` and `cmd_add()`
- and causes `setjmp()` to return 1



## Automatic, Register, Volatile Variables

---

- After `longjmp()`, what are the values of the automatic and register variables?
  - Rolled back
  - Left alone
- Standards: indeterminate
- Volatile variables: don't rollback values
- Global, static variables: leave alone

## Program 7.13: longjmp() ...

```
#include "apue.h"
#include <setjmp.h>
static void f1(int, int, int, int);
static void f2(void);
static jmp_buf jmpbuffer;
static int globval;
int main(void) {
    int autoval; register int regival; volatile int volaval; static int statval;
    globval = 1; autoval = 2; regival = 3; volaval = 4; statval = 5;
    if (setjmp(jmpbuffer) != 0) {
        printf("after longjmp:\n");
        printf("globval = %d, autoval = %d, regival = %d," " volaval = %d, statval = %d\n", globval, autoval, regival, volaval, statval);
        exit(0);
    } /* * Change variables after setjmp, but before longjmp. */
}
```

## Program 7.13: longjmp() ...

```
globval = 95; autoval = 96; regival = 97; volaval = 98; statval = 99;
    f1(autoval, regival, volaval, statval); /* never returns */
    exit(0);
}
static void f1(int i, int j, int k, int l) {
    printf("in f1():\n");
    printf("globval = %d, autoval = %d, regival = %d," " volaval = %d,
    statval = %d\n", globval, i, j, k, l);
    f2();
}
static void f2(void) {
    longjmp(jmpbuffer, 1);
}
```



## Program 7.13: results

**\$ cc testjmp.c**

**\$ ./a.out**

**in f1():**

**globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99**

**after longjmp:**

**globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99**

**compile without  
any optimization**

**\$ cc -O testjmp.c**

**\$ ./a.out**

**in f1():**

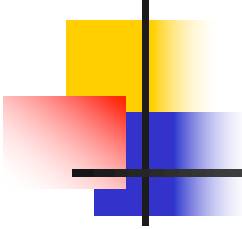
**globval = 95, autoval = 96, regival = 97, volaval = 98, statval = 99**

**after longjmp:**

**globval = 95, autoval = 2, regival = 3, volaval = 98, statval = 99**

**compile with  
full optimization**

# Program 7.14: Incorrect usage of automatic variables



```
#include <stdio.h>

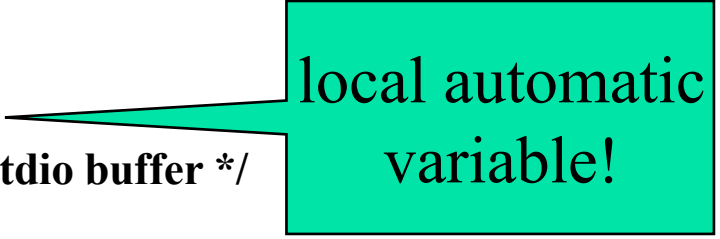
#define DATAFILE      "datafile"

FILE *
open_data(void)
{
    FILE *fp;
    char databuf[BUFSIZ];
        /* setvbuf makes this the stdio buffer */

    if ( (fp = fopen(DATAFILE, "r")) == NULL)
        return(NULL);

    if (setvbuf(fp, databuf, BUFSIZ, _IOLBF) != 0)
        return(NULL);

    return(fp);
        /* error */
}
```



local automatic variable!





## 11. getrlimit(), setrlimit()

---

- Every process has resource limits

```
#include <sys/time.h>
```

```
#include <sys/resource.h>
```

```
int getrlimit(    int resource,  
                  struct rlimit *rlptr    );
```

```
int setrlimit(    int resource,  
                  const struct rlimit *rlptr );
```

Return: 0 if OK, nonzero on error



# Resource Limits

---

■ struct rlimit {

■ rlim\_t rlim\_cur; /\* soft limit: curr limit \*/


■ rlim\_t rlim\_max; /\* hard limit: max \*/

};

■ **Soft limit:** can be changed by any process to  
 $\leq$  **hard limit**

■ **Hard limit:** can be changed by any process to  
 $\geq$  **soft limit** (irreversible!)

■ can be raised only by superuser process



# Resource Limits Example

---

- `RLIMIT_AS`
  - The maximum size in bytes of a process' s total available memory. This affects the `sbrk` function (Section 1.11) and the `mmap` function (Section 14.8).
- `RLIMIT_NPROC`
  - The maximum number of child processes per real user ID.



# Program 7.16: resource limits

```
#include <sys/types.h>
#include <sys/time.h>
#include <sys/resource.h>
#include "apue.h"

#define doit(name)          pr_limits(#name, name)

static void                pr_limits(char *, int);

int main(void) {
    doit(RLIMIT_CORE);
    doit(RLIMIT_CPU);
    doit(RLIMIT_DATA);
    doit(RLIMIT_FSIZE);
#ifdef RLIMIT_MEMLOCK
    doit(RLIMIT_MEMLOCK);
#endif
#ifdef RLIMIT_NOFILE      /* SVR4 name */
    doit(RLIMIT_NOFILE);
#endif
#ifdef RLIMIT_OFILE       /* 44BSD name */
    doit(RLIMIT_OFILE);
#endif
#ifdef RLIMIT_NPROC
    doit(RLIMIT_NPROC);
#endif
}
```

```

#ifdef RLIMIT_RSS
    doit(RLIMIT_RSS);
#endif
    doit(RLIMIT_STACK);
#ifdef RLIMIT_VMEM
    doit(RLIMIT_VMEM);
#endif
    exit(0);
}
static void pr_limits(char *name, int resource) {
    struct rlimit      limit;

    if (getrlimit(resource, &limit) < 0)
        err_sys("getrlimit error for %s", name);
    printf("%-14s  ", name);
    if (limit.rlim_cur == RLIM_INFINITY)
        printf("(infinite)  ");
    else
        printf("%10ld  ", limit.rlim_cur);
    if (limit.rlim_max == RLIM_INFINITY)
        printf("(infinite)\n");
    else
        printf("%10ld\n", limit.rlim_max);
}

```



## Program 7.16: 4.3 BSD results

---

■ \$ a.out		
■ RLIMIT_CORE	(infinite)	(infinite)
■ RLIMIT_CPU	(infinite)	(infinite)
■ RLIMIT_DATA	8388608	16777216
■ RLIMIT_FSIZE	(infinite)	(infinite)
■ RLIMIT_MEMLOCK	(infinite)	(infinite)
■ RLIMIT_OFILE	64	(infinite)
■ RLIMIT_NPROC	40	(infinite)
■ RLIMIT_RSS	27070464	27070464
■ RLIMIT_STACK	524288	16777216