

# Chapter 8.

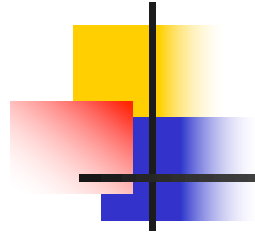
# Process Control



---

朱金辉

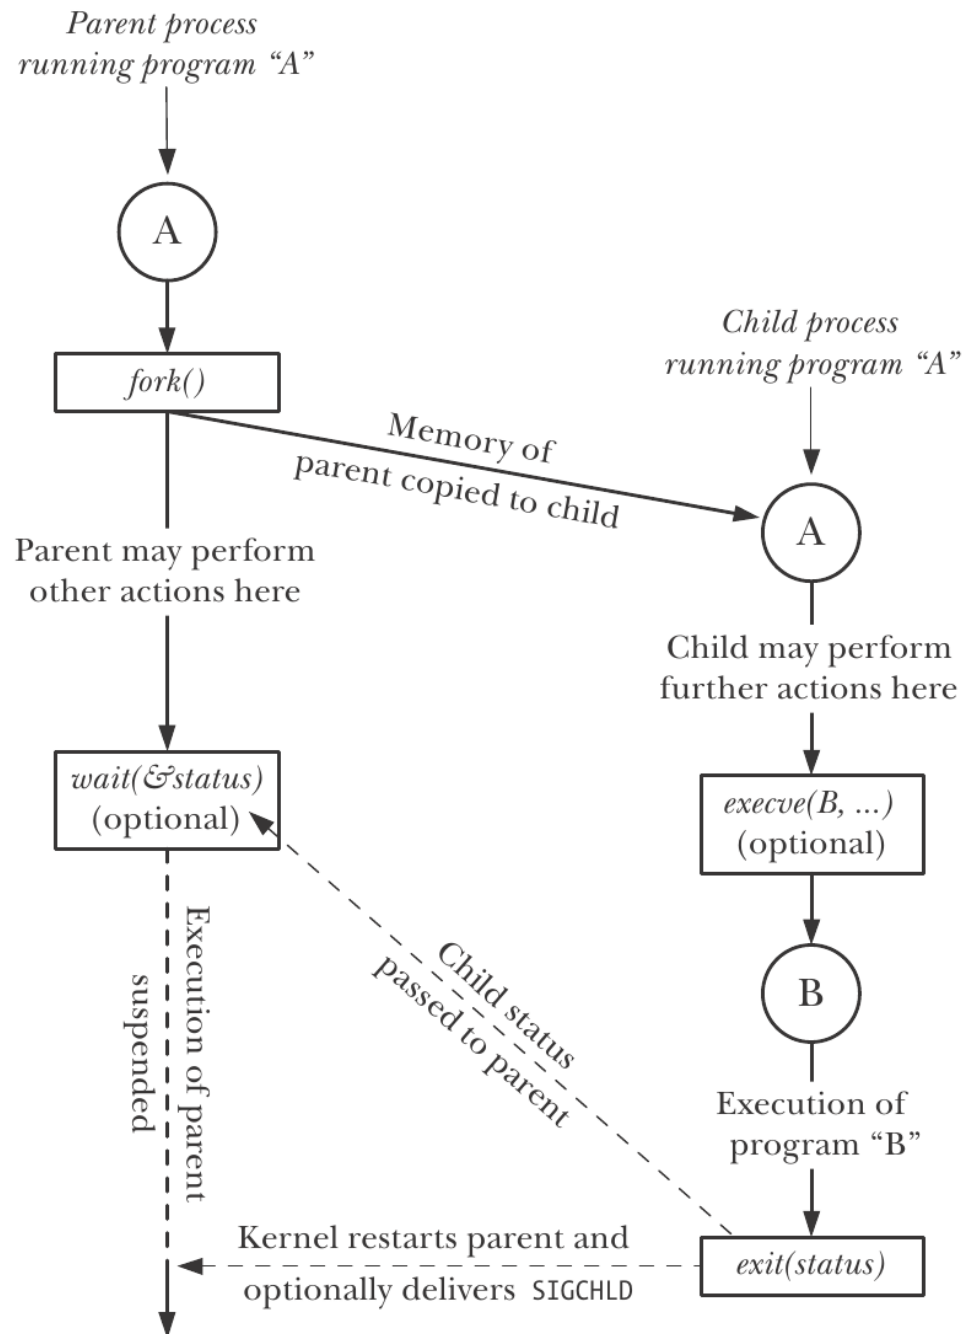
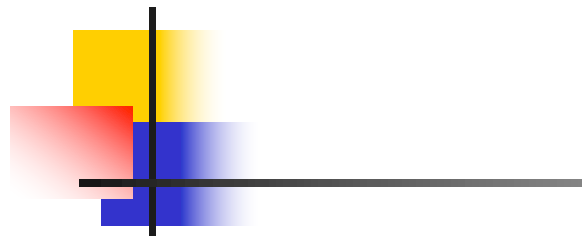
华南理工大学软件学院



# 1. Introduction

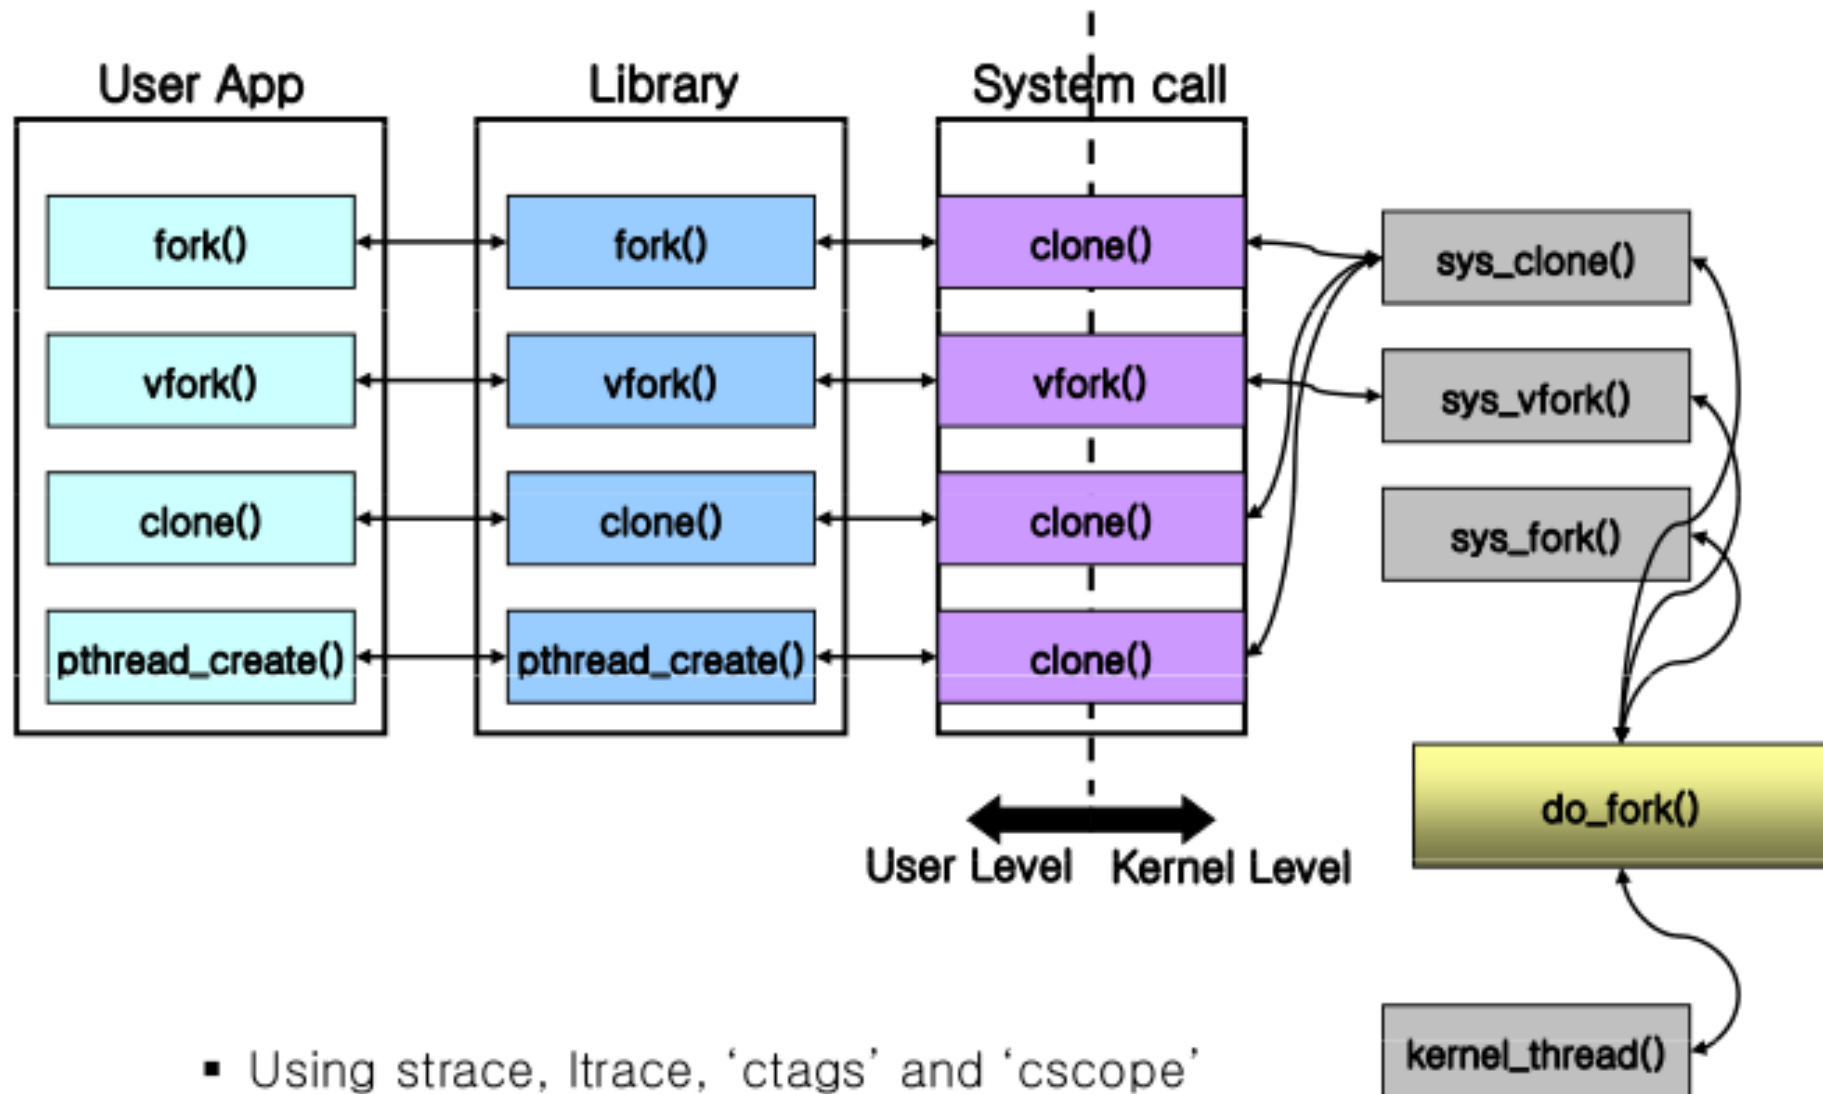
---

- Creation of new processes (fork)
- Process termination
- Executing programs (exec)
- IDs
- `system()`
- Process accounting



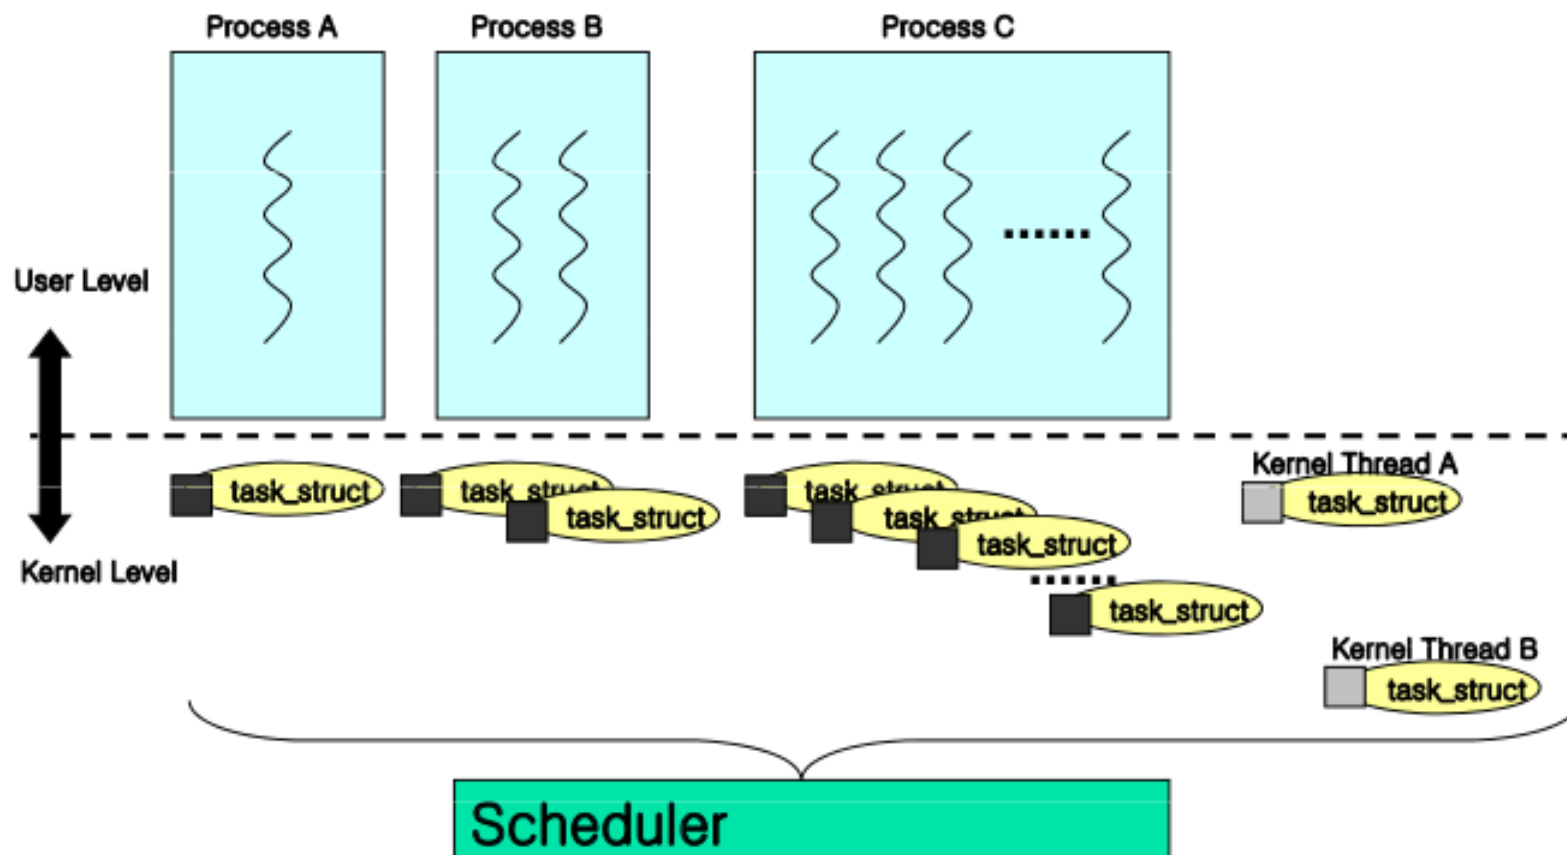
# Fork API

- Flow controls

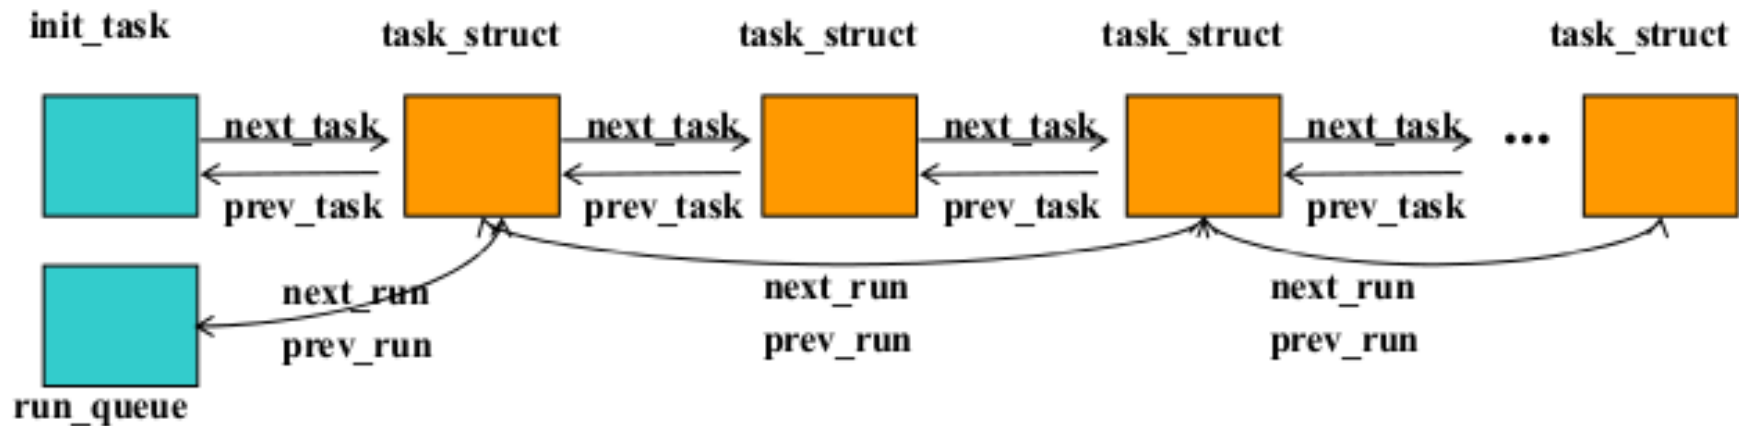


# Linux task model

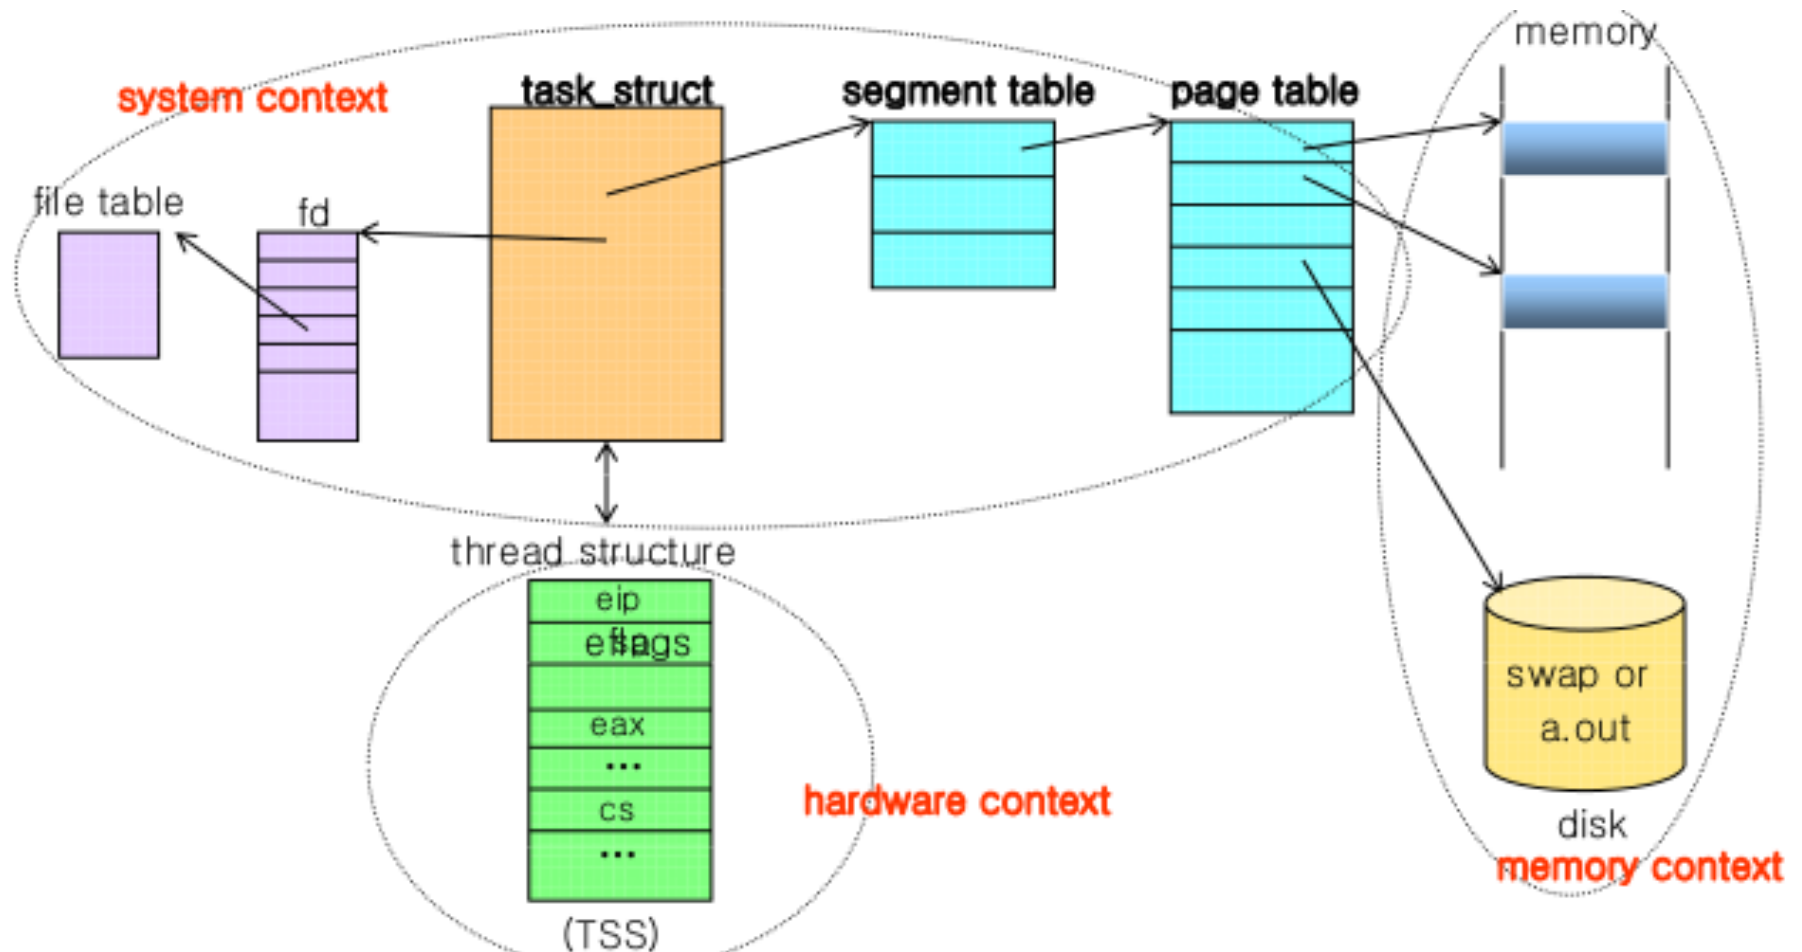
## ■ Internal structure



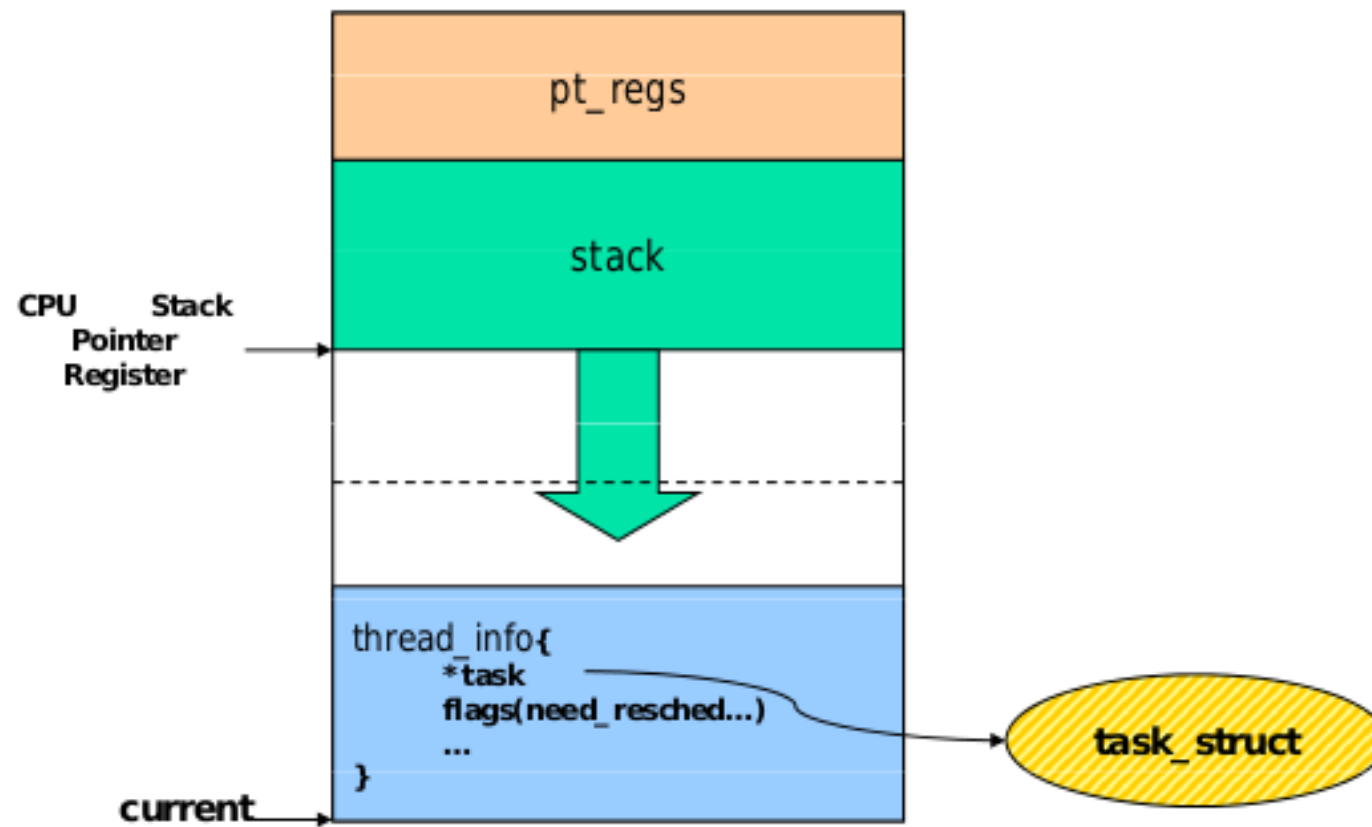
# Task queue



# Task context

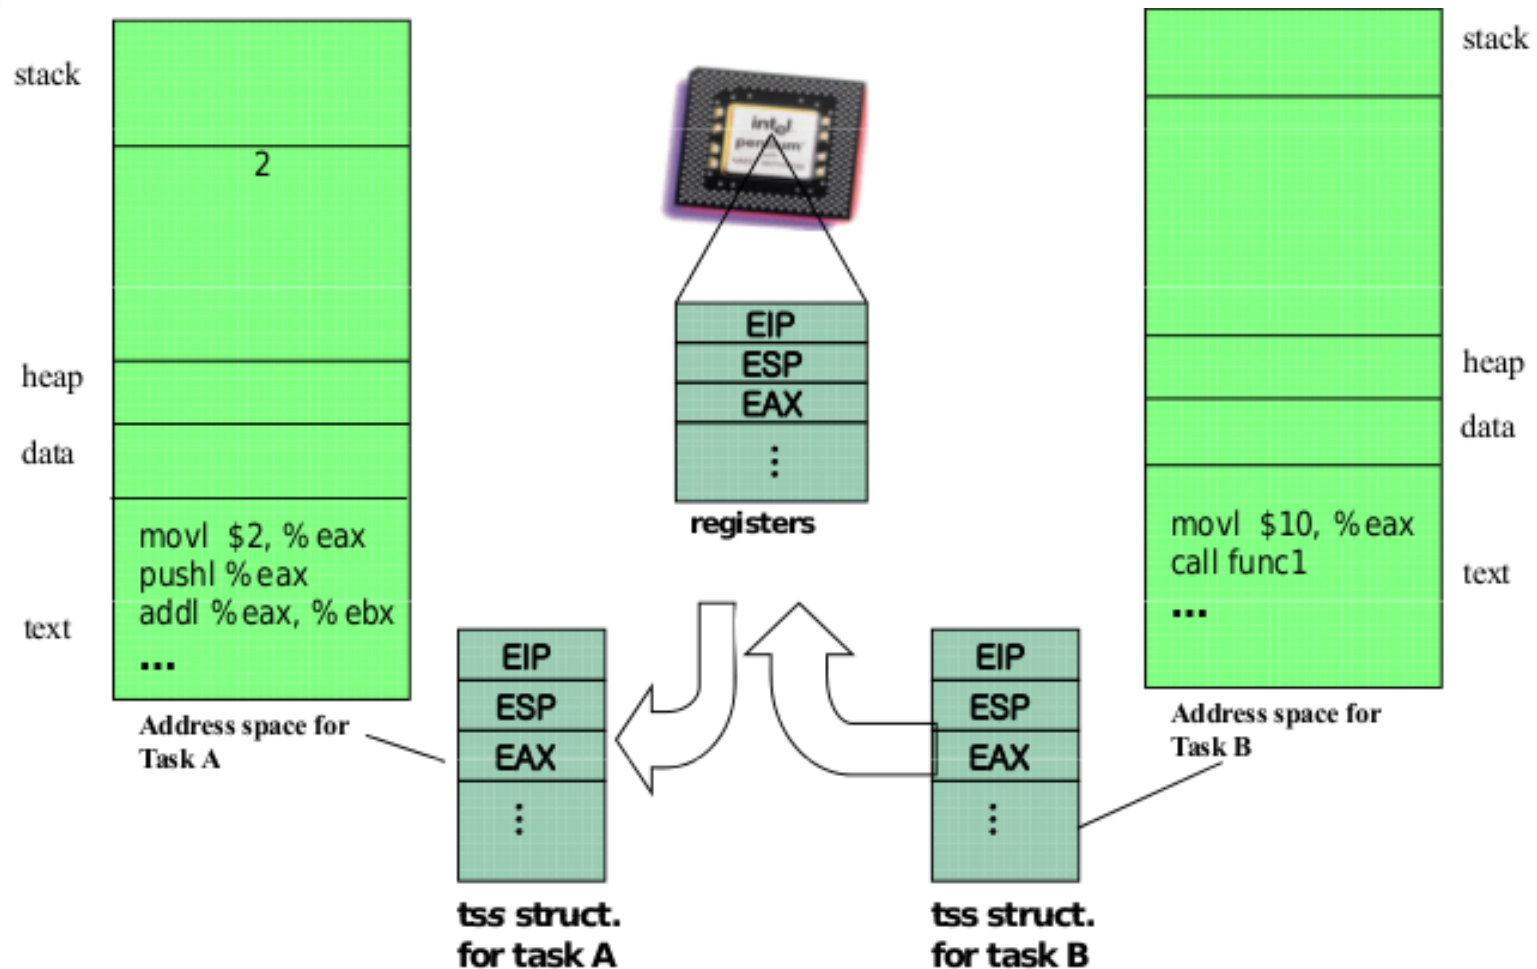


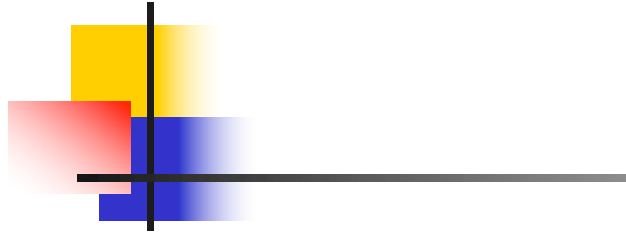
# Kernel stack



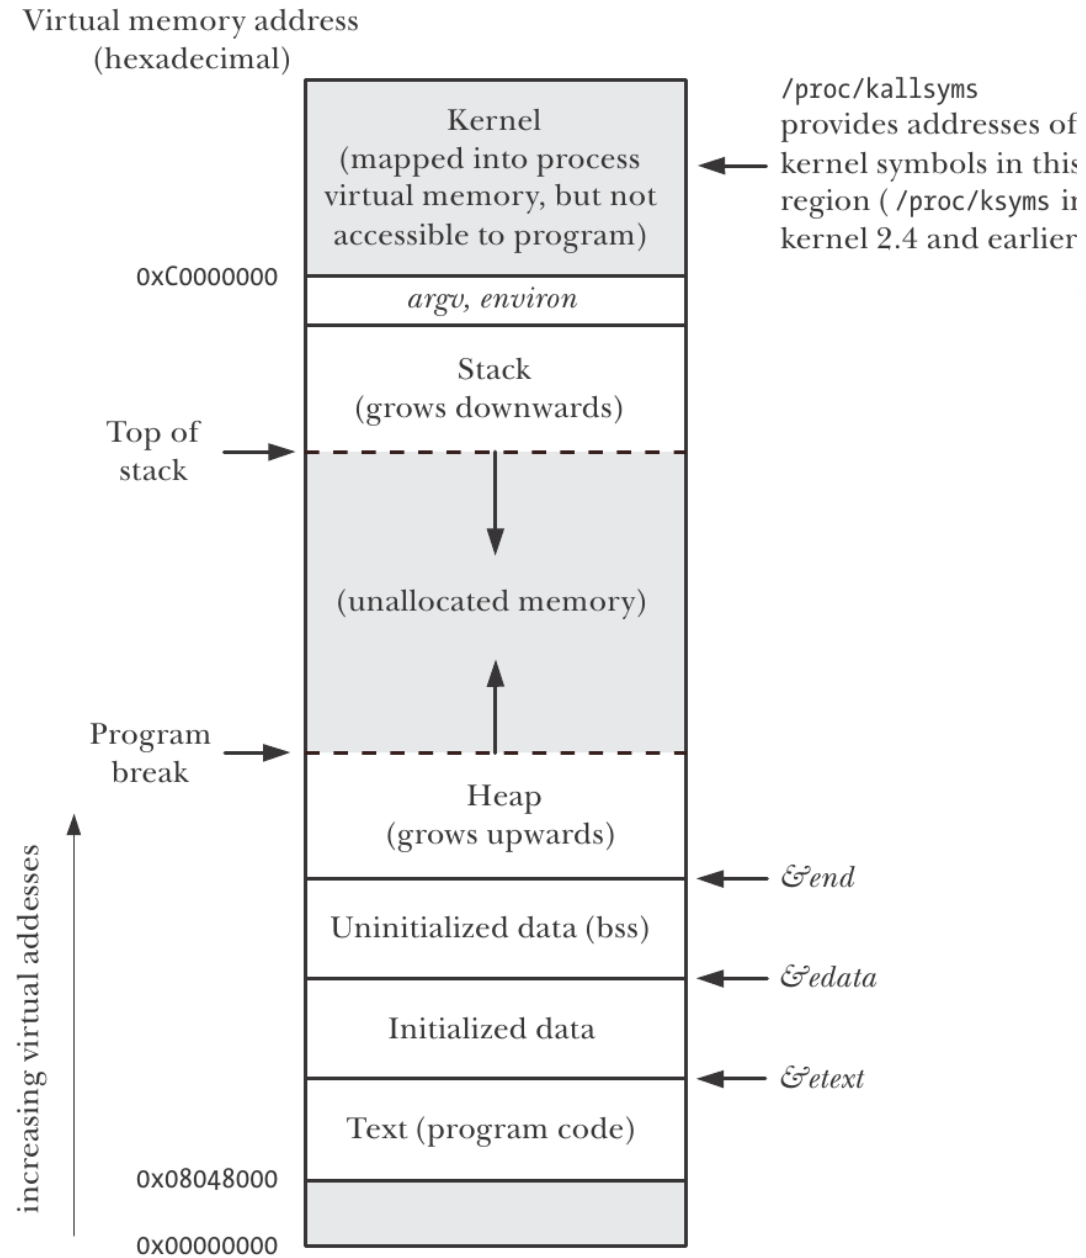


# Hardware context switch



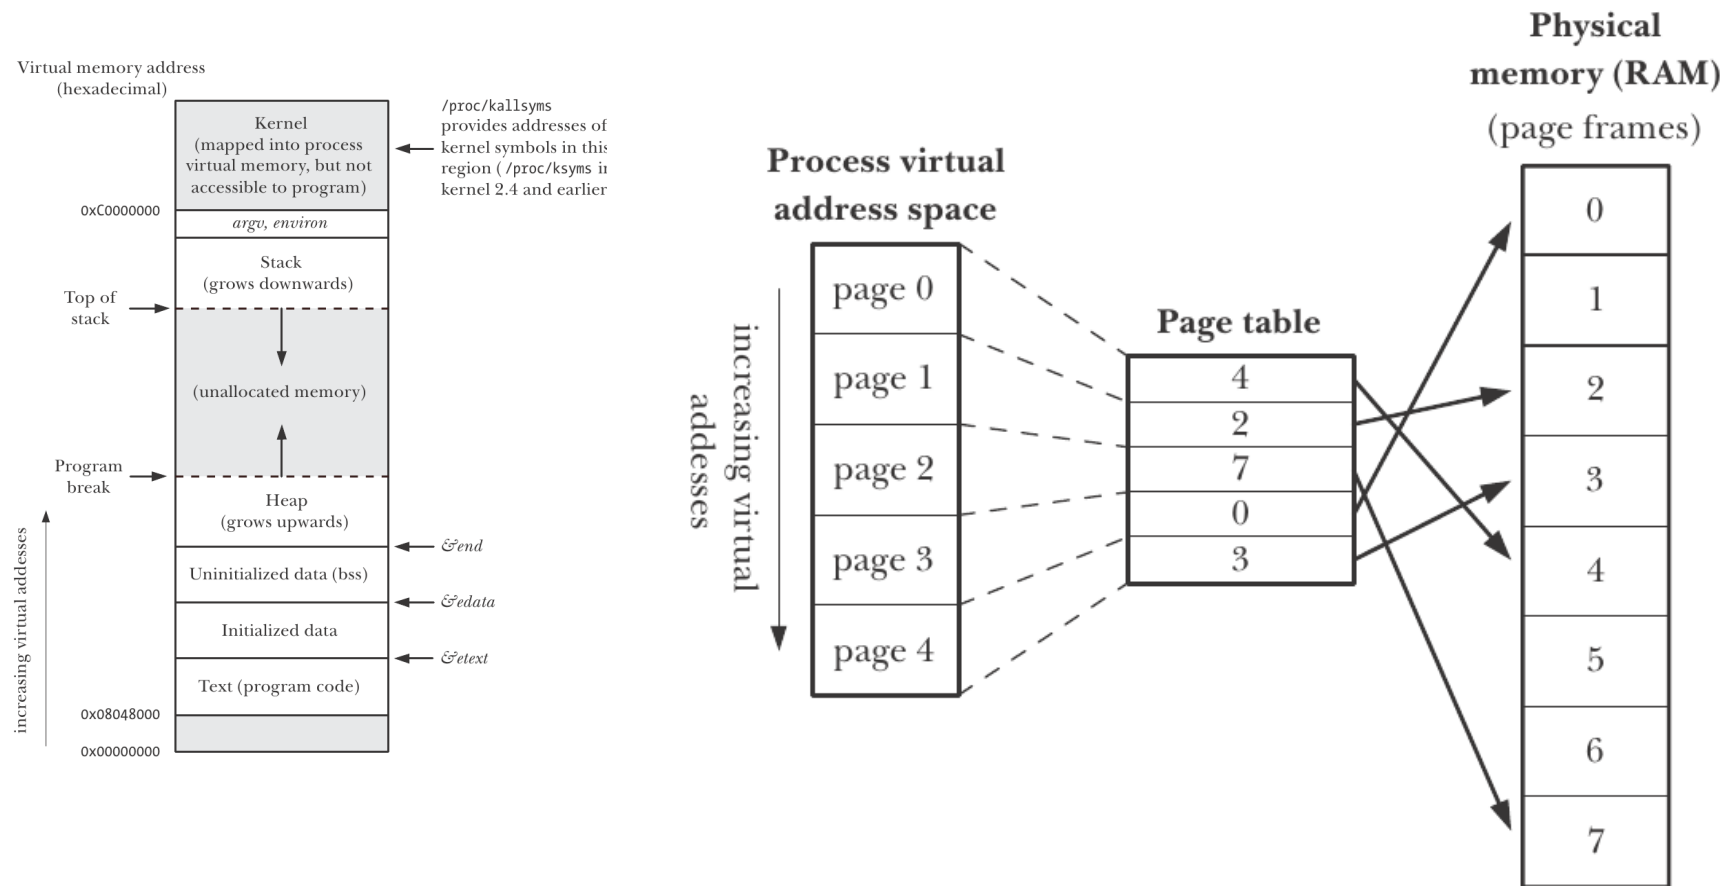


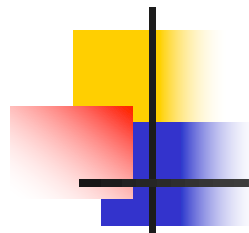
# Virtual Memory



Typical memory layout of a process on Linux/x86-32

# Virtual mem & Physical mem





# Example

Physical Memory	
00x	H E L L
01x	R L D !
02x	O W O
03x	H A V E
04x	F U N
05x	L O T
06x	S O F
07x	; - )

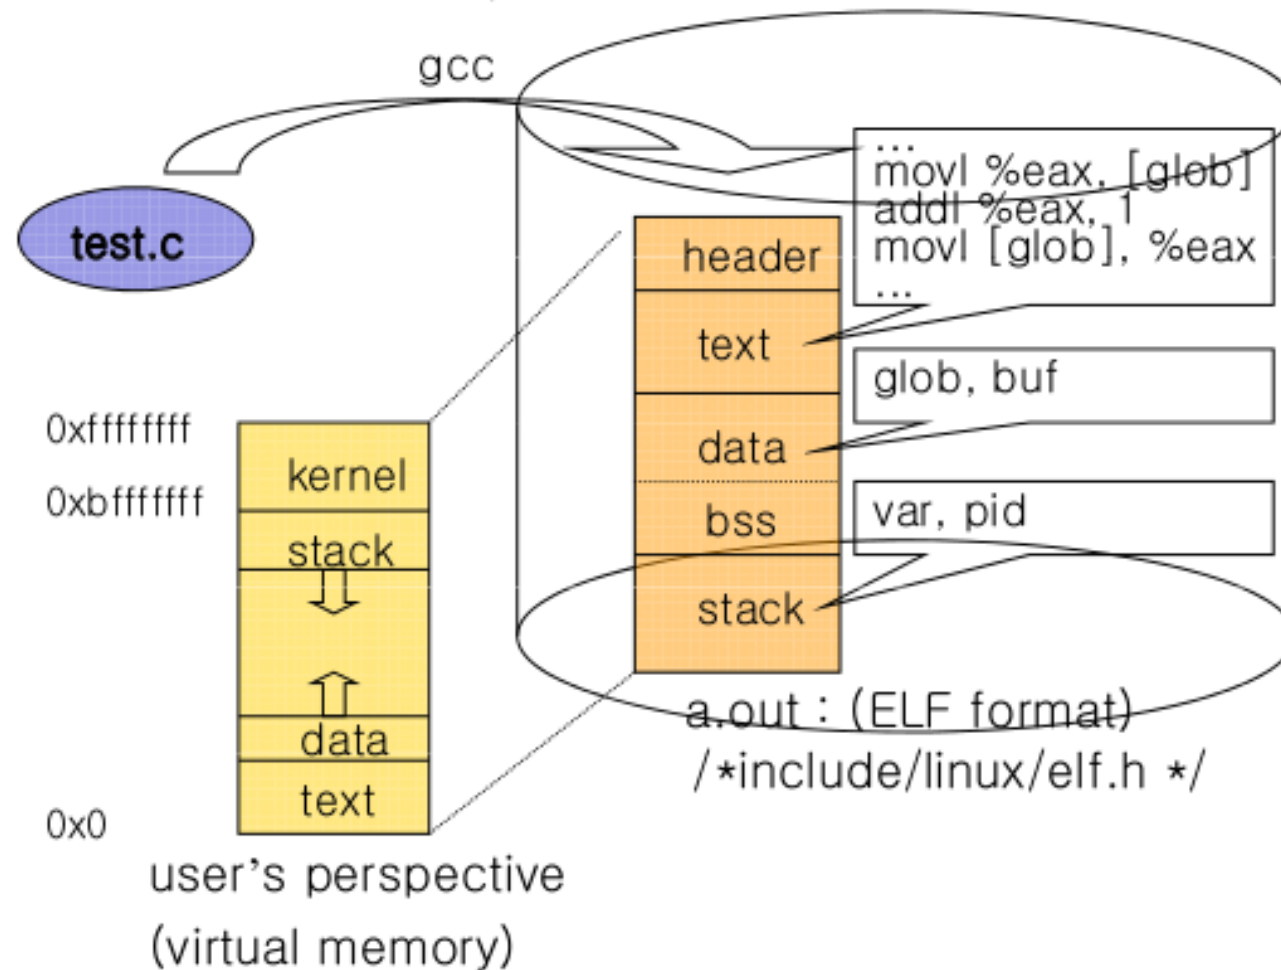
Process A			
Page Table		Virtual Memory	
00x	00	00x	H E L L
01x	02	01x	O W O
02x	01	02x	R L D !
03x	n.a.	03x	#####
04x	n.a.	04x	#####
05x	07	05x	; - )

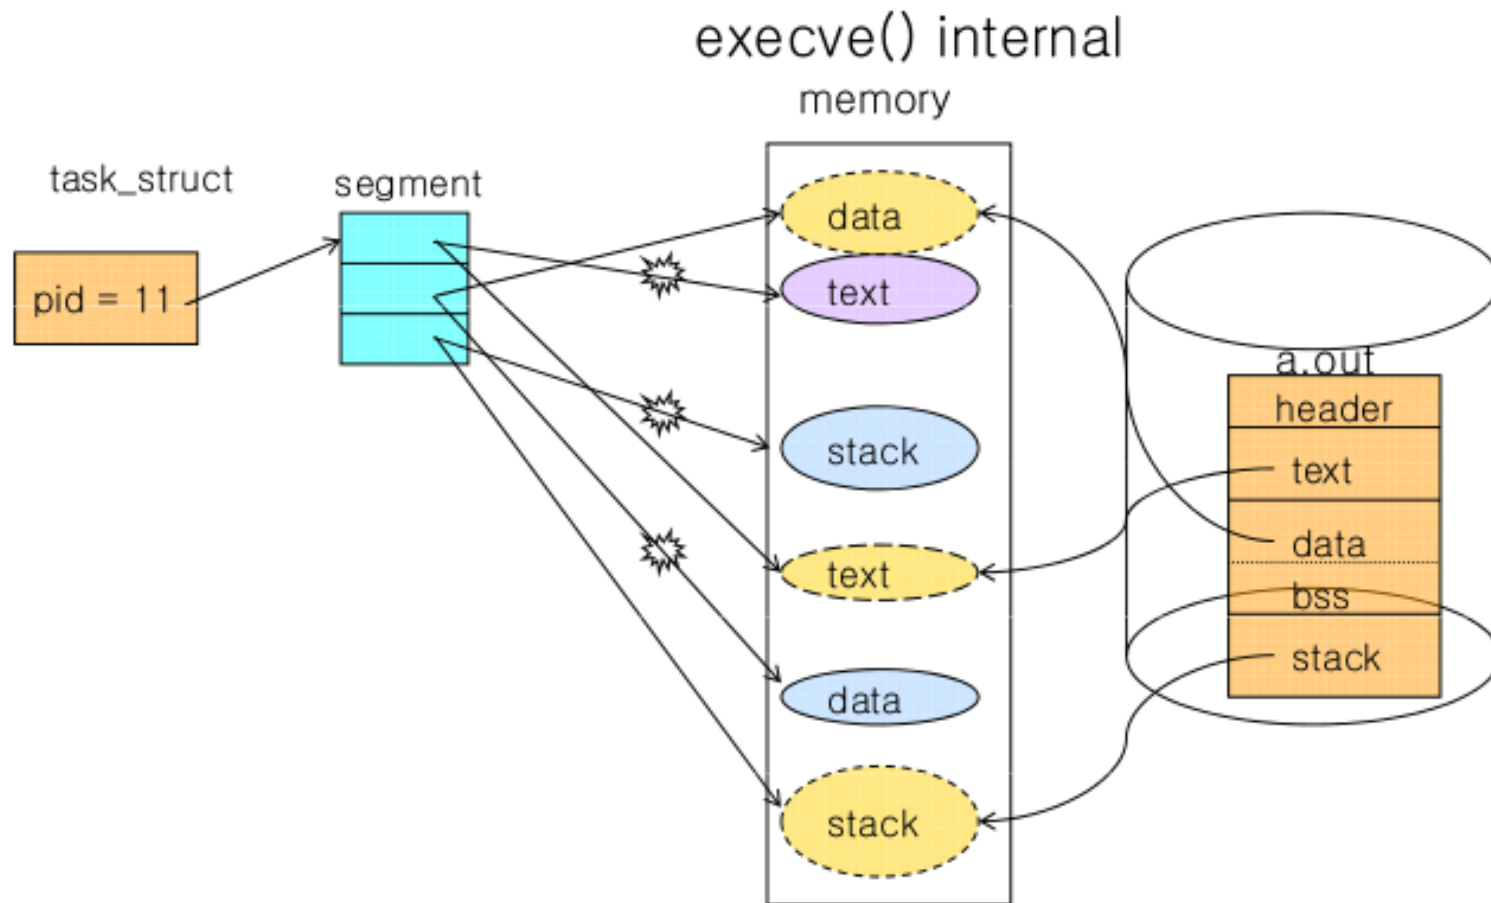
Process B			
Page Table		Virtual Memory	
00x	03	00x	H A V E
01x	05	01x	L O T
02x	06	02x	S O F
03x	04	03x	F U N
04x	n.a.	04x	#####
05x	07	05x	; - )

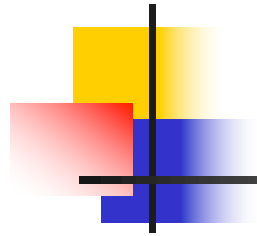
# execve

- fork internal : compile results



# execve





## 2. Process Identifiers

- Process ID = a nonnegative integer

PID	Process
0	<b>swapper</b> (scheduler)
1	<b>init</b> (/sbin/init)
2	<b>pagedaemon</b> (virtual memory paging)
3, 4, ...	other processes



# Identifier functions

---

- `#include <sys/types.h>`
- `#include <unistd.h>`
- `pid_t getpid(void);` return PID
- `pid_t getppid(void);` return **parent** PID
- `uid_t getuid(void);` return **real** UID
- `uid_t geteuid(void);` return **effective** UID
- `gid_t getgid(void);` return real GID
- `gid_t getegid(void);` return effective GID





## 3. fork Function

---

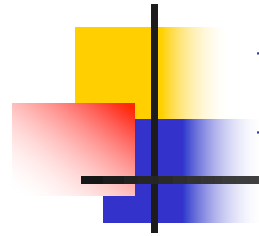
- `fork()` is the ONLY way to create a process in Unix kernel by user

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t fork(void);
```

- Returns: 0 in child, child PID in parent, -1 on error



# Parent / Child Processes

---

- Parent and child continue executing instructions following the `fork()` call
- Child gets a **copy** of parent's data space, heap, and stack
- Often, read-only text segment is shared
- Often, `fork()` is followed by `exec()`
- **Waste** of space and time for setting up child's program space!!!

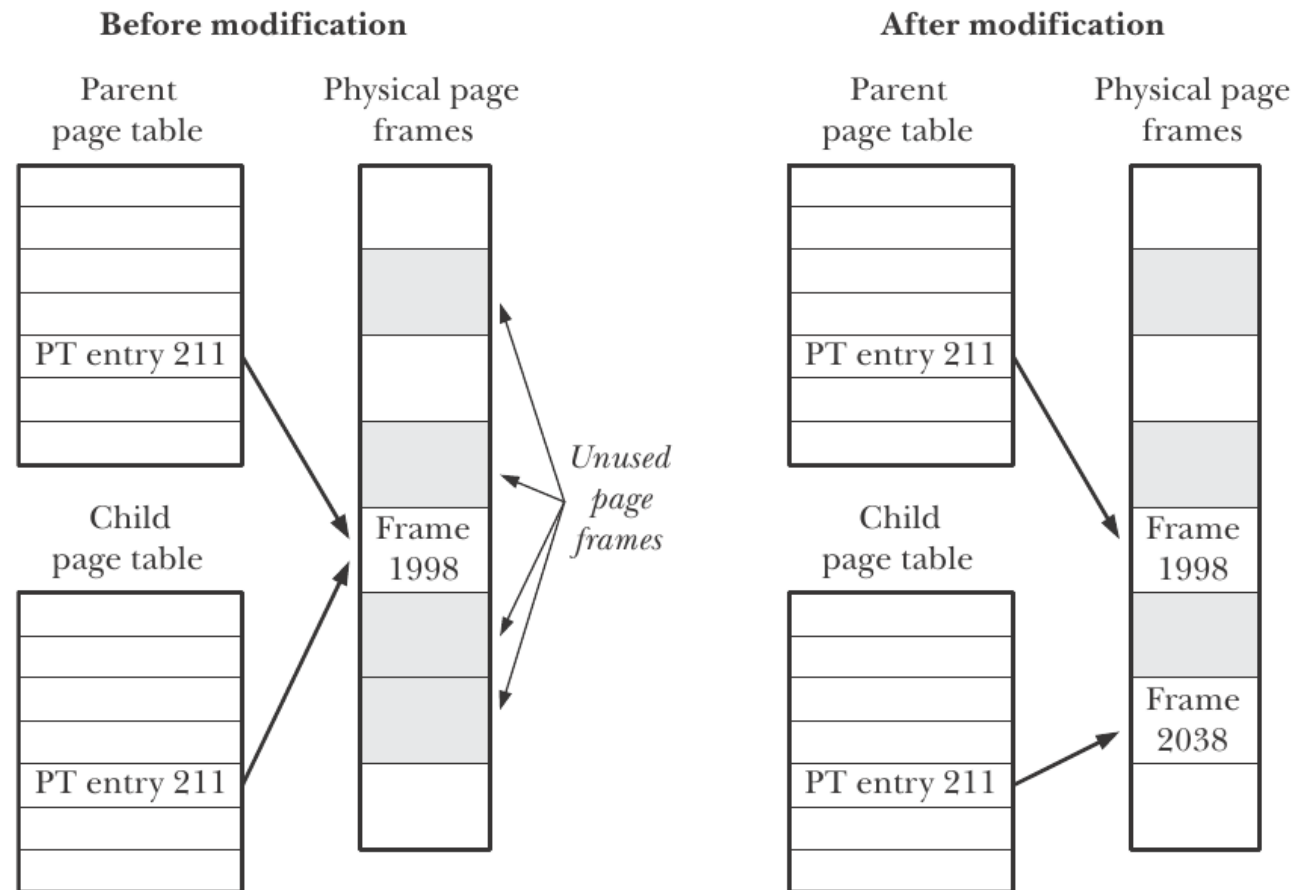


# Copy-On-Write (COW)

---

- Memory regions are **read-only** and **shared** by parent and child
- If either process wants to write, kernel makes a **copy** of that memory only for that process.
- Saves space and time!

# COW example



Page tables before and after modification of a shared copy-on-write page

# Program 8.1: fork()

```
#include <sys/types.h>
#include "apue.h"
int glob = 6;      /* external variable in initialized data */
char    buf[] = "a write to stdout\n";
int main(void) {
    int    var;      /* automatic variable on the stack */
    pid_t pid;
    var = 88;
    if (write(STDOUT_FILENO, buf, sizeof(buf)-1) != sizeof(buf)-1)
        err_sys("write error");
    printf("before fork\n"); /* we don't flush stdout */
    if ( (pid = fork()) < 0)    err_sys("fork error");
    else if (pid == 0) {      /* child */
        glob++;               /* modify variables */
        var++;
    } else
        sleep(2);             /* parent */
    printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
    exit(0);
}
```



## Program 8.1: results

---

**\$ a.out**

a write to stdout

**before fork**

pid = 430, glob = 7, var = 89

pid = 429, glob = 6, var = 88

**\$ a.out > temp.out**

**\$ cat temp.out**

a write to stdout

**before fork**

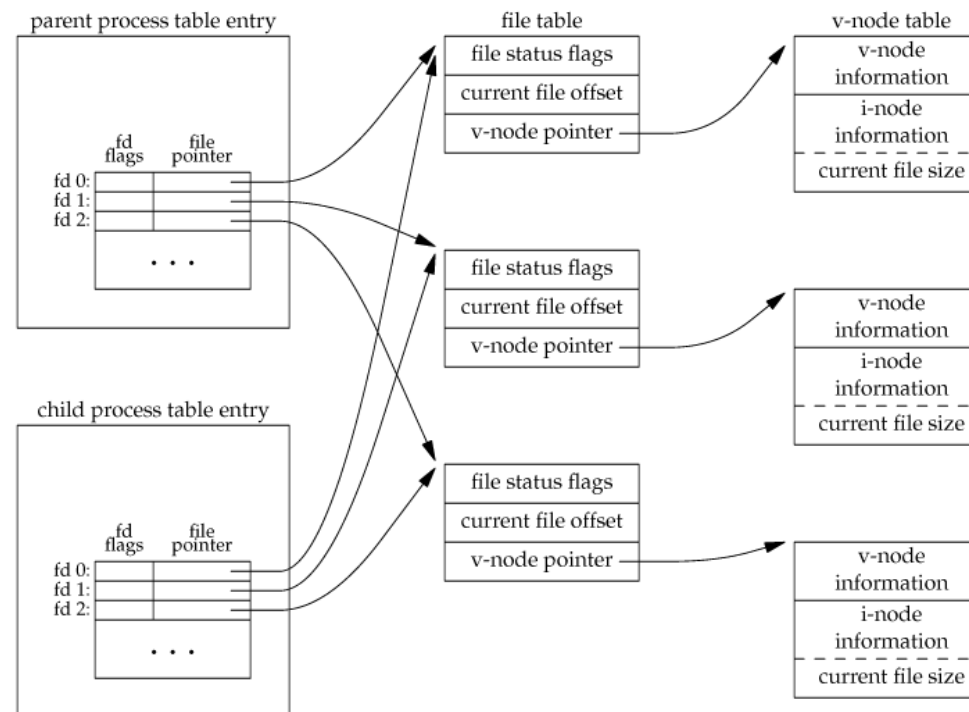
pid = 432, glob = 7, var = 89

**before fork**

pid = 431, glob = 6, var = 88

# File Sharing

- Parent and child share the **same file descriptors**
- Parent and child share the **same file offset**, otherwise overwrite
- **Intermixed output** from parent and child





## 4. vfork Function

---

- Creates a new process only to ‘exec’ a new program
- **No copy of parent’s address space** for child (not needed!)
- Before exec, child runs in “address space of parent”
- Efficient in paged virtual memory
- **Child runs first**
- Parent waits until child ‘exec’ or ‘exit’





## Program 8.3: vfork()

```
#include <sys/types.h>
#include "apue.h"

int      glob = 6;          /* external variable in initialized data */

int main(void) {
    int      var;           /* automatic variable on the stack */
    pid_t pid;

    var = 88;
    printf("before vfork\n");    /* we don't flush stdio */

    if ( (pid = vfork()) < 0)
        err_sys("vfork error");
    else if (pid == 0) {        /* child */
        glob++;                /* modify parent's variables */
        var++;
        _exit(0);              /* child terminates */
    }
    /* parent */
    printf("pid = %d, glob = %d, var = %d\n", getpid(), glob, var);
    exit(0);
}
```



## Program 8.3: results

---

- \$ a.out

before vfork

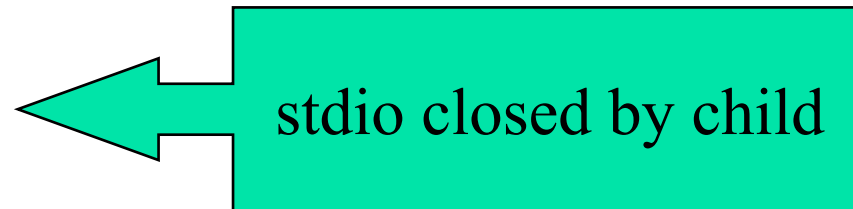
pid = 607, glob = 7, var = 89

- increments by child appear in parent address space

- Instead of `_exit()` → `exit()`, results in:

\$ a.out

before vfork





## 5. Child Termination

---

Termination status:

- normal: exit status
- abnormal: kernel indicates reason
- What if child terminates before parent?
  - Child returns termination status to parent
- What if parent terminates before child?
  - Parent PID (of orphaned child) = 1 (init)



# SIGCHLD

---

- Child terminates →  
Kernel sends **SIGCHLD** signal to parent
- Default action for SIGCHLD signal: ignore it
- Signal handlers can be defined by users  
(Chapter 10)



## 6. wait(), waitpid()

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t wait(int *statloc);
```

block wait for  
any one child to  
terminate

```
pid_t waitpid(    pid_t pid, int *statloc,  
                 int options);
```

place for storing  
termination status  
NULL → no need!

- Return: PID if OK, 0, -1 on error



## wait3 and wait4

---

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

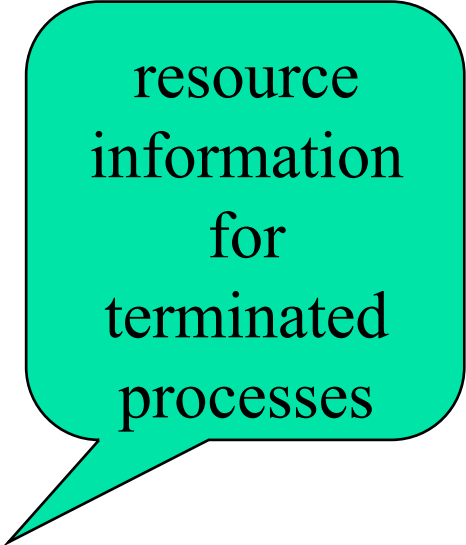
```
#include <sys/time.h>
```

```
#include <sys/resource.h>
```

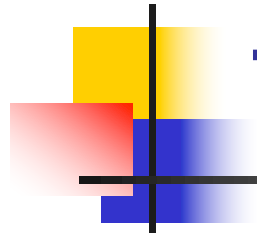
```
pid_t wait3(int *statloc, int options,  
            struct rusage *rusage);
```

```
pid_t wait4(pid_t pid, int *statloc, int options,  
            struct rusage *rusage);
```

- Return: PID if OK, 0 or -1 on error



resource  
information  
for  
terminated  
processes



# Termination Status Macros

Macro	Description
<code>WIFEXITED (status)</code>	<p>True if status was returned for a child that terminated normally. In this case, we can execute</p> <p><code>WEXITSTATUS (status)</code></p> <p>to fetch the low-order 8 bits of the argument that the child passed to <code>exit</code>, <code>_exit</code>, or <code>_Exit</code>.</p>
<code>WIFSIGNALED (status)</code>	<p>True if status was returned for a child that terminated abnormally, by receipt of a signal that it didn't catch. In this case, we can execute</p> <p><code>WTERMSIG (status)</code></p> <p>to fetch the signal number that caused the termination.</p> <p>Additionally, some implementations (but not the Single UNIX Specification) define the macro</p> <p><code>WCOREDUMP (status)</code></p> <p>that returns true if a core file of the terminated process was generated.</p>
<code>WIFSTOPPED (status)</code>	<p>True if status was returned for a child that is currently stopped. In this case, we can execute</p> <p><code>WSTOPSIG (status)</code></p> <p>to fetch the signal number that caused the child to stop.</p>
<code>WIFCONTINUED (status)</code>	<p>True if status was returned for a child that has been continued after a job control stop (XSI option; <code>waitpid</code> only).</p>



# Program 8.5: print exit status

```
#include <sys/types.h>
#include <sys/wait.h>
#include "apue.h"

void pr_exit(int status) {
    if (WIFEXITED(status))
        printf("normal termination, exit status = %d\n",
               WEXITSTATUS(status));
    else if (WIFSIGNALED(status))
        printf("abnormal termination, signal number = %d%s\n",
               WTERMSIG(status),
#ifdef WCOREDUMP
               WCOREDUMP(status) ? " (core file generated)" : "");
#else
               "");
#endif
    else if (WIFSTOPPED(status))
        printf("child stopped, signal number = %d\n",
               WSTOPSIG(status));
}
```





## Program 8.6: demo exit status

```
#include    <sys/types.h>
#include    <sys/wait.h>
#include    "apue.h"

int main(void) {
    pid_t    pid;
    int      status;

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0)          /* child */
        exit(7);
    if (wait(&status) != pid)  /* wait for child */
        err_sys("wait error");
    pr_exit(status);           /* and print its status */
}
```



## Program 8.6 (II Part)

---

```
if ( (pid = fork()) < 0)
    err_sys("fork error");
else if (pid == 0)      /* child */
    abort();            /* generates SIGABRT */

if (wait(&status) != pid) /* wait for child */
    err_sys("wait error");
pr_exit(status);        /* and print its status */
```



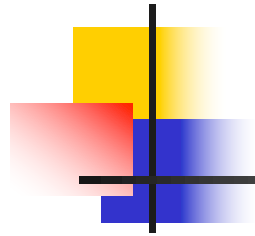
## Program 8.6 (III part)

---

```
if ( (pid = fork()) < 0)
    err_sys("fork error");
else if (pid == 0)                /* child */
    status /= 0; /* divide by 0 generates SIGFPE */

if (wait(&status) != pid)         /* wait for child */
    err_sys("wait error");
pr_exit(status);                 /* and print its status */

exit(0);
}
```



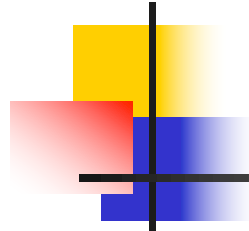
## Program 8.6: results

---

- **\$ a.out**
- normal termination, exit status = 7
- abnormal termination, signal number = 6  
(core file generated)
- abnormal termination, signal number = 8  
(core file generated)

SIGABRT

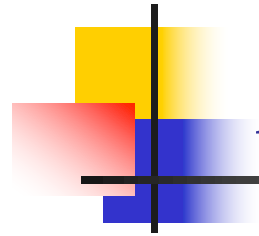
SIGFPE



# Zombie process

---

- Suppose **child terminates first**  
**&& parent don't wait child**
- Zombie: minimal info of dead child process  
(pid, termination status, CPU time)



# Avoiding zombie processes

---

- A process forks a child
- It does not wait for the child to complete
- It does not want child to become zombie
- How to do this?
- Answer: fork twice! (Program 8.8)

## Program 8.8: Avoid Zombie

```
int main(void) {
    pid_t    pid;
    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) {                /* first child */
        if ( (pid = fork()) < 0)
            err_sys("fork error");
        else if (pid > 0) /* parent from second fork */
            exit(0);     /* == first child */
        /* second child; parent becomes init */
        sleep(2);
        printf("second child, parent pid = %d\n", getppid());
        exit(0);
    }
    if (waitpid(pid, NULL, 0) != pid) /* wait for first child */
        err_sys("waitpid error");

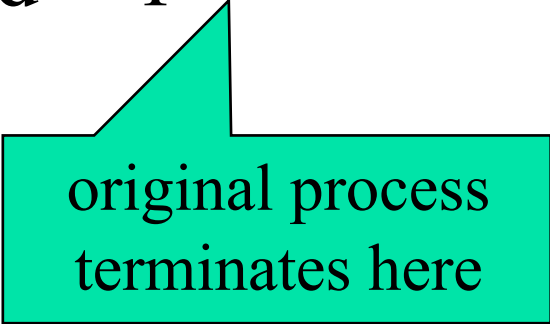
    /* We're the parent (the original process) */
    ...
    exit(0);
}
```



## Program 8.8: results

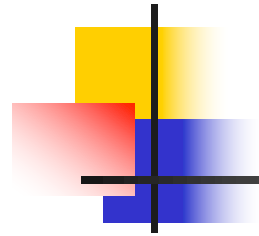
---

- **\$ a.out**
- second child, parent pid = 1



original process  
terminates here





## 9. Race Conditions

---

- Multiple processes share some data
- Outcome depends on the order of their execution (i.e. RACE)
- After `fork()`, we **cannot predict** if the parent or the child runs first!
- The order of execution depends on:
  - system load
  - Kernel's scheduling algorithm



# Program 8.12: Race Condition

---

```
#include <sys/types.h>
#include "apue.h"

static void charatime(char *);

int main(void) {
    pid_t pid;

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) {
        charatime("output from child\n");
    } else {
        charatime("output from parent\n");
    }
    exit(0);
}
```



## Program 8.12 (continued)

---

```
static void  
charatime(char *str){  
    char        *ptr;  
    int         c;  
  
    setbuf(stdout, NULL);    /* set unbuffered */  
    for (ptr = str; c = *ptr++; )  
        putc(c, stdout);  
}
```



## Program 8.12: results

---

- **\$ a.out**
  - output from child
  - output from parent
- 
- **\$ a.out**
  - ooutput from parent
  - utput from child



# Race Conditions

---

- Race condition problems are hard to detect because they work “most of the time” !
  1. For parent to wait for child
    - call wait, waitpid, wait3, wait4
  2. For child to wait for parent
    - **while (getppid() != 1) sleep(1);**

polling!  
wastes CPU  
time!

use signals or  
other IPC methods



# Race Conditions

---

- After fork
  - parent and child both need to do something on its own
  - For example, **parent**: write a record in a log file and **child**: creates a log file
- Parent and child need to:
  - **TELL** each other when its initial set of operations are done, and
  - **WAIT** for each other to complete



## Program 8.13: No race condition

---

```
#include <sys/types.h>
#include "apue.h"
static void charatime(char *);
int main(void) {
    pid_t pid;
    TELL_WAIT();
    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) {
        WAIT_PARENT();          /* parent goes first */
        charatime("output from child\n");
    } else {
        charatime("output from parent\n");
        TELL_CHILD(pid);
    }
    exit(0);
}
```



## 10. exec Functions

---

```
#include <unistd.h>
int execl(const char *pathname,
          const char *arg0, ... /* (char *)0 */);
int execv(const char *pathname,
          char *const argv[]);
int execl_e(const char *pathname, const char *arg0, ... /*
          (char *)0, char *const envp[] */);
int execve(const char *pathname,
           char *const argv[], char *const envp[]);
int execlp(const char *filename,
           const char *arg0, ... /* (char *)0 */);
int execvp(const char *filename,
           char *const argv[]);

Return -1 on error, no return on success.
```



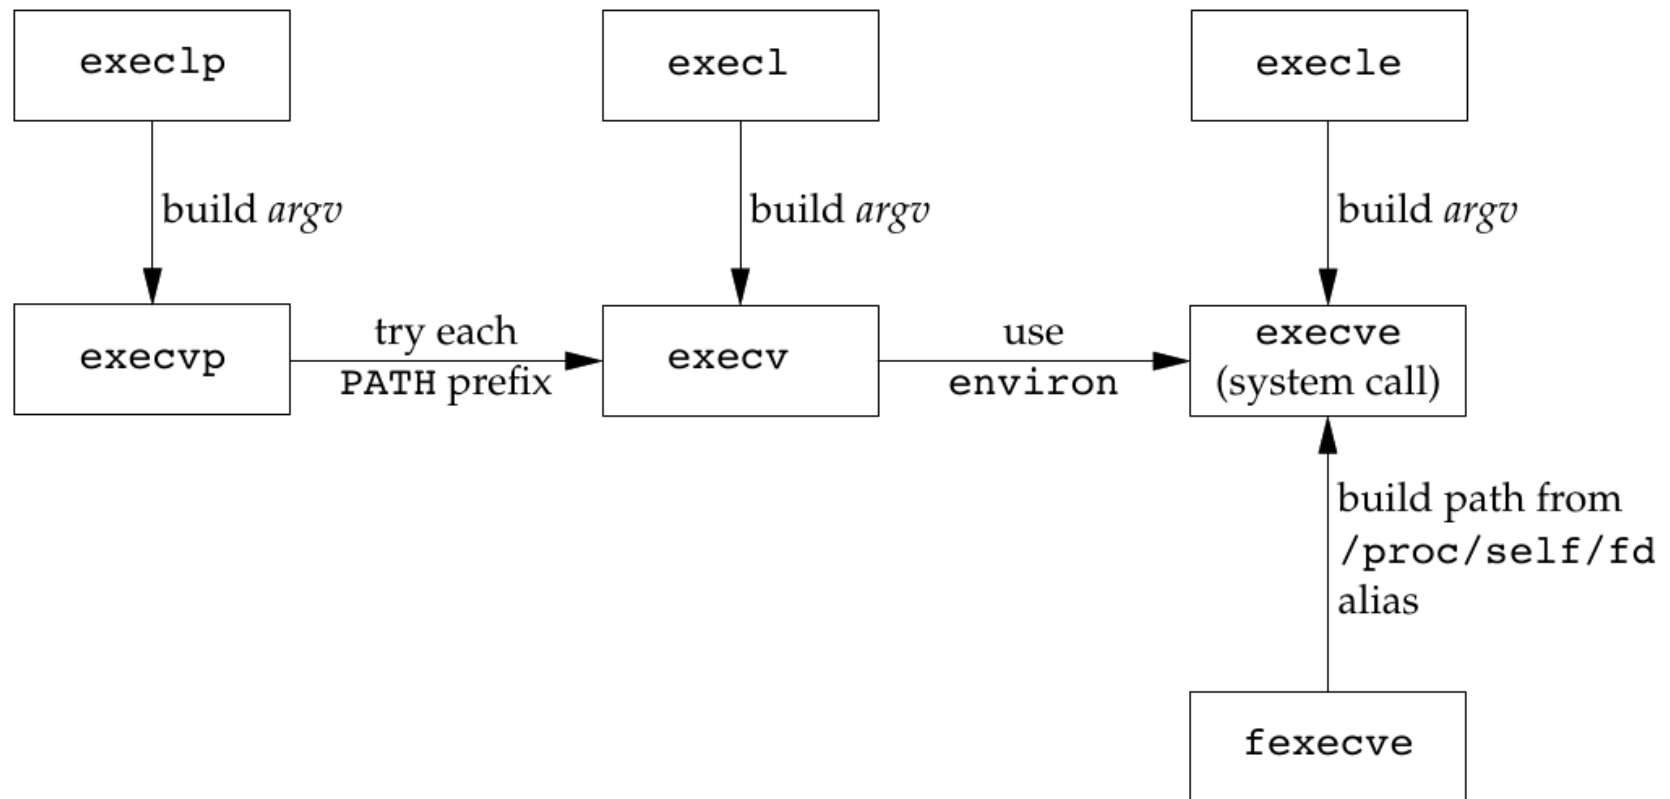


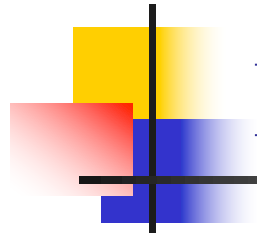
# Differences among exec functions

---

- **filename** (execl**p**, execvp**p**: uses PATH) v/s pathname (others: does not use PATH)
- list (**l**) v/s vector (**v**)
  - list of arguments (execl, execl, execlp)
  - array of pointers to arguments (execv, execve, execvp)
- pointer to an array of pointers to **e**nvironment strings (execl, execve) v/s environ (others)

# Relationship of 6 exec functions





# Differences among exec functions

Function	<i>pathname</i>	<i>filename</i>	<i>fd</i>	Arg list	<i>argv[ ]</i>	<i>environ</i>	<i>envp[ ]</i>
execl	•			•		•	
execlp		•		•		•	
execle	•			•			•
execv	•				•	•	
execvp		•			•	•	
execve	•				•		•
fexecve			•		•		•
(letter in name)		p	f	l	v		e

## Program 8.16: exec functions

```
#include <sys/types.h>
#include <sys/wait.h>
#include "apue.h"

char      *env_init[] = { "USER=unknown", "PATH=/tmp", NULL };

int main(void) {
    pid_t pid;

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) { /* specify pathname, specify environment */
        if (execl("/home/stevens/bin/echoall",
                  "echoall", "myarg1", "MY ARG2", (char *) 0,
                  env_init) < 0)
            err_sys("execl error");
    }
    if (waitpid(pid, NULL, 0) < 0)
        err_sys("wait error");

    if ( (pid = fork()) < 0)
        err_sys("fork error");
    else if (pid == 0) { /* specify filename, inherit environment */
        if (execlp("echoall", "echoall", "only 1 arg", (char *) 0) < 0)
            err_sys("execlp error"); }
    exit(0); }
```



# Inheritance by child from parent after exec

---

- PID, Parent PID
- Real UID, Real GID
- Supplementary GIDs
- Process GID
- Session ID
- Controlling Terminal
- Time Left Until Alarm Clock
- Current Working Directory
- Root Directory
- File Mode Creation Mask
- File Locks
- Process Signal Mask
- Pending Signals
- Resource Limits
- Nice value
- tms\_utime, tms\_stime, tms\_cutime, tms\_ustime values



# 11. Changing UIDs and GIDs

---

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
int setuid(uid_t uid);
```

```
int setgid(gid_t gid);
```

- Return: 0 if OK, -1 on error

- 1) If **Superuser**

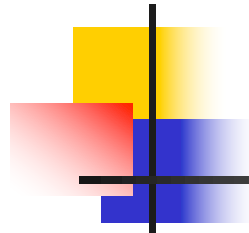
real, effective, saved set-UID := uid

- 2) If **real Or saved set-UID == uid**

effective := uid

else

errno := EPERM; return error;



# seteuid(), setegid()

---

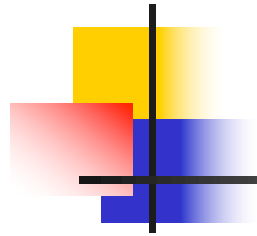
```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
int seteuid(uid_t uid);
```

```
int setegid(gid_t gid);
```

- Return: 0 if OK, -1 on error
- Only **effective** UID or GID is changed



## “saved set-UID” Example: at

- Example: **at**, executes commands at a specified time.

```
→ ~ ps -ef | grep atd
daemon      695      1  0 14:29 ?                00:00:00 /usr/sbin/atd -f
```

```
→ ~ ls -l /usr/bin/at /usr/sbin/atd
-rwsr-sr-x 1 daemon daemon 51464 1月 15 2016 /usr/bin/at
-rwxr-xr-x 1 root   root   26632 1月 15 2016 /usr/sbin/atd
```

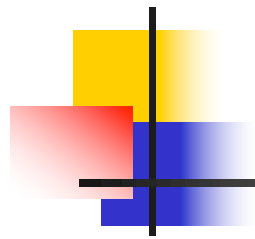


```
→ ~ at 5:00
warning: commands will be executed using /bin/sh
at> date >tmp.txt
at> <EOT>
job 22 at Tue Oct 10 05:00:00 2017
```

```
→ ~ sudo ls -l /var/spool/cron
total 12
drwxrwx--T 2 daemon daemon 4096 10月 9 16:40 atjobs
```

```
→ ~ sudo ls -l /var/spool/cron/atjobs
total 16
-rwx----- 1 zhu daemon 5979 10月 9 16:34 a00015017f65d8
-rwx----- 1 zhu daemon 5978 10月 9 16:41 a00016017f65ec
```

```
→ ~ sudo cat /var/spool/cron/atjobs/a00016017f65ec
#!/bin/sh
# atrun uid=1000 gid=1000
# mail zhu 0
umask 2
XDG_SEAT=seat0; export XDG_SEAT
ROS_MASTER_URI=http://localhost:11311; export ROS_MASTER_URI
ROSLISP_PACKAGE_DIRECTORIES=; export ROSLISP_PACKAGE_DIRECTORIES
cd /home/zhu || {
    echo 'Execution directory inaccessible' >&2
    exit 1
}
date >tmp.txt
```



## “saved set-UID” Example: At

- Example: **at**, executes commands at a specified time.
1. assuming owned by daemon, set-UID bit is SET, after **exec**:
    - Real UID = our own UID
    - Effective UID = **daemon**
    - Saved set-UID = daemon

```
→ ~ ls -l /usr/bin/at
-rwsr-sr-x 1 daemon daemon 51464 1月 15 2016 /usr/bin/at
```

```
→ ~ at 5:00
warning: commands will be executed using /bin/sh
at> date >tmp.txt
at> <EOT>
job 22 at Tue Oct 10 05:00:00 2017
```

**In general, we try to use the *least-privilege* model when we design our applications.**

2. **reduce** privileges. it calls `seteuid()`, only e-UID is changed:

- Real UID = our own UID (unchanged!)
- Effective UID = **our own UID**
- Saved set-UID = **daemon** (unchanged!)

Runs with our own UID as effective UID, can access only our own normally accessed files.

No additional permissions.

```
→ ~ sudo ls -l /var/spool/cron
total 12
drwxrwx--T 2 daemon daemon 4096 10月 9 16:40 atjobs
```

3. **Increase** privileges to access the configuration files that control which commands are to be run and the time at which they need to run. These files are owned by the daemon that will run the. The at command calls **seteuid** to set the effective user ID to daemon. **This call is allowed because** the argument to seteuid equals the saved set-user-ID. (This is why we need the saved set-user-ID.) After this, we have:

- Real UID = our own UID (unchanged!)
- Effective UID = **daemon**
- Saved set-UID = daemon (unchanged!)

4. After the files are modified to record the commands to be run and the time at which they are to be run, the at command **lowers** its privileges by calling **seteuid** to set its effective user ID to our user ID. This prevents any accidental misuse of privilege. At this point, we have

- Real UID = our own UID (unchanged!)
- Effective UID = **our own UID**
- Saved set-UID = daemon (unchanged!)

```
→ ~ ls -l /usr/sbin/atd  
-rwxr-xr-x 1 root root 26632 1月 15 2016 /usr/sbin/atd
```

```
→ ~ ps -U root | grep atd  
695 ?          00:00:00 atd  
→ ~ ps -u daemon | grep atd  
695 ?          00:00:00 atd
```

**Real uid = root**

**Effective uid = daemon**

5. The `daemon atd` starts out running with root privileges. To run commands on our behalf, the daemon calls `fork` and the child calls `setuid` to change its user ID to our user ID. Because the child is running with root privileges, this **changes all of the IDs**. We have

- Real UID = our own UID
- Effective UID = **our own UID**
- Saved set-UID = our own UID

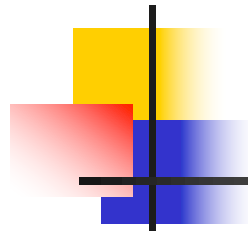
Now the daemon can safely **execute** commands.



# Changing 3 UIDs

ID	exec		setuid( <i>uid</i> )	
	set-user-ID bit off	set-user-ID bit on	superuser	unprivileged user
real user ID	unchanged	unchanged	set to <i>uid</i>	unchanged
effective user ID	unchanged	set from user ID of program file	set to <i>uid</i>	set to <i>uid</i>
saved set-user ID	copied from effective user ID	copied from effective user ID	set to <i>uid</i>	unchanged

```
→ ~ ls -l /usr/bin/at
-rwsr-sr-x 1 daemon daemon 51464 1月 15 2016 /usr/bin/at
```



# setreuid(), setregid()

---

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
int setreuid(uid_t ruid, uid_t euid);
```

```
int setregid(gid_t rgid, gid_t egid);
```

- Return: 0 if OK, -1 on error
- Sets the real user ID of the process to *ruid* and the effective user ID to *eu*id.
- If argument is -1, leave that ID unchanged.



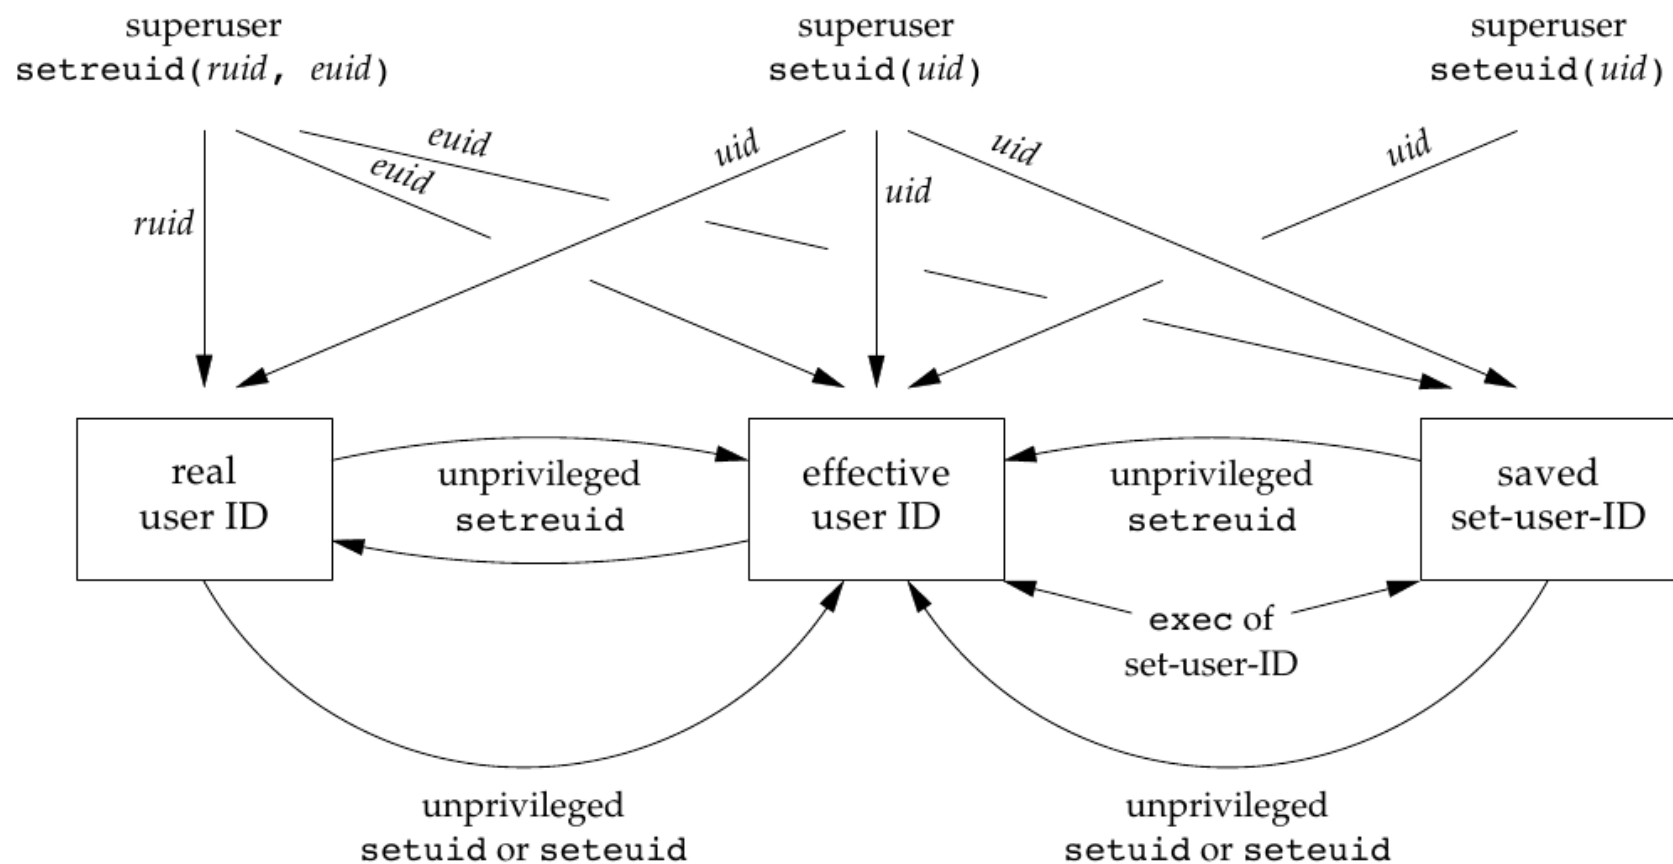


# setreuid

---

- `int setreuid(uid_t ruid, uid_t euid);`
- The **setreuid()** function *has been* used to *swap* the real and effective user IDs in set-user-ID programs to temporarily relinquish the set-user-ID value.
- This purpose is *now* better served by the use of the **seteuid()** function.

# Summary of set ID functions





## 13. system Function

---

```
#include <stdlib.h>
```

```
int system ( const char *cmdstring );
```

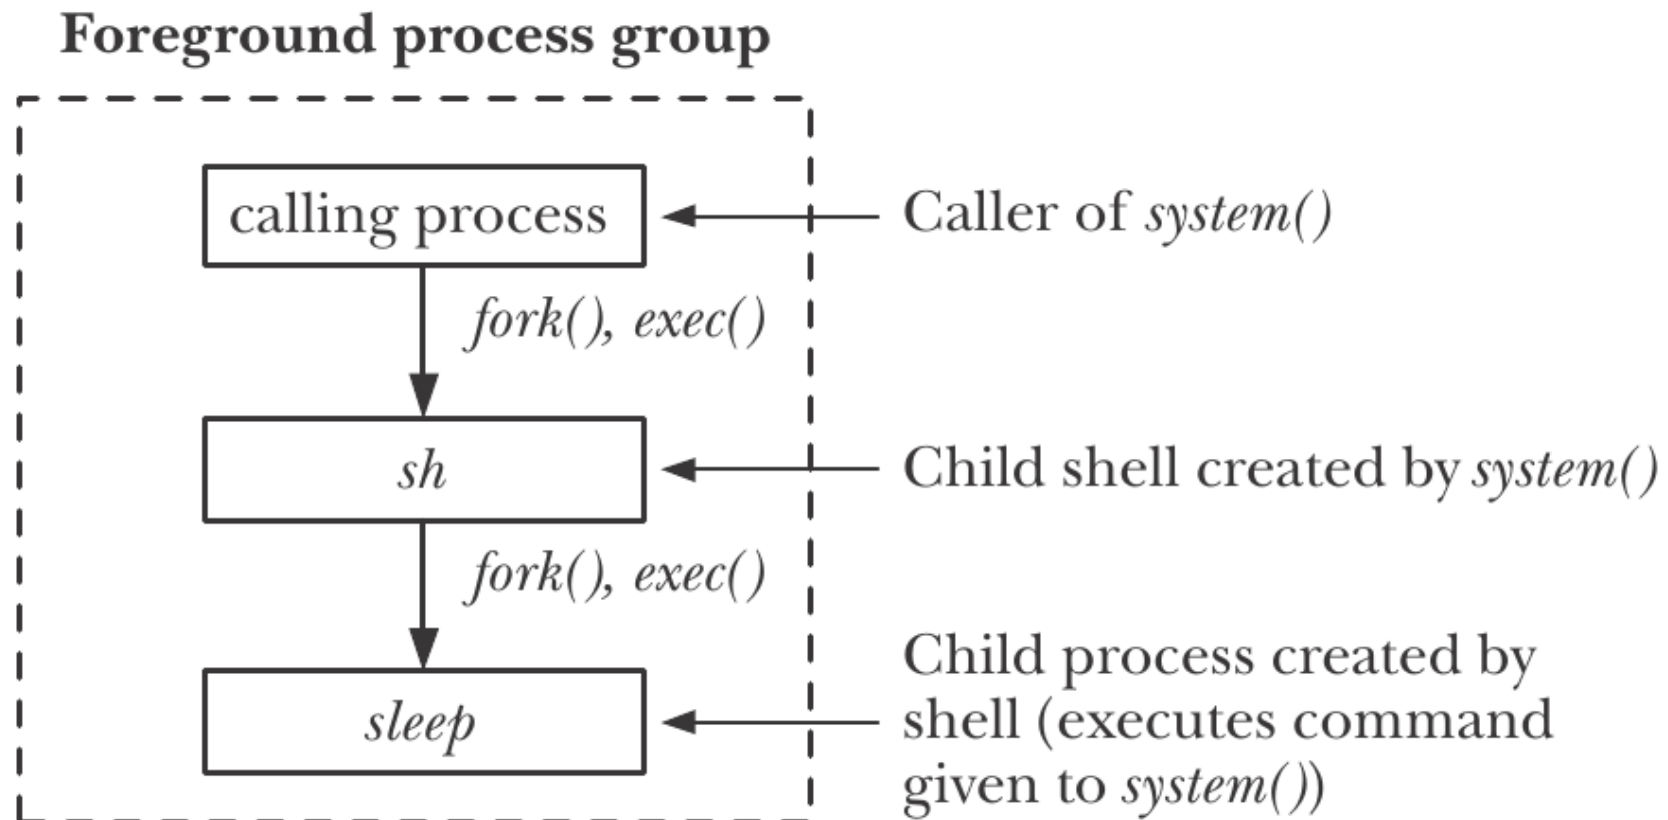
uses **fork** to create a child process that executes the shell command using **execl** :

```
execl("/bin/sh", "sh", "-c", command, (char *) 0);
```

**system()** returns after the command has been completed.

# System example

Arrangement of processes during execution of `system( "sleep 20" )`





## Program 8.22: system implement

---

```
#include      <sys/types.h>
#include      <sys/wait.h>
#include      <errno.h>
#include      <unistd.h>

int system(const char *cmdstring)
    /* version without signal handling */
{
    pid_t      pid;
    int        status;

    if (cmdstring == NULL)
        return(1); /* always a command processor with Unix */

    if ( (pid = fork()) < 0) {
        status = -1; /* probably out of processes */
    }
}
```



## Program 8.22

---

```
    else if (pid == 0) {                /* child */
        execl("/bin/sh", "sh", "-c", cmdstring, (char *) 0);
        _exit(127);                    /* execl error */
    } else {                            /* parent */
        while (waitpid(pid, &status, 0) < 0)
            if (errno != EINTR) {
                status = -1; /* error other than EINTR */
                break;
            }
    }
    return(status);
}
```

to prevent child  
flushing buffer



## Program 8.23: calling system

---

```
#include    <sys/types.h>
#include    <sys/wait.h>
#include    "apue.h"
int main(void) {
    int      status;
    if ( (status = system("date")) < 0 )
        err_sys("system() error");
    pr_exit(status);
    if ( (status = system("nosuchcommand")) < 0 )
        err_sys("system() error");
    pr_exit(status);
    if ( (status = system("who; exit 44")) < 0 )
        err_sys("system() error");
    pr_exit(status);
    exit(0);
}
```



## Program 8.23: results

**\$ a.out**

{ Thu Aug 29 14:24:19 MST 1991

normal termination, exit status = 0 **for date**

{ sh: nosuchcommand: not found

normal termination, exit status = 1 **for nosuchcommand**

{ stevens console Aug 25 11:49

stevens tty0 Aug 29 05:56

stevens tty1 Aug 29 05:56

stevens tty2 Aug 29 05:56

normal termination, exit status = 44 **for exit**





# Problem

---

- What happens if we call **system** from a **set-user-ID program**?
- A security hole!
- Should never be done!

Compile into program **tsys**

## Program 8.24: system from cmd

```
#include "apue.h"
int main(int argc, char *argv[]) {
    int status;
    if (argc < 2)
        err_quit("command-line argument required");
    if ( (status = system(argv[1])) < 0)
        err_sys("system() error");
    pr_exit(status);
    exit(0);
}
```

Compile into program **printuids**



## Program 8.25: print UIDs

```
#include "apue.h"

int
main(void)
{
    printf("real uid = %d, effective\n", getuid(), geteuid());
    exit(0);
}
```



# Program 8.24, 8.25: results

**\$ tsys printuids**

real uid = 224, **effective uid = 224**

normal termination, exit status = 0

**make tsys set-user-ID**

**\$ su**

Password:

**# chown root tsys**

**# chmod u+s tsys**

**# ls -l tsys**

-rwsrwxr-x 1 root 105737 Aug 18 11:21 tsys

**# exit**

**\$ tsys printuids**

real uid = 224, **effective uid = 0**

**this is a security hole**

normal termination, exit status = 0



# 17. Process Times

---

```
#include <sys/times.h>
```

```
clock_t times(struct tms *buf);
```

- Returns: elapsed wall clock time in clock ticks if OK, -1 on error

```
struct tms {
```

```
    clock_t tms_utime; /* user CPU time */
```

```
    clock_t tms_stime; /* system CPU time */
```

```
    clock_t tms_cutime; /* sum of user time for terminated children */
```

```
    clock_t tms_cstime; /* sum of system time for terminated children */
```

```
};
```



## Program 8.30 (main())

---

```
#include    <sys/times.h>
#include    "apue.h"

static void    pr_times(clock_t, struct tms *, struct tms *);
static void    do_cmd(char *);

int main(int argc, char *argv[]) {
    int        i;

    for (i = 1; i < argc; i++)
        do_cmd(argv[i]);    /*once each command-line arg */
    exit(0);
}
```



## Program 8.30 do\_cmd()

---

```
static void do_cmd(char *cmd)    /*execute and time the "cmd" */ {
    struct tms    tmsstart, tmsend;
    clock_t      start, end;
    int          status;
    fprintf(stderr, "\ncommand: %s\n" , cmd);
    if ( (start = times(&tmsstart)) == -1)    /* starting values */
        err_sys("times error");
    if ( (status = system(cmd)) < 0)/* execute command */
        err_sys( "system() error" );
    if ( (end = times(&tmsend)) == -1)        /* ending values */
        err_sys( "times error" );
    pr_times(end-start, &tmsstart, &tmsend);
    pr_exit(status);
}
```

# Program 8.30 pr\_times()

```
static void pr_times(clock_t real, struct tms *tmsstart, struct tms *tmsend)
{
    static long          clktck = 0;
    if (clktck == 0) /* fetch clock ticks per second first time */
        if ( (clktck = sysconf(_SC_CLK_TCK)) < 0)
            err_sys("sysconf error");
    fprintf(stderr, " real:  %7.2f\n", real / (double) clktck);
    fprintf(stderr, " user:  %7.2f\n",
        (tmsend->tms_utime - tmsstart->tms_utime) / (double) clktck);
    fprintf(stderr, " sys:   %7.2f\n",
        (tmsend->tms_stime - tmsstart->tms_stime) / (double) clktck);
    fprintf(stderr, " child user: %7.2f\n",
        (tmsend->tms_cutime - tmsstart->tms_cutime) / (double) clktck);
    fprintf(stderr, " child sys:  %7.2f\n",
        (tmsend->tms_cstime - tmsstart->tms_cstime) / (double) clktck);
}
```





## Program 8.30: results

---

```
→ proc git:(master) X ./times1 "sleep 3"
```

```
command: sleep 3
```

```
real:    3.00
```

```
user:    0.00
```

```
sys:     0.00
```

```
child user:    0.00
```

```
child sys:     0.00
```

```
normal termination, exit status = 0
```

```
→ proc git:(master) X ./times1 "find /usr/include -type f -exec wc {} \; >/dev/null"
```

```
command: find /usr/include -type f -exec wc {} \; >/dev/null
```

```
real:    36.68
```

```
user:    0.00
```

```
sys:     0.00
```

```
child user:    0.74
```

```
child sys:     2.69
```

```
normal termination, exit status = 0
```