

Chapter 15.

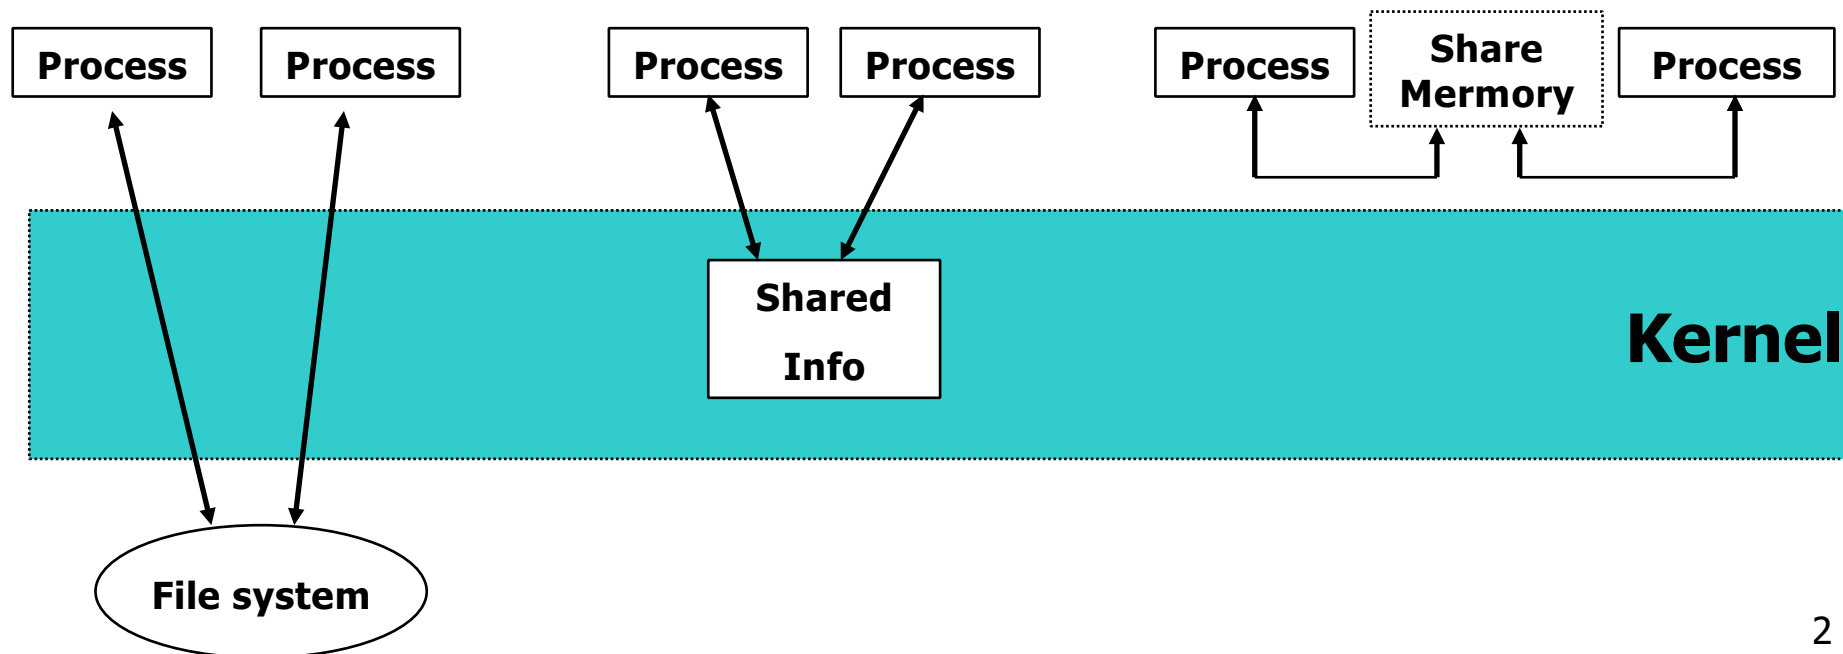
Interprocess Communication

朱金辉

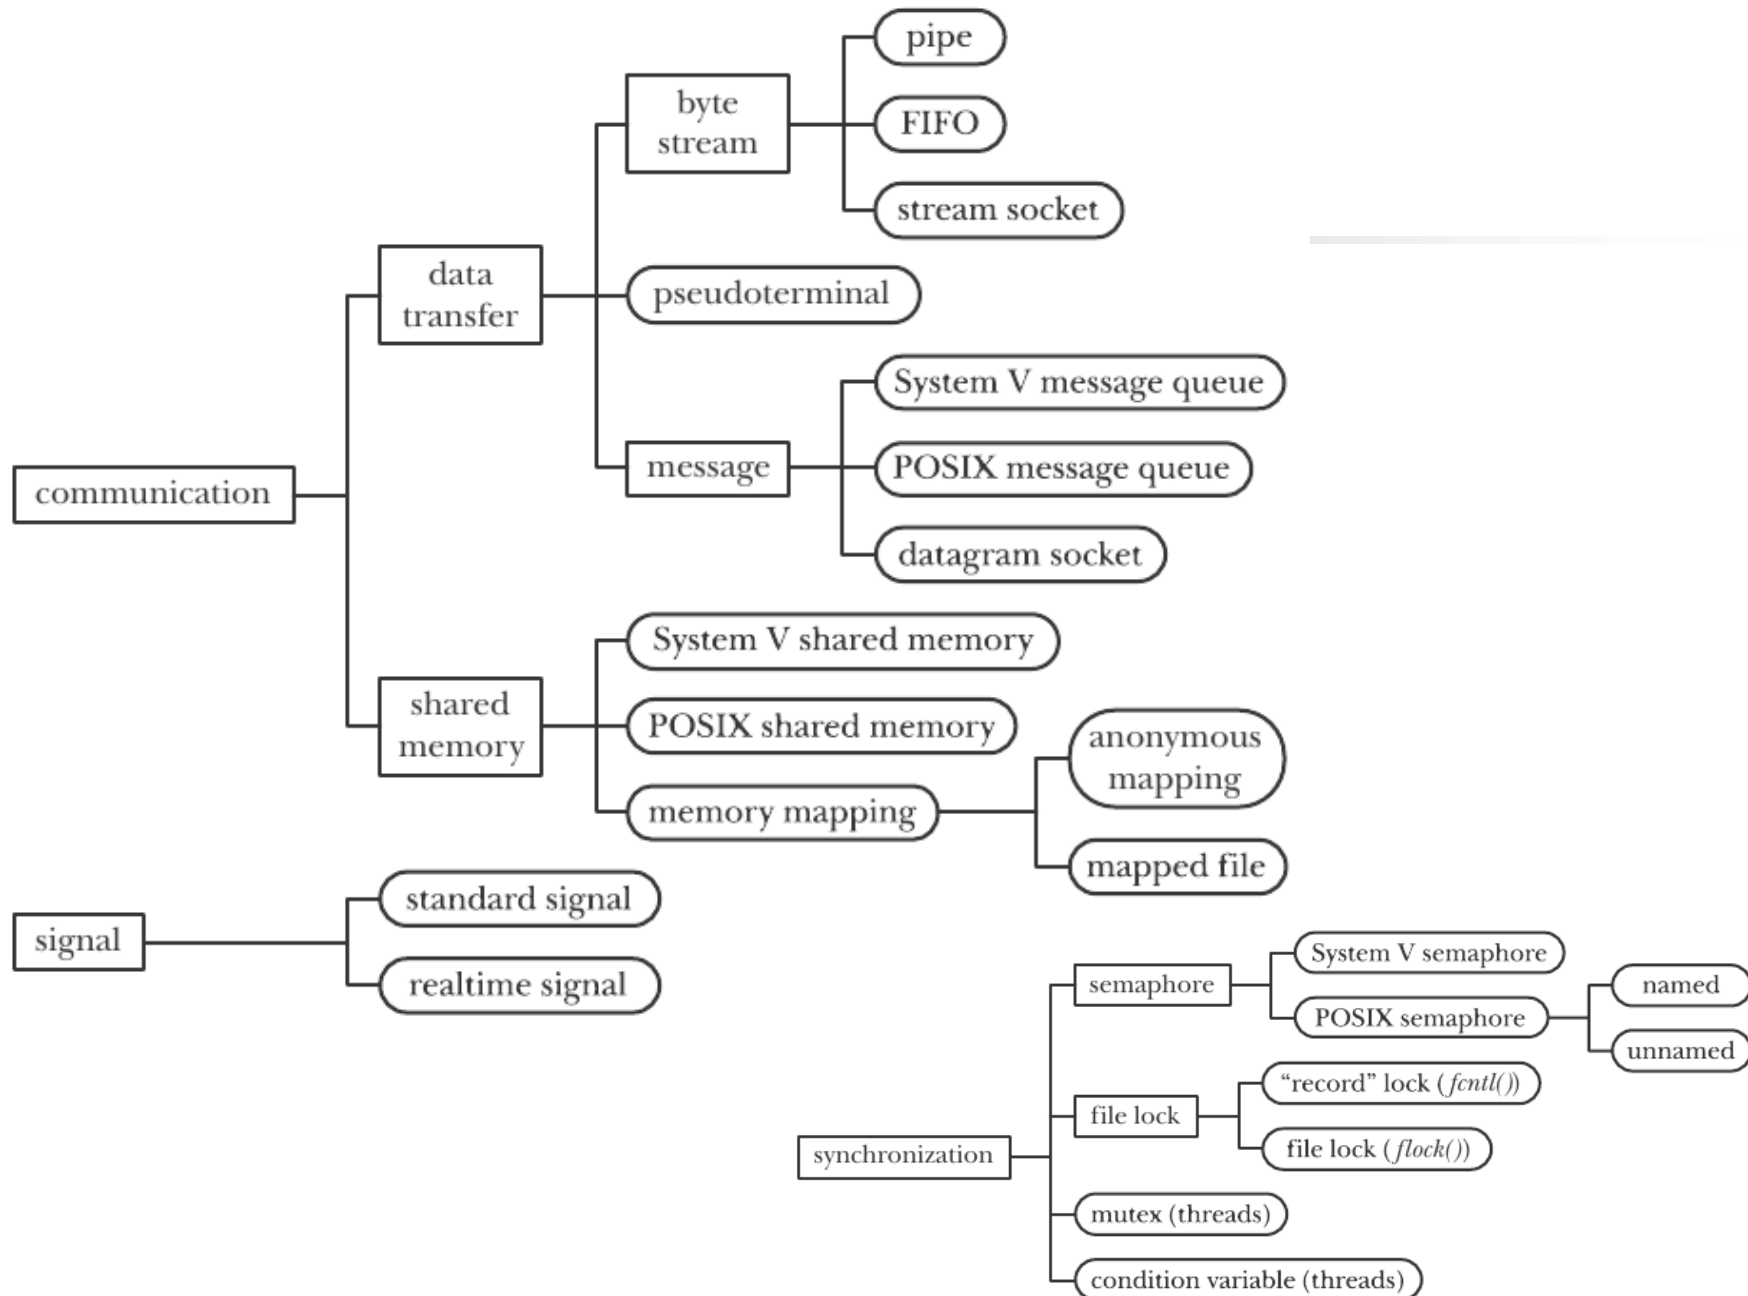
华南理工大学软件学院

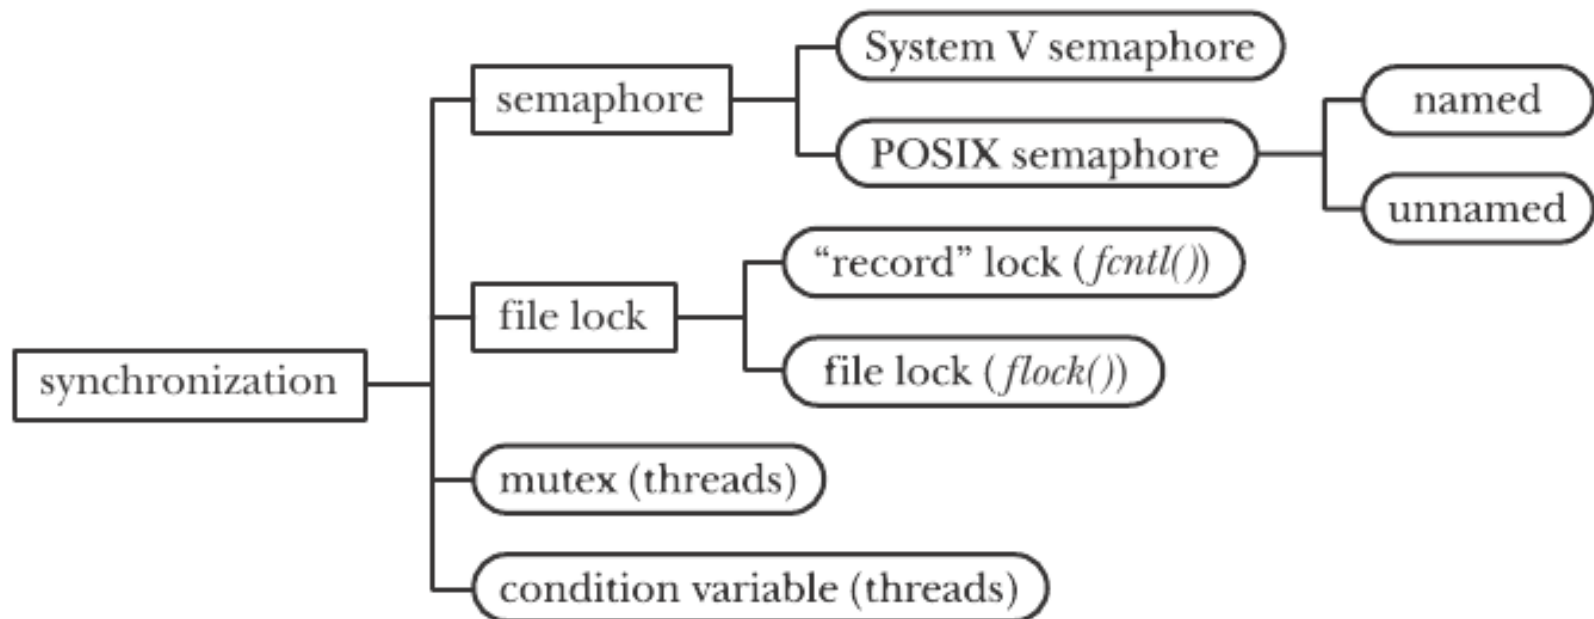
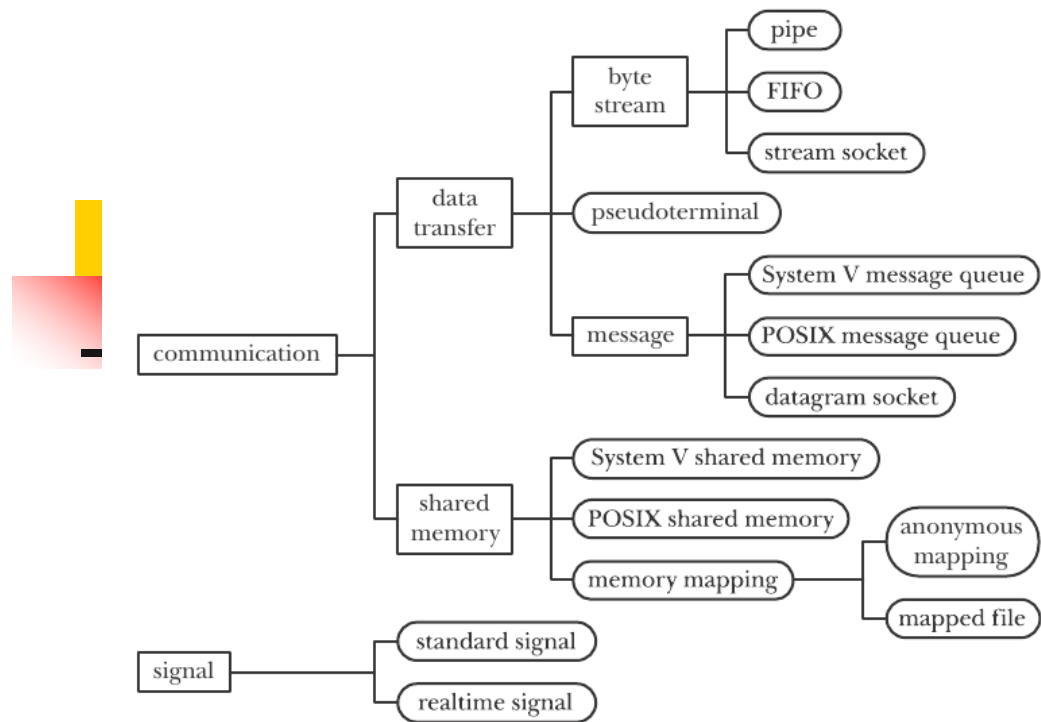
UNIX IPC

- IPC stands for interprocess communication
- Message passing between different process that are running on some operating system
- Need some forms of synchronization



A taxonomy of UNIX IPC facilities








Contents

- Pipes
 - FIFO: name pipes
- Message Queues
- Semaphores
- Shared Memory
- Memory mapping

1. Pipes



Exchanging data between two processes using a pipe



Pipes


- Oldest form
- half-duplex (only 1-direction data flow)
- no names
- process creates a pipe, forks, and uses the pipe to communicate with child

```
#include <unistd.h>


int pipe (int fieldes[2]);
```
- Returns: 0 if OK, -1 on error

Pipes


- `filedes[0]`: open for reading
- `filedes[1]`: open for writing
- Output of `filedes[1]` is input of `filedes[0]`






or



Half-duplex pipe after a fork







Program 15.5: Parent→Child Pipe

```
int main(void) {
    int          n, fd[2];
    pid_t        pid;
    char line[MAXLINE];
    if (pipe(fd) < 0)
        err_sys("pipe error");
    if ( (pid = fork()) < 0) err_sys("fork error");
    else if (pid > 0) {          /* parent */
        close(fd[0]);
        write(fd[1], "hello world\n", 12);
    } else {                    /* child */
        close(fd[1]);
        n = read(fd[0], line, MAXLINE);
        write(STDOUT_FILENO, line, n);
    }
    exit(0);
}
```




Pipes

- **Read** from a pipe if write end closed:
 - if data, data is read, and
 - when no data, returns 0 (➔ EOF)
- **Write** to a pipe if read end closed:
 - **SIGPIPE** generated,
 - write() returns errno = EPIPE
- PIPE_BUF: #bytes in kernel's pipe buffer size

Using 2 pipes for parent/child synchronization

Figure 15.8. Using two pipes for parentchild synchronization





Program 15.7: Pipe sync

```
#include    "apue.h"

static int  pfd1[2], pfd2[2];

void
TELL_WAIT() {
    if (pipe(pfd1) < 0 || pipe(pfd2) < 0)
        err_sys("pipe error");
}

Void TELL_PARENT(pid_t pid) {
    if (write(pfd2[1], "c", 1) != 1)
        err_sys("write error");
}
```



Program 15.7: Pipe sync

```
void WAIT_PARENT(void) {
    char    c;
    if (read(pfd1[0], &c, 1) != 1) err_sys("read error");
    if (c != 'p') err_quit("WAIT_PARENT: incorrect data");
}

void TELL_CHILD(pid_t pid) {
    if (write(pfd1[1], "p", 1) != 1) err_sys("write error");
}

void WAIT_CHILD(void) {
    char    c;
    if (read(pfd2[0], &c, 1) != 1) err_sys("read error");
    if (c != 'c')
        err_quit("WAIT_CHILD: incorrect data");
}
```



popen & pclose Functions

- All dirty work:
 - create a pipe
 - fork a child
 - close unused ends of pipe
 - exec a shell to execute a command
 - wait for command to terminate
- #include <stdio.h>
`FILE *popen(const char *cmdstring, const char *type);`
- Returns: file pointer if OK, NULL on error
`int pclose(FILE *fp);`
- Returns: termination status of cmdstring, or -1 on error

popen: type = r or w



Figure 15.9 Result of `fp = popen(cmdstring, "r")`



Figure 15.10 Result of `fp = popen(cmdstring, "w")`

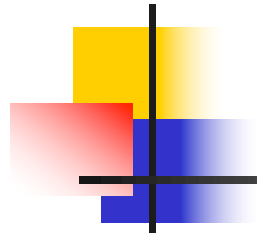


Program 15.11: pager (popen)

```
#include    <sys/wait.h>
#include    "apue.h"

#define     PAGER "${PAGER:-more}" /* environment
    variable, or default */

int main(int argc, char *argv[]) {
    char    line[MAXLINE];
    FILE    *fpin, *fpout;
    if (argc != 2)
        err_quit("usage: a.out <pathname>");
    if ( (fpin = fopen(argv[1], "r")) == NULL)
        err_sys("can't open %s", argv[1]);
```

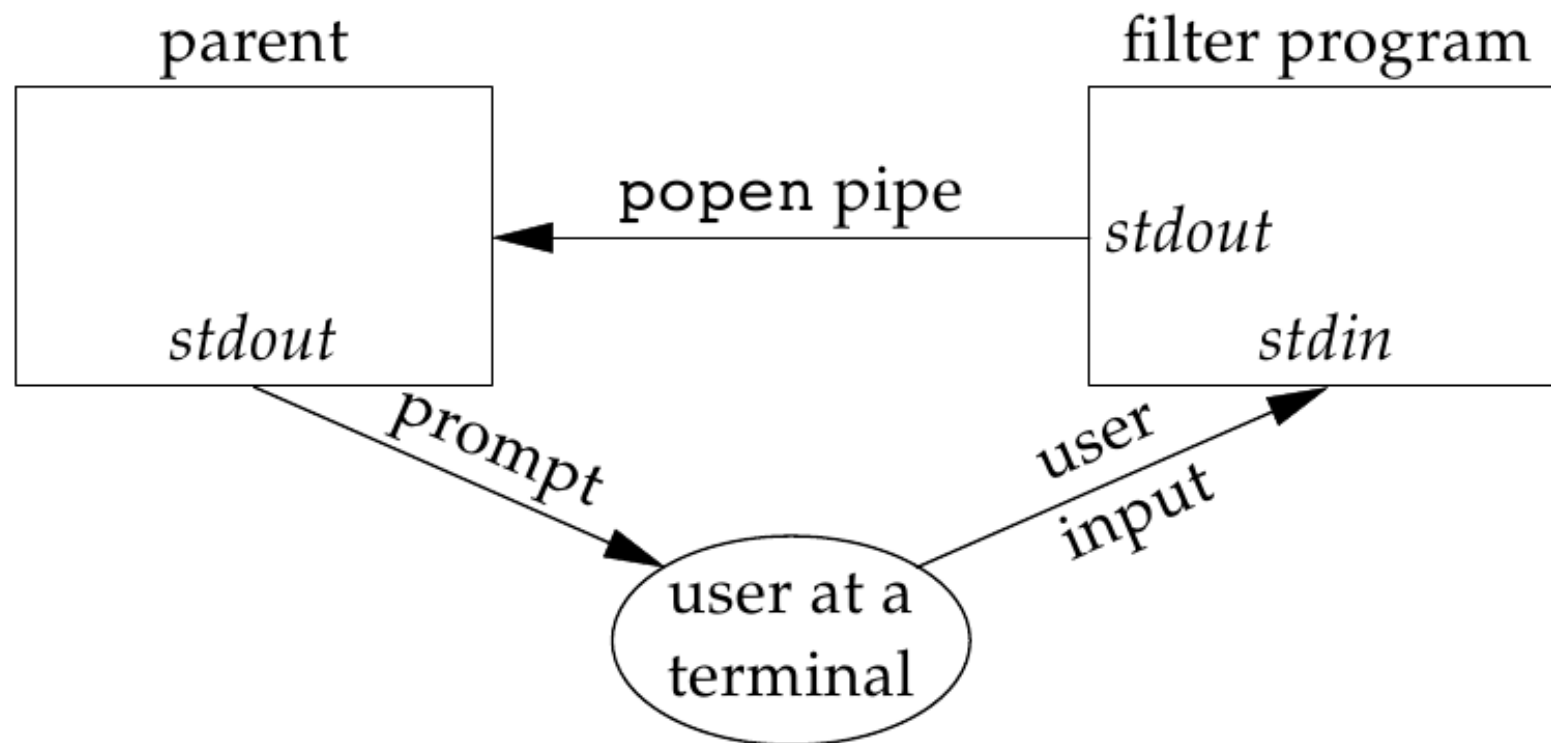


Program 15.11: pager (popen)

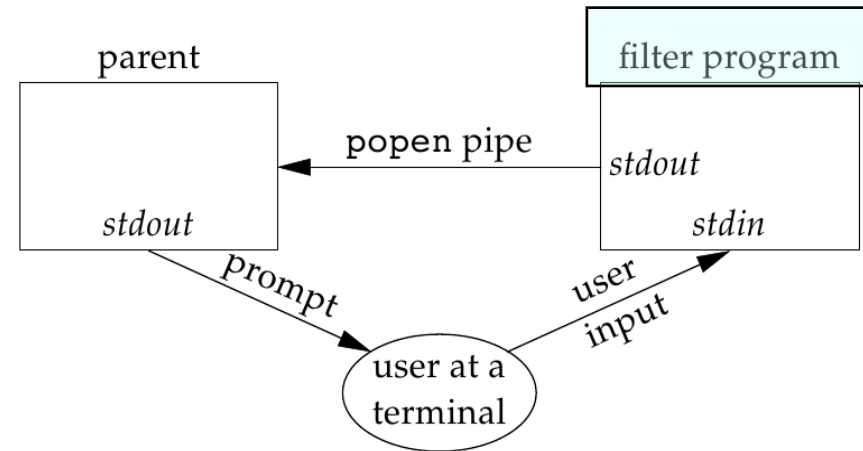
```
if ( (fpout = popen(PAGER, "w")) == NULL)
    err_sys("popen error");

    /* copy argv[1] to pager */
while (fgets(line, MAXLINE, fpin) != NULL) {
    if (fputs(line, fpout) == EOF)
        err_sys("fputs error to pipe");
    }
if (ferror(fpin))
    err_sys("fgets error");
if (pclose(fpout) == -1)
    err_sys("pclose error");
exit(0);
}
```

Filter Program (with popen)



Program 15.14 (filter)

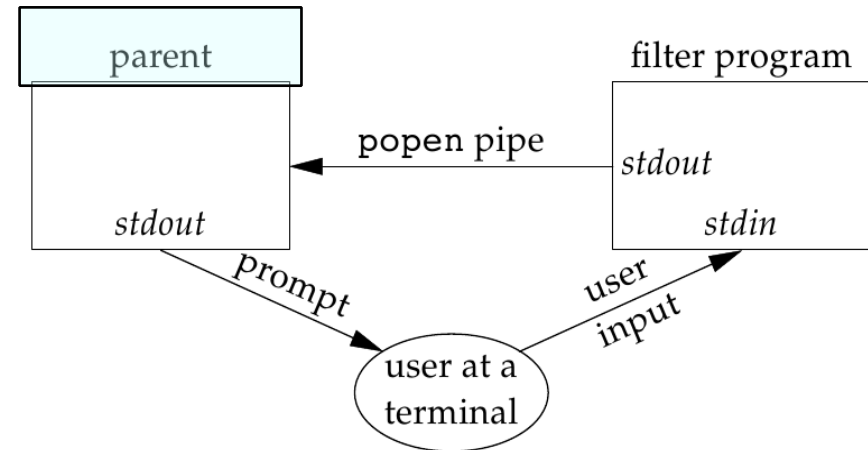


```
#include    <ctype.h>
#include    "apue.h"

int main(void) {
    int      c;
    while ( (c = getchar()) != EOF) {
        if (isupper(c))
            c = tolower(c);
        if (putchar(c) == EOF)
            err_sys("output error");
        if (c == '\n')
            fflush(stdout);
    }
    exit(0);
}
```

Program 15.15 (using filter)

```
#include <sys/wait.h>
int main(void) {
    char line[MAXLINE];
    FILE *fpin;
    if ( (fpin = popen("myucl", "r")) == NULL)
        err_sys("popen error");
    for ( ; ; ) {
        fputs("prompt> ", stdout);
        fflush(stdout);
        if (fgets(line, MAXLINE, fpin) == NULL)        /* read from pipe */
            break;
        if (fputs(line, stdout) == EOF)
            err_sys("fputs error to pipe");
    }
    if (pclose(fpin) == -1) err_sys("pclose error");
    putchar('\n');
    exit(0);
}
```





2. FIFOs

- Also called “named pipes”
- Can be used by unrelated processes
- File type: `stat.st_mode = FIFO`
- Test with `S_ISFIFO` macro
- Similar to creating a file : `mkfifo`
- Pathname exists
- `open`, `close`, `read`, `write`, `unlink`, etc



FIFOs

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```


```
int mkfifo (const char *pathname,  
mode_t mode);
```

- Returns: 0 if OK, -1 on error
- mode: same as for open function


Using FIFOs to Duplicate Output Streams

- `mkfifo fifo1`
- `prog3 < fifo1 &`
- `prog1 < infile | tee fifo1 | prog2`

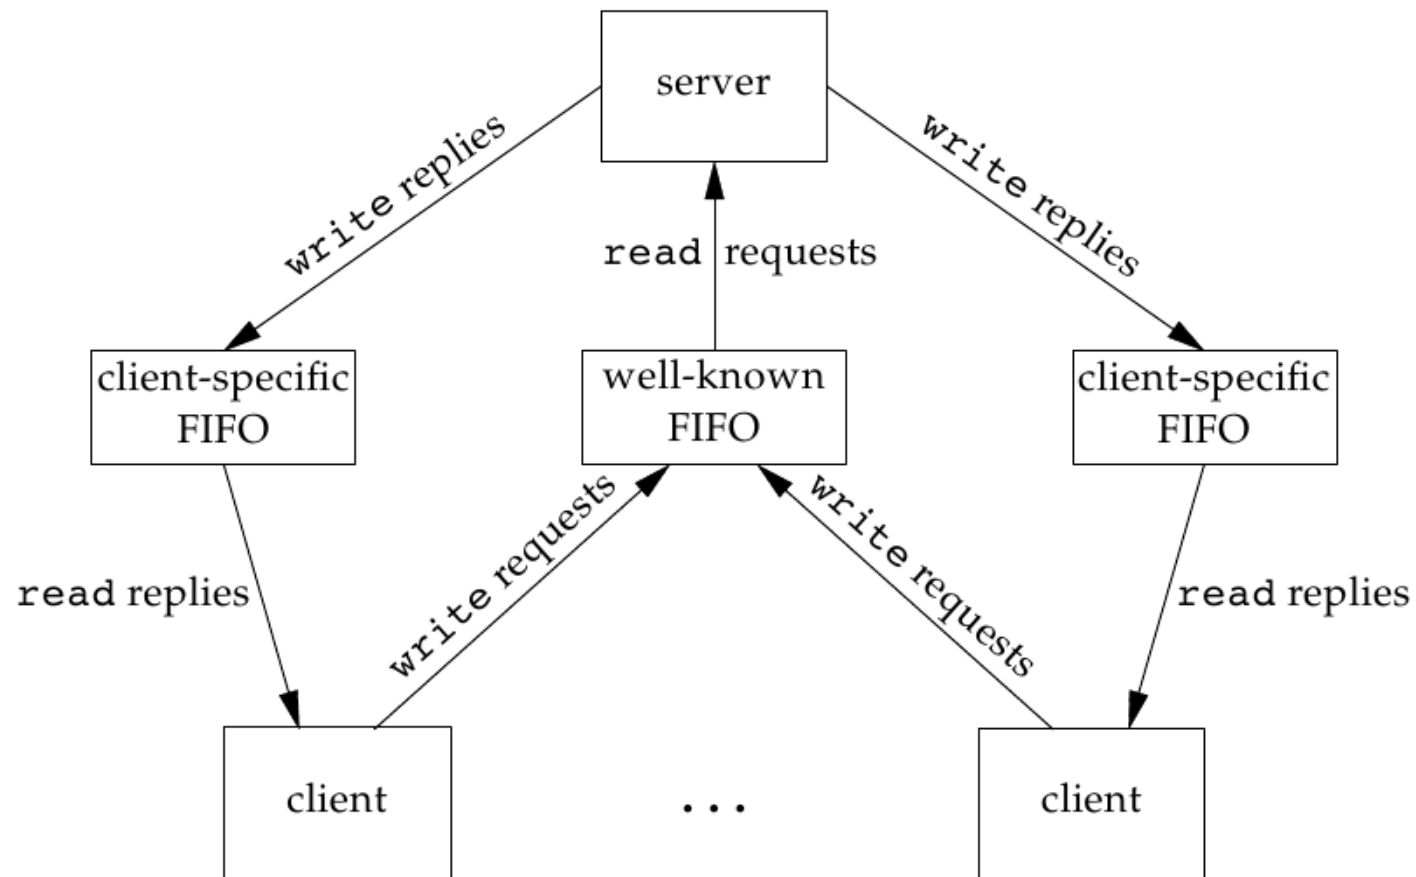
**Copies stdin to
stdout and file**



Client-Server Communication Using a FIFO



Client-Server Communication Using a FIFO





Client-Server Communication Using a FIFO

- Solution: Create a FIFO for each client such that server can reply using the client-specific FIFO
 - **E.g.** /tmp/serv1.XXXX, where XXXXX is client's process ID
 - Impossible for server to know if a client has crashed, FIFOs left in system,
 - server must catch SIGPIPE (FIFO with 1 writer, no reader)

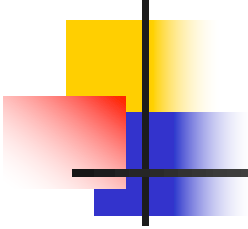


3. XSI IPC

XSI : **X**/Open **S**ystem **I**nterface extension

The XSI IPC functions are based closely on the **System V IPC** functions.

- ✓ Message Queues
- ✓ Semaphores
- ✓ Shared Memory
- When creating an IPC structure, a **key** must be specified (type: **key_t**)
- Each IPC structure has a nonnegative integer **identifier** (large!)
- Command: **ipcs**, **ipcrm**



Client-Server Rendezvous at same IPC structure (1)

- Server creates a new IPC structure using `key = IPC_PRIVATE`
- Server stores returned identifier in some file for client to obtain
- Disadvantage: file I/O!



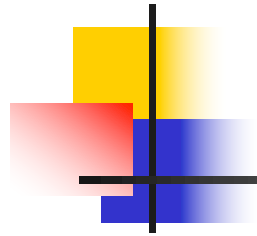
Client-Server Rendezvous at same IPC structure (2)

- Define a key in a common header
- Client and server agree to use that key
- Server creates a new IPC structure using that key
- **Problem:** key exists? (msgget, semget, shmget returns error)
- **Solution:** delete existing key, create a new one again!



Client-Server Rendezvous at same IPC structure (3)

- Client and server agree on
 - a pathname
 - a project ID (char between 0 ~ 255)
- `ftok()` converts the 2 values into a key
- Client and server use that key
- **Disadvantage:** `ftok` → a function call!



Permission Structure for IPC

XSI IPC associates an ipc_perm structure with each IPC structure.

```
struct ipc_perm {  
    uid_t uid; /* owner's EUID */  
    gid_t gid; /* owner's EGID */  
    uid_t cuid; /* creator's EUID */  
    gid_t cgid; /* creator's EGID */  
    mode_t mode; /* access mode */  
    ulong seq; /* slot usage seq number */  
    key_t key; /* key */  
};
```



Summary of XSI IPC functions

	Message queues	Semaphores	Shared Memory
Header	<sys/msg.h>	<sys/sem.h>	<sys/shm.h>
create or open	msgget	semget	shmget
Control operations	msgctl	semctl	shmctl
IPC operations	msgsnd msgrcv	semop	shmat shmdt



(1) Message Queues

- Linked list of messages
- Stored in kernel
- Identified by message queue identifier
- msgget: create new or open existing q
- msgsnd: add new msg to a queue
- msgrcv: receive msg from a queue
- Fetching order: based on type of msg



Message Queues

- creates a new message queue or obtains the identifier of an existing queue
- `#include <sys/types.h>`
`#include <sys/ipc.h>`
`#include <sys/msg.h>`
`int msgget(key_t key, int flag);`
- Returns: msg queue ID if OK, -1 on error

```
id = msgget(key, IPC_CREAT | S_IRUSR | S_IWUSR);  
if (id == -1)  
    errExit("msgget");
```



Operations on queue

```
#include <sys/types.h>
```

```
#include <sys/ipc.h>
```

```
#include <sys/msg.h>
```

```
int msgctl(int msqid, int cmd, struct  
msqid_ds *buf);
```

IPC_STAT: fetch into buf
IPC_SET: set from buf
IPC_RMID: remove from
queue

- Returns: 0 if OK, -1 on error



Message Queues data structure

```
struct msqid_ds {
    struct ipc_perm  msg_perm;           /* see Section 15.6.2 */
    msgqnum_t        msg_qnum;           /* # of messages on queue */
    msglen_t          msg_qbytes;        /* max # of bytes on queue */
    pid_t             msg_lspid;         /* pid of last msgsnd() */
    pid_t             msg_lrpid;         /* pid of last msgrcv() */
    time_t            msg_stime;         /* last-msgsnd() time */
    time_t            msg_rtime;         /* last-msgrcv() time */
    time_t            msg_ctime;         /* last-change time */
    :
};
```



Message Structure

```
struct mymesg {  
    long mtype;           /* positive msg type */  
    char mtext[512];      /* message data */  
}
```



Place data on message queue

```
#include <sys/types.h>
```

```
#include <sys/ipc.h>
```

```
#include <sys/msg.h>
```

```
int msgsnd(int msqid, const void *ptr,  
           size_t nbytes, int flag);
```

- Returns: 0 if OK, -1 on error

**Pointer to:
type + data (nbytes data)**

IPC_NOWAIT or 0



Retrieve Message from Queue

```
#include <sys/types.h>
```

```
#include <sys/ipc.h>
```

```
#include <sys/msg.h>
```

```
int msgrcv(int msqid, void *ptr,  
           size_t nbytes, long type, int flag);
```

type == 0: 1st msg

type > 0: 1st msg of type

**type < 0: first msg with type
≤ |type| and smallest**

- Returns: data size in message if OK, -1 on error



Example(1/4)

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#include <stdio.h>
#define MSGSZ 128
struct msgbuf {
    long mtype;
    char mtext[MSGSZ];
};

int main () {
    int msgid;
    key_t key;
    struct msgbuf sbuf, rbuf;
    key=4587;
```



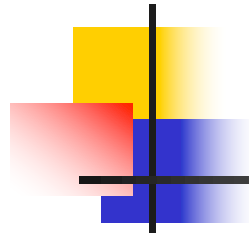
Example(2/4)

```
    if ((msgid=msgget(key, IPC_CREAT|0666))<0) {
        printf("msgget");
        exit(1);
    }
    if (msgrcv(msgid,&rbuf,MSGSZ,0,0)<0) {
        printf("msgrcv");
        exit(1);
    }
    sbuf.mtype=2;
    sprintf(sbuf.mtext,"I received the message.");
    if (msgsnd(msgid,&sbuf,strlen(sbuf.mtext)+1,0)<0) {
        printf("msgsnd");
        exit(1);
    }
    exit(0);
}
```



Example(3/4)

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#include <stdio.h>
#define MSGSZ 128
struct msgbuf {
    long mtype;
    char mtext[MSGSZ];
};
int main () {
    int msgid;
    key_t key;
    struct msgbuf sbuf, rbuf;
    key=4587;
```



Example(4/4)

```
if ((msgid=msgget(key, 0666)) < 0) {  
    printf("msgget");  
    exit(1);  
}  
sbuf.mtype=1;  
sprintf(sbuf.mtext, "Did you get the Message ?");  
if (msgsnd(msgid, &sbuf, strlen(sbuf.mtext)+1, 0) {  
    printf("msgsnd");  
    exit(1);  
}  
if (msgrcv(msgid, &rbuf, MSGSZ, 2, 0) < 0) {  
    printf("msgrcv");  
    exit(1);  
}  
printf("%s\n", rbuf.mtext);  
exit(0);  
}
```



(2) Semaphores

- Semaphores are intended to let multiple processes **synchronize** their operations.
- A **counter** to provide access to shared data object for multiple processes
- To object a shared resource:
 - Test semaphore controlling resource
 - If value > 0 , value--, grant use
 - If value $== 0$, sleep until value > 0
 - Release resource, value ++

**Need to
be
ATOMIC**



Complication

- Three features contribute to this unnecessary complication.
 1. A semaphore is defined as **a set of** one or more **semaphore values**
 2. Creation (semget) is independent of initialization (semctl). **Cannot atomically create** a new semaphore set and **initialize** all the values in the set.
 3. All IPCs exist even if no process is using them. Need **worry about** process terminating without releasing semaphore.



Semaphore structure

```
struct semid_ds {
    struct ipc_perm  sem_perm; /* see Section 15.6.2 */
    unsigned short   sem_nsems; /* # of semaphores in set */
    time_t           sem_otime; /* last-semop() time */
    time_t           sem_ctime; /* last-change time */
    :
};
```

Each semaphore is represented by an anonymous structure containing at least the following members:

```
struct {
    unsigned short   semval; /* semaphore value, always >= 0 */
    pid_t            sempid; /* pid for last operation */
    unsigned short   semncnt; /* # processes awaiting semval>curval */
    unsigned short   semzcnt; /* # processes awaiting semval==0 */
    :
};
```




Obtain a semaphore

- `#include <sys/types.h>`
- `#include <sys/ipc.h>`
- `#include <sys/sem.h>`
- `int semget (key_t key, int nsems, int flag);`
- Returns: sem ID if OK, -1 on error



Semaphore control

```
#include <sys/types.h>
```

```
#include <sys/ipc.h>
```

```
#include <sys/sem.h>
```

```
int semctl (int semid, int semnum, int cmd, union semun arg);
```

■ **cmd**: stat, set, rmid, getval, setval, getall, setall ...

```
union semun {  
    int          val;      /* for SETVAL */  
    struct semid_ds *buf;   /* for IPC_STAT and IPC_SET */  
    unsigned short *array; /* for GETALL and SETALL */  
};
```



Semaphore operation(1)

- `#include <sys/types.h>`
- `#include <sys/ipc.h>`
- `#include <sys/sem.h>`
- `int semop(int semid, struct sembuf semoparray[], size_t nops);`
- Return: 0 if ok, -1 on error. This is an atomic operation

```
struct sembuf {
    unsigned short sem_num; /* member # in set (0, 1, ..., nsems-1) */
    short          sem_op;  /* operation (negative, 0, or positive) */
    short          sem_flg; /* IPC_NOWAIT, SEM_UNDO */
};
```



Struct of sembuf

```
struct sembuf {  
    ushort sem_num;  
    /*member # in set (0,1,...,nsems-1)*/  
    short sem_op;  
    /*operation (negative, 0, or positive) */  
    short sem_flg;  
    /*IPC_NOWAIT,SEM_UNDO */  
}
```



Semaphore operation(2)

- `sem_op` is **positive**: release the resource hold by process, value of `sem_op` added to semaphore's value
- `sem_op` is **0**: the calling process wants to wait until the semaphore's value becomes 0.
- `sem_op` is **negative**: process want to obtain resources



Example(1/5)

```
int main {  
    int sem_id,i,creat=0,pause_time;  
    char *cp;  
    srand((unsigned int )getpid());  
    sem_id=open_semaphore_set((key_t)1234,1);  
    if (argc>1&&strcmp(argv[1],"1")) {  
        init_a_semaphore(sem_id,0,1);  
        creat=1;  
        sleep(2);  
    }
```



Example(2/5)

```
for(i=0;i<argc;i++) {
    cp=argv[i];
    if (!semaphore_P(sem_id)) exit(-1);
    printf("process %d:",getpid());
    fflush(stdout);
    while(*cp) {
        printf("%c",*cp); fflush(stdout);
        sleep(rand()%3);
        cp++;
    }
    printf("\n");
    if (semaphore_V(sem_id)) exit(-1);
    sleep(rand()%2);
}
```



Example(3/5)

```
printf("\n%d –finished\n",getpid());
if (creat==1) {
    sleep(10);
    rm_semaphore(sem_id);
}
exit(0);
}
int open_semaphore_set(key_t keyval,int numsems) {
    int sid;
    if (!numsems) return(-1);
    if ((sid=semget(keyval,numsems,IPC_CREAT|0660))== -1) return (-1);
    else return(sid);
}
```




Example(4/5)

```
int semaphore_P(int sem_id) {
    struct sembuf sb; sb.sem_num=0;
    sb.sem_op=-1; sb.sem_flag=SEM_UNDO;
    if (semop(sem_id,&sb,1)==-1 {
        fprintf(stderr,"semaphore_P failed\n"); return(0);
    }
    return(1);
}

int semaphore_V(int sem_id) {
    struct sembuf sb; sb.sem_num=0;
    sb.sem_op=1; sb.sem_flag=SEM_UNDO;
    if (semop(sem_id,&sb,1)==-1 {
        fprintf(stderr,"semaphore_V failed\n"); return(0);
    }
    return(1);
}
```



Example(5/5)

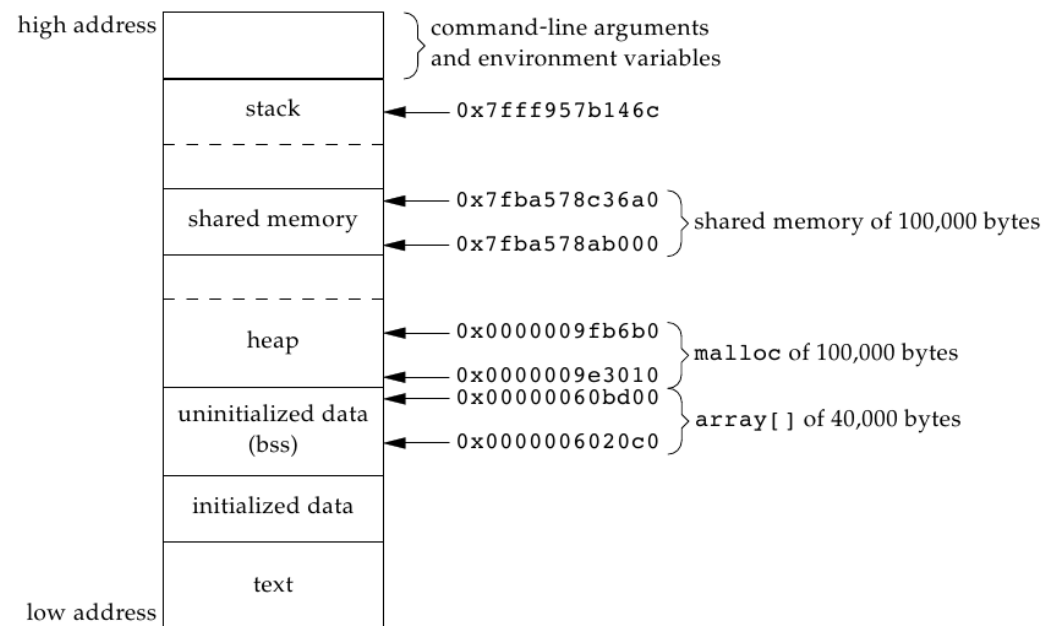
```
union semun {
    int val;
    struct semid_ds *buf;
    unsigned short *array;
};

void init_a_semaphore(int sid,int semnum,int initval) {
    union semun semopts;
    semopts.val=initval;
    semctl(sid,semnum,SETVAL,semopts);
}

int rm_semaphore(int sid) {
    return(semctl(sid,0,IPC_RMID,0));
}
```

(3) Shared Memory


- Fastest form of IPC:
no need of data
copying between
client & server
- Must synchronize
access to a shared
memory segment
- Semaphores are used





Obtain a shared memory id

- `#include <sys/types.h>`
- `#include <sys/ipc.h>`
- `#include <sys/shm.h>`
- `int shmget (key_t key, int size, int flag);`
- Returns: shared memory ID if OK, -1 on error



Shared Memory Operations

- `#include <sys/types.h>`
- `#include <sys/ipc.h>`
- `#include <sys/shm.h>`
- `int shmctl (int shmid, int cmd, struct shm_id *buf);`
- Returns: 0 if OK, -1 on error



Shared Memory Segment Structure

```
struct shmid_ds {
    struct ipc_perm    shm_perm;    /* see Section 15.6.2 */
    size_t             shm_segsz;    /* size of segment in bytes */
    pid_t              shm_lpid;     /* pid of last shmop() */
    pid_t              shm_cpid;     /* pid of creator */
    shmatt_t           shm_nattch;   /* number of current attaches */
    time_t             shm_atime;     /* last-attach time */
    time_t             shm_dtime;     /* last-detach time */
    time_t             shm_ctime;     /* last-change time */
    :
};
```



Attaching to a process

- `#include <sys/types.h>`
- `#include <sys/ipc.h>`
- `#include <sys/shm.h>`
- `void *shmat (int shmid, void *addr, int flag);`
- Returns: **pointer to shared memory segment** if OK, -1 on error
- `addr` is NULL, the system chooses a suitable (unused) address at which to attach the segment.
- `addr` isn't NULL and `SHM_RND` is specified in `flag`, the attach occurs at the address equal to `addr` rounded down to the nearest multiple of `SHMLBA`.



Detaching from a process

- `#include <sys/types.h>`
- `#include <sys/ipc.h>`
- `#include <sys/shm.h>`
- `int shmdt (void *addr);`
- Returns: 0 if OK, -1 on error



Program 15.31: data storage

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#include "apue.h"

#define  ARRAY_SIZE      40000
#define  MALLOC_SIZE     100000
#define  SHM_SIZE        100000
#define  SHM_MODE        (SHM_R | SHM_W) /* user read/write */

char    array[ARRAY_SIZE];      /* uninitialized data = bss */
int main(void) {
    int    shmid;
    char  *ptr, *shmptr;
```

```

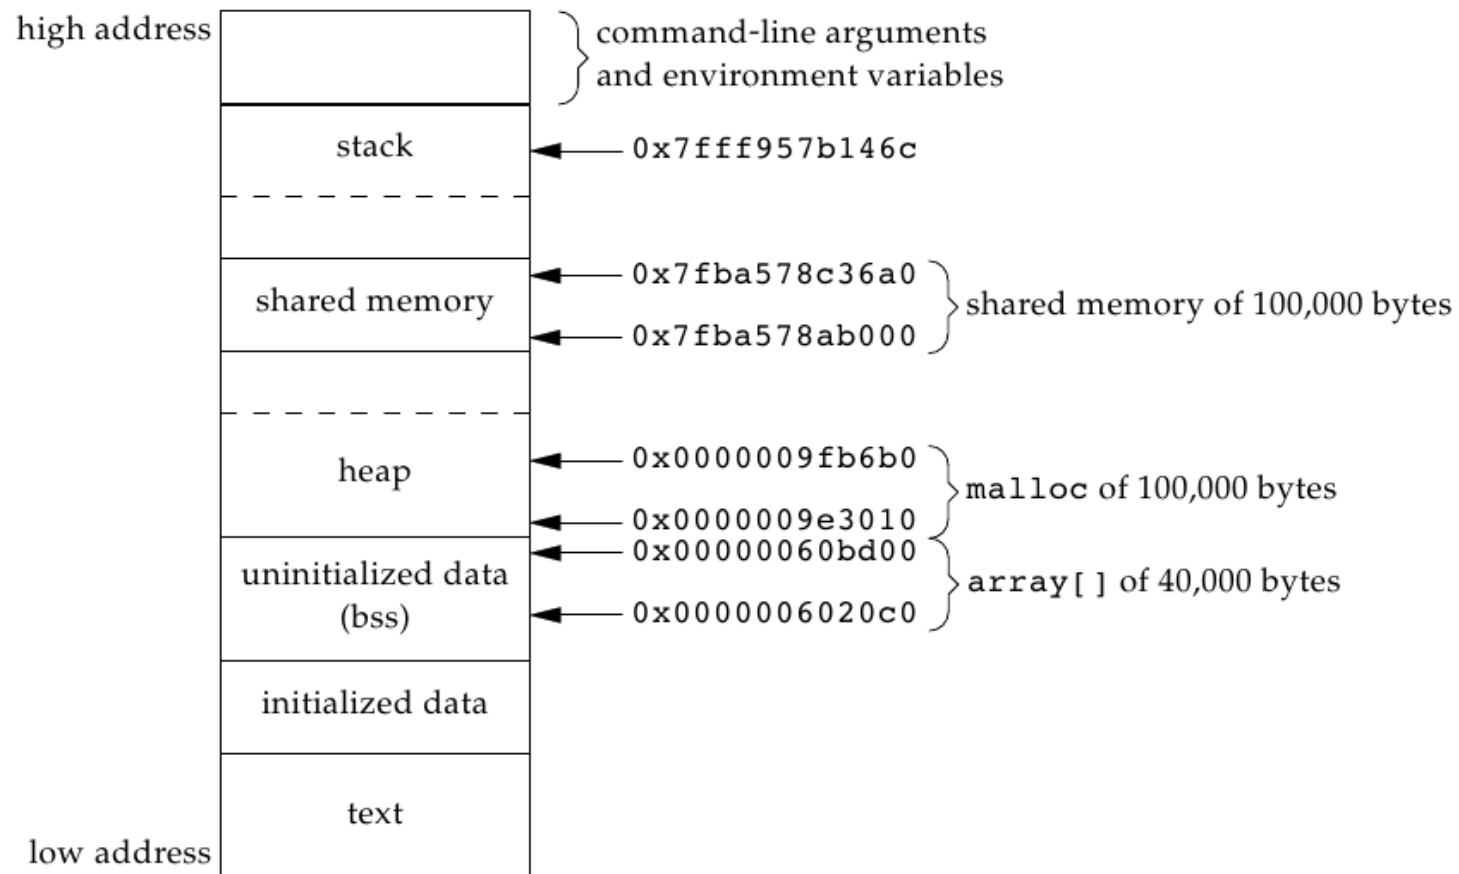
printf("array[] from %x to %x\n", &array[0], &array[ARRAY_SIZE]);
printf("stack around %x\n", &shmid);
if ( (ptr = malloc(MALLOC_SIZE)) == NULL)
    err_sys("malloc error");
printf("malloced from %x to %x\n", ptr, ptr+MALLOC_SIZE);

if ( (shmid = shmget(IPC_PRIVATE, SHM_SIZE, SHM_MODE)) < 0)
    err_sys("shmget error");
if ( (shmptr = shmat(shmid, 0, 0)) == (void *) -1)
    err_sys("shmat error");
printf("shared memory attached from %x to %x\n",
        shmptr, shmptr+SHM_SIZE);
if (shmctl(shmid, IPC_RMID, 0) < 0) err_sys("shmctl error");
exit(0);
}

```



Memory Layout (Program 15.31)





4. memory mapped


- The `mmap()` system call creates a new memory mapping in the calling process's virtual address space.
 - **File mapping:** A file mapping maps a region of a file directly into the calling process's virtual memory.
 - **Anonymous mapping:** An anonymous mapping doesn't have a corresponding file. Instead, the pages of the mapping are initialized to 0.



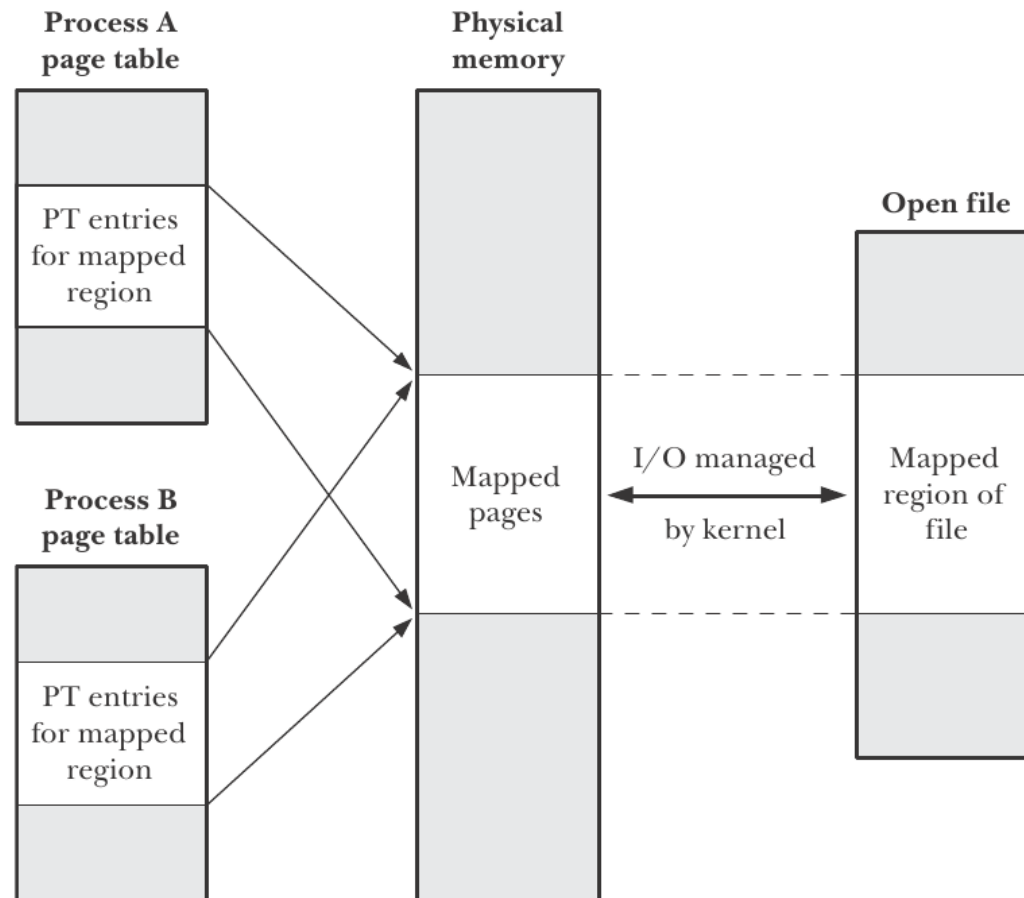
mmap

- include <sys/mman.h>
- void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset);
- Returns starting address of mapping on success, or MAP_FAILED on error
- int munmap(void *addr, size_t length);

(1) File mapping



Two processes with a **shared** mapping of the same region of a file



```
addr = mmap(NULL, MEM_SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);
```



(2) Anonymous mapping

- An anonymous mapping is one that doesn't have a corresponding file.
 - Specify `MAP_ANONYMOUS` in *flags* and specify *fd* as `-1`.
 - Open the `/dev/zero` device file and pass the resulting file descriptor to *mmap()*.




15.33 IPC between parent and child using memory mapped I/O of /dev/zero

```
#include "apue.h"
#include <fcntl.h>
#include <sys/mman.h>

#define NLOOPS      1000
#define SIZE        sizeof(long)    /* size of shared memory area */

static int
update(long *ptr)
{
    return((*ptr)++);    /* return value before increment */
}
```



```
int
main(void)
{
    int      fd, i, counter;
    pid_t    pid;
    void      *area;

    if ((fd = open("/dev/zero", O_RDWR)) < 0)
        err_sys("open error");
    if ((area = mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED,
        fd, 0)) == MAP_FAILED)
        err_sys("mmap error");
    close(fd);          /* can close /dev/zero now that it's mapped */
}
```

```

    TELL_WAIT();

    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid > 0) {                /* parent */
        for (i = 0; i < NLOOPS; i += 2) {
            if ((counter = update((long *)area)) != i)
                err_quit("parent: expected %d, got %d", i, counter);

            TELL_CHILD(pid);
            WAIT_CHILD();
        }
    } else {                            /* child */
        for (i = 1; i < NLOOPS + 1; i += 2) {
            WAIT_PARENT();

            if ((counter = update((long *)area)) != i)
                err_quit("child: expected %d, got %d", i, counter);

            TELL_PARENT(getppid());
        }
    }

    exit(0);
}

```



Anonymous Memory Mapping

- To modify the program in Figure 15.33 to use this facility, we make three changes:
- (a) remove the open of /dev/zero
- (b) remove the close of fd
- (c) change the call to mmap to the following:

```
if ((area = mmap(0, SIZE, PROT_READ | PROT_WRITE,  
    MAP_ANON | MAP_SHARED, -1, 0)) == MAP_FAILED)
```

5. Posix IPC

Interface	Message queues	Semaphores	Shared memory
Header file	<code><mqqueue.h></code>	<code><semaphore.h></code>	<code><sys/mman.h></code>
Object handle	<code>mqd_t</code>	<code>sem_t *</code>	<code>int</code> (file descriptor)
Create/open	<code>mq_open()</code>	<code>sem_open()</code>	<code>shm_open() + mmap()</code>
Close	<code>mq_close()</code>	<code>sem_close()</code>	<code>munmap()</code>
Unlink	<code>mq_unlink()</code>	<code>sem_unlink()</code>	<code>shm_unlink()</code>
Perform IPC	<code>mq_send()</code> , <code>mq_receive()</code>	<code>sem_post()</code> , <code>sem_wait()</code> , <code>sem_getvalue()</code>	operate on locations in shared region
Miscellaneous operations	<code>mq_setattr()</code> —set attributes <code>mq_getattr()</code> —get attributes <code>mq_notify()</code> —request notification	<code>sem_init()</code> —initialize unnamed semaphore <code>sem_destroy()</code> —destroy unnamed semaphore	(none)



Comparison of System V IPC and POSIX IPC

- POSIX IPC has the following general **advantages** when compared to System V IPC:
 - ✓ **interface** is simpler
 - ✓ the use of **names** instead of keys, and the open, close, and unlink functions—is more consistent with the traditional UNIX file model
 - ✓ POSIX IPC objects are **reference counted**. This simplifies object deletion, because we can unlink a POSIX IPC object, knowing that it will be destroyed only when all processes have closed it.



Comparison of System V IPC and POSIX IPC

- However, there is one notable advantage in favor of System V IPC: [portability](#).
- System V IPC is specified in SUSv3 and supported on nearly every UNIX implementation. By contrast, each of the POSIX IPC mechanisms is an optional component in SUSv3. Some UNIX implementations don't support (all of) the POSIX IPC mechanisms.